

C++ now

2025

# C++ Memory Safety in WebKit

Geoffrey Garen

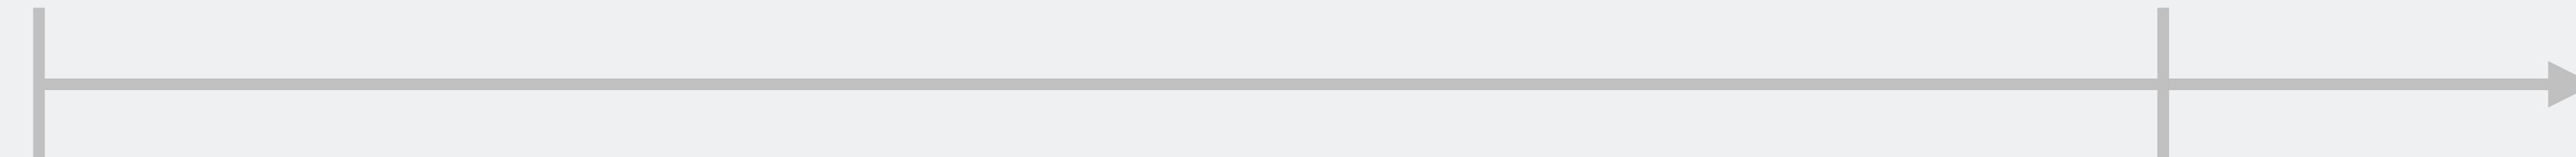
# C++ Memory Safety in WebKit

Geoffrey Garen  
WebKit Architect at Apple

[webkit.slack.com](https://webkit.slack.com) | ggaren

# The Giant Rewrite Proposal

CPPNow



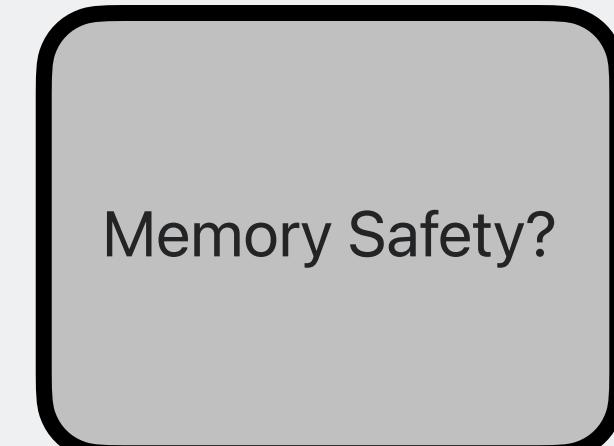
2025

2075

# The Isolated Component Proposal



My Giant C++ Project



Memory Safety?

# WebKit



# WebKit



## Platforms

- Apple platforms
- Linux
- Windows
- Blender

# WebKit



## Composition

- 9MM+ LoC
- Mostly C++

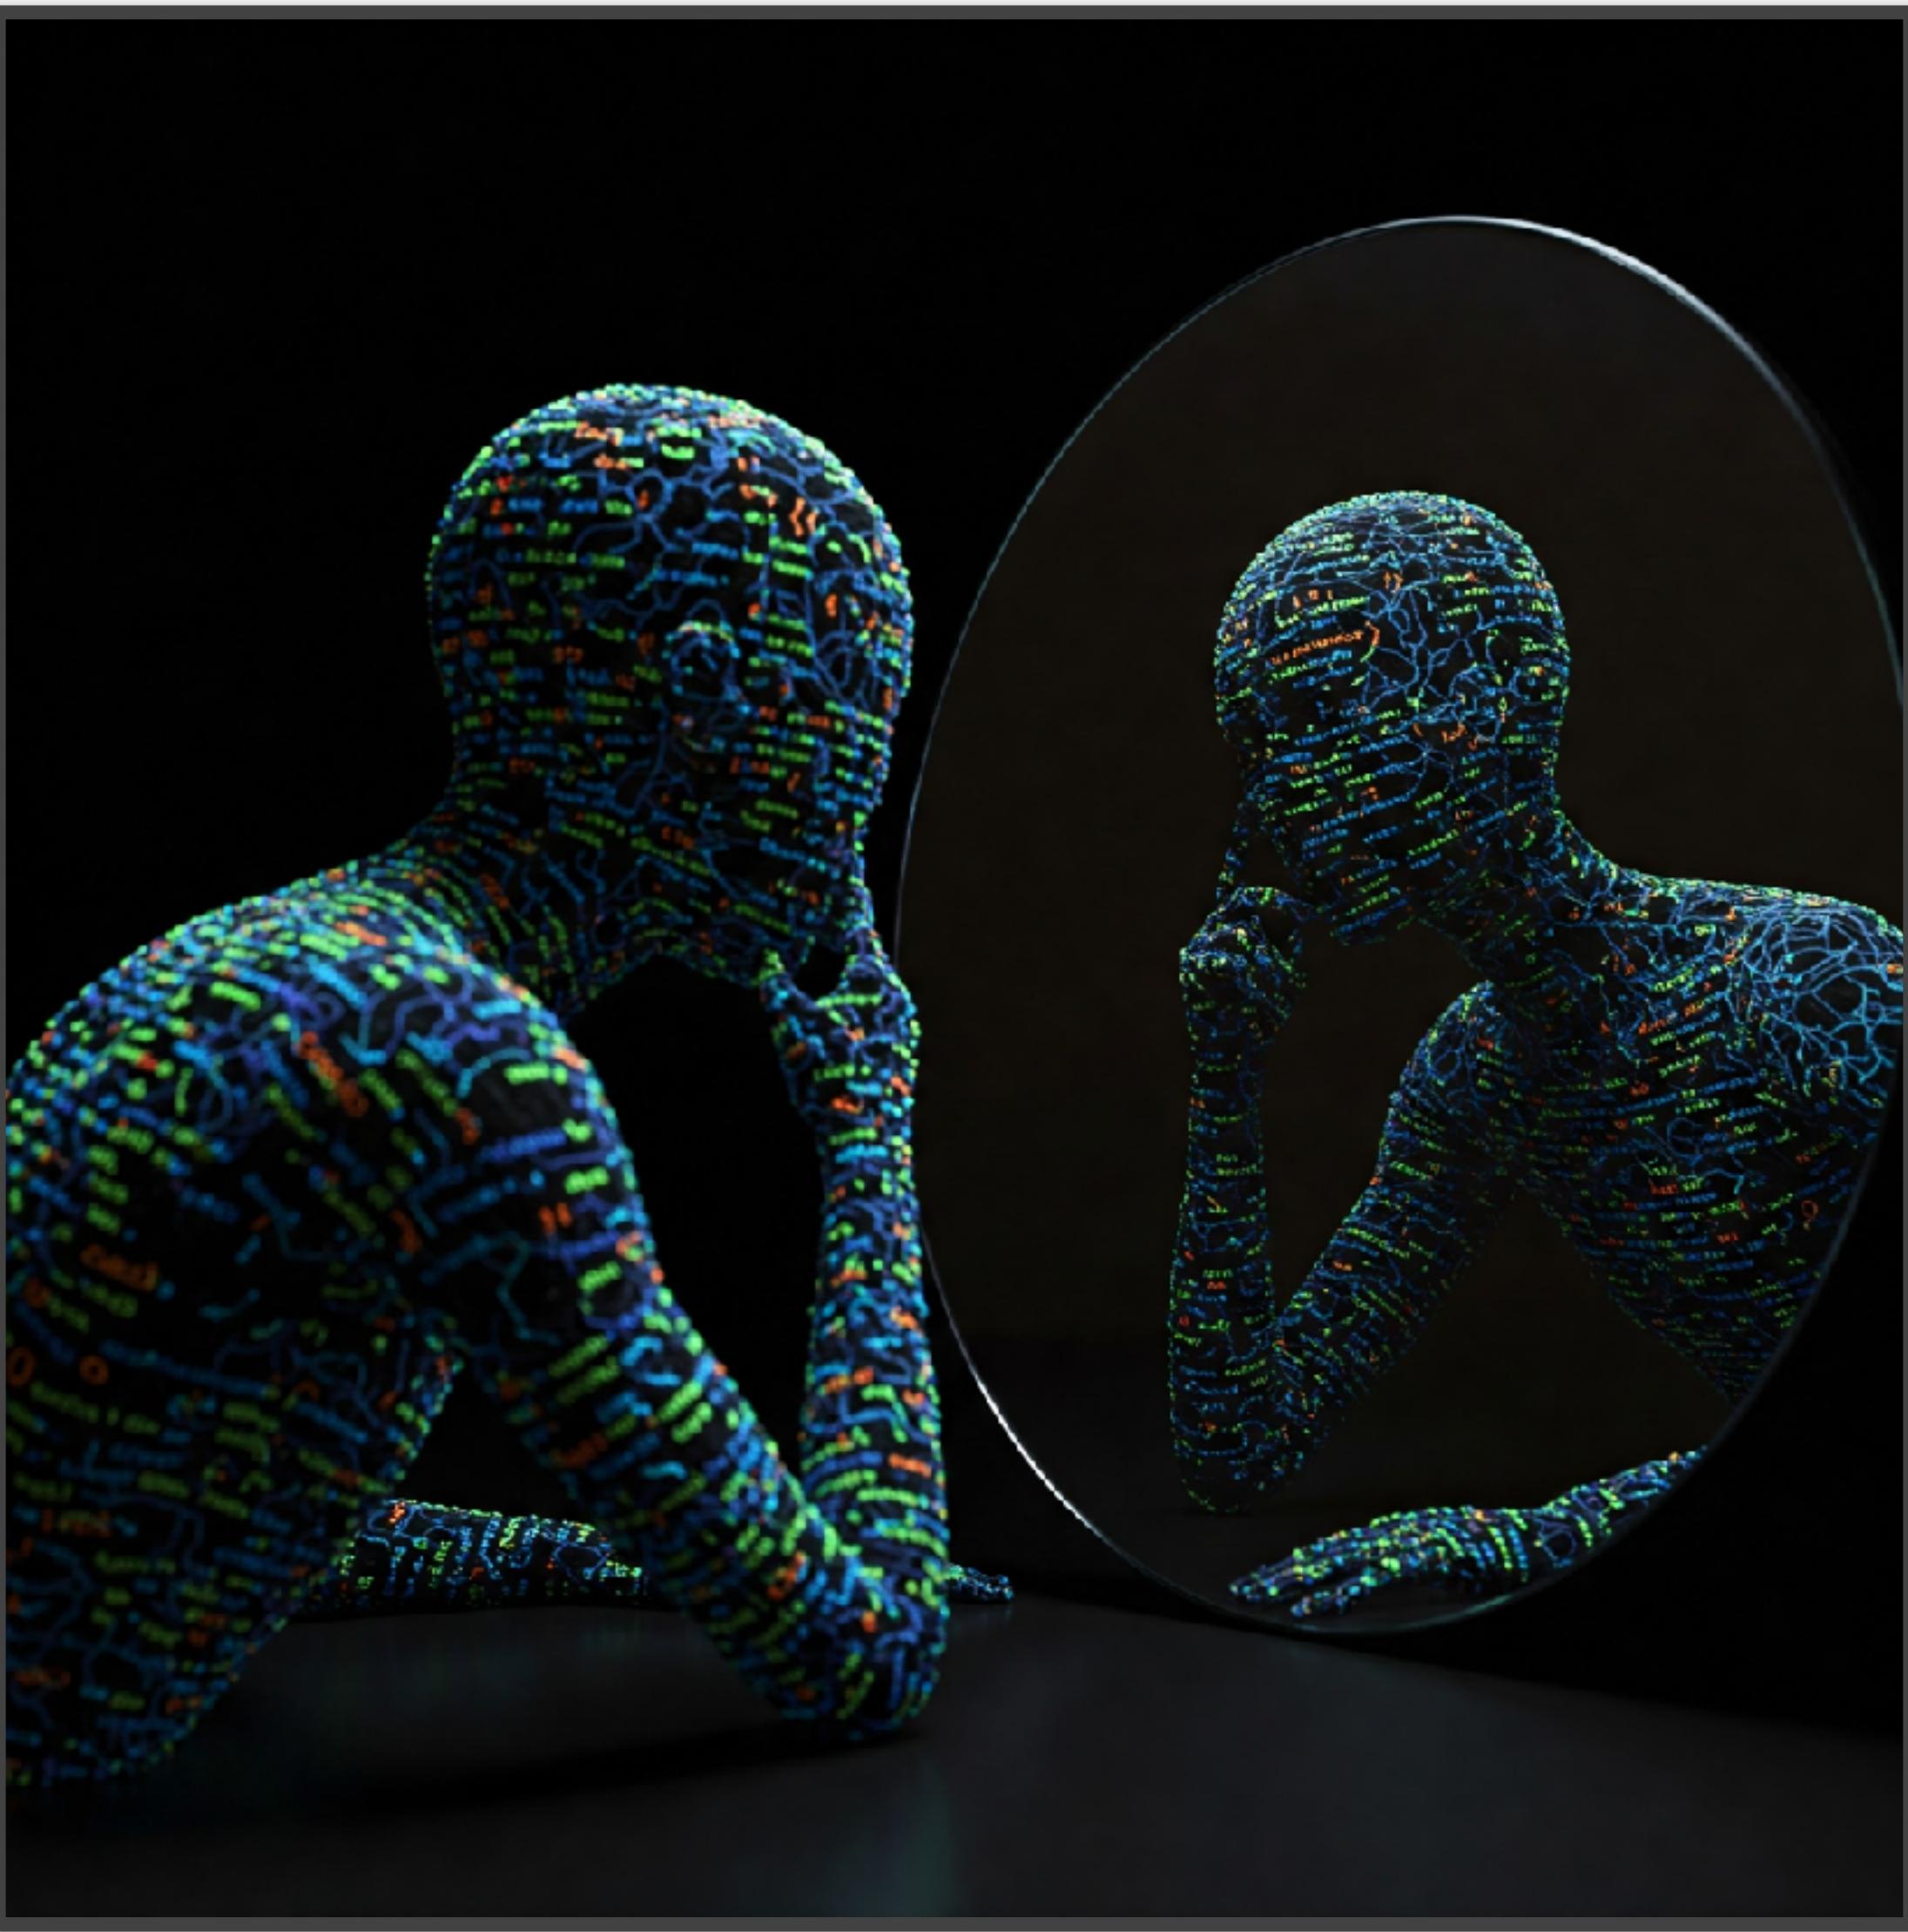
# WebKit



## Goals

- Security
- Performance

C++ <=> C++



# Test Environment



# Agenda

Project Maintainer



Standards Participant

Compiler Author

# Memory Un-Safety

*A bug that can create a weird machine*



# Memory Un-Safety – Categories

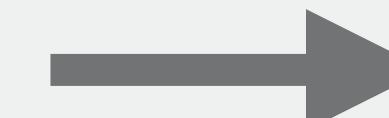
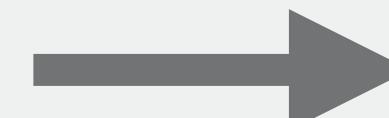
- Pointer Bugs
  - Lifetime
  - Bounds
  - Type

# Memory Un-Safety – Categories

- Undefined behavior 😬
  - Null pointer dereference
  - Infinite loop
  - etc. (~175 cases)

# Memory Safety – Ideal Definition

*!(A bug that can create a weird machine)*



# Memory Safety – Practical Definition

*A programming paradigm that rules out memory un-safety — incrementally, at compile time*

# Bounds

```
// Unsafe pointer arithmetic

void ShadowBlur::blurLayerImage(
    unsigned char* imageData, const IntSize& size, int rowStride) {
    ...
    limit = (side1 < dim) ? side1 : dim;
    for (i = 0; i < limit; ptr += stride, next += stride, ++i, ++ofs) {
        *ptr = (sum * invCount) >> blurSumShift;
        sum += ((ofs < dim) ? *next : alpha2) - alpha1;
    }
    ...
}
```

```
// Unsafe pointer arithmetic [-Wunsafe-buffer-usage, -Werror]

void ShadowBlur::blurLayerImage(
    unsigned char* imageData, const IntSize& size, int rowStride) {
    ...
    limit = (side1 < dim) ? side1 : dim;
    for (i = 0; i < limit; ptr += stride, next += stride, ++i, ++ofs) {
        ^
        ^
        ^

Error: unsafe pointer arithmetic [-Wunsafe-buffer-usage, -Werror]

    *ptr = (sum * invCount) >> blurSumShift;
    sum += ((ofs < dim) ? *next : alpha2) - alpha1;
}
...
}
```

```
// Unsafe pointer arithmetic - FIXED

void ShadowBlur::blurLayerImage(
    std::span<char> imageData, const IntSize& size, int rowStride) {
    ...
}
```

```
// Unsafe library function
```

```
memcpy(mtlBuffer.contents, ipc, tooManyBytes);
```

^ ^ ^

```
Error: function 'memcpy' is unsafe [-Wunsafe-buffer-usage-in-libc-call, -Werror]
```

```
// API Boundaries

template<typename T, std::size_t Extent = std::dynamic_extent>
inline constexpr auto unsafeMakeSpan(T* ptr, size_t size)
{
    ALLOW_UNSAFE_BUFFER_USAGE_BEGIN
        return std::span<T, Extent> { ptr, size };
    ALLOW_UNSAFE_BUFFER_USAGE_END
}

template<typename T, std::size_t TExtent, typename U, std::size_t UExtent>
void memcpySpan(std::span<T, TExtent> destination, std::span<U, UExtent> source)
{
    RELEASE_ASSERT(destination.size() >= source.size());
    ALLOW_UNSAFE_BUFFER_USAGE_BEGIN
        memcpy(destination.data(), source.data(), source.size_bytes());
    ALLOW_UNSAFE_BUFFER_USAGE_END
}
```

# Enforcement

- `-Wunsafe-buffer-usage`
- `-Werror`
- `CLANG_CXX_STANDARD_LIBRARY_HARDENING = extensive`
- `#pragma clang diagnostic ignored "-Wunsafe-buffer-usage"`

Adoption (1 Year)

94%

# Performance Cost

0%  
In Benchmarks  
End-to-End  
After Optimization

Type

```
// Unsafe Downcast

class ScrollingStateFrameScrollingNode final : public ScrollingStateScrollingNode {...};

auto& frameScrollingNode = static_cast<ScrollingStateFrameScrollingNode>(*node);
```

```
// Unsafe Downcast – FIXED

class ScrollingStateFrameScrollingNode final : public ScrollingStateScrollingNode {...};

auto& frameScrollingNode = downcast<ScrollingStateFrameScrollingNode>(*node);
```

# Enforcement

A screenshot of a code editor window titled "build.webkit.org". The code is written in C++ and defines a class `GeolocationController` with various methods and member variables. A specific line of code is highlighted with a yellow background and a tooltip: `static GeolocationController\* from(Page\* page) { return static\_cast<GeolocationController\*>(Supplement<Page>::from(page, supplementName())); }`. A tooltip above the cursor reads "Unsafe cast from base type 'Supplement' to derived type 'GeolocationController'". The code includes several annotations in green text, such as `WEBCORE\_EXPORT` and comments explaining variable types.

```
53     void removeObserver(Geolocation&);  
54  
55     void requestPermission(Geolocation&);  
56     void cancelPermissionRequest(Geolocation&);  
57  
58     WEBCORE_EXPORT void positionChanged(const std::optional<GeolocationPositionData>&);  
59     WEBCORE_EXPORT void errorOccurred(GeolocationError&);  
60  
61     std::optional<GeolocationPositionData> lastPosition();  
62  
63     GeolocationClient& client();  
64  
65     WEBCORE_EXPORT static ASCIIILiteral supplementName();  
66     static GeolocationController* from(Page* page) { return static_cast<GeolocationController*>(Supplement<Page>::from(page, supplementName())); } Unsafe cast from base type 'Supplement' to derived type 'GeolocationController'  
67  
68     void revokeAuthorizationToken(const String&);  
69  
70     void didNavigatePage();  
71  
72 private:  
73     WeakRef<Page> m_page;  
74     CheckedPtr<GeolocationClient> m_client; // Only becomes null in the class destructor  
75  
76     void activityStateDidChange(OptionSet<ActivityState> oldActivityState, OptionSet<ActivityState> newActivityState) override;  
77  
78     std::optional<GeolocationPositionData> m_lastPosition;  
79  
80     bool needsHighAccuracy() const { return !m_highAccuracyObservers.isEmpty(); }  
81  
82     void startUpdatingIfNecessary();  
83     void stopUpdatingIfNecessary();  
84  
85     typedef HashSet<Ref<Geolocation>> ObserversSet;  
86     // All observers; both those requesting high accuracy and those not.  
87     ObserversSet m_observers;  
88     ObserversSet m_highAccuracyObservers;  
89  
90     // While the page is not visible, we pend permission requests.  
91     HashSet<Ref<Geolocation>> m_pendingPermissionRequest;  
92  
93     RegistrableDomain m_registrableDomain;  
94     bool m_isUpdating { false };  
95 };  
96
```

# Enforcement

- Uses libAnalysis (clang)
- Analyzes the AST (no data flow)
- Enforced in CI (not at desk)
- Enforced incrementally (per-file skip lists and per-line suppress annotations)

# Code That's Already “Safe”

commit ff08cabb2102fb7af646608a942999bd03f94e44

Author: Chris Dumez <cdumez@apple.com>

Date: Tue Feb 11 22:07:50 2025 -0800

Address some of the Safer CPP warnings in DNSResolveQueueCFNet.cpp

[https://bugs.webkit.org/show\\_bug.cgi?id=287469](https://bugs.webkit.org/show_bug.cgi?id=287469)

Reviewed by Darin Adler.

```
- if (address->sa_family == AF_INET)
-     return WebCore::IPAddress { reinterpret_cast<const sockaddr_in*>(address)-
>sin_addr };
+     if (auto* addressV4 = dynamicCastToIPv4SocketAddress(*address))
+         return WebCore::IPAddress { addressV4->sin_addr };
```

# Adoption (1 Year)

98%

# Performance Cost

0%

In Benchmarks

End-to-End

After Optimization

# Lifetime (Reference Counted)

```
// Reference Counting Types in WebKit

// Base classes for single-threaded reference counting

template<typename T> class RefCounted;
template<typename T> class CanMakeWeakPtr;

// Smart pointers

template<typename T> class RefPtr;
template<typename T> class WeakPtr;
```

```
// Missing Reference Count

class RTCPeerConnection
  : public RefCounted<RTCPeerConnection>
  , public CanMakeWeakPtr<RTCPeerConnection> { ... };

class RTCIceTransport {
  ...
  WeakPtr<RTCPeerConnection> m_connection;
};

void RTCIceTransport::onStateChanged()
{
  ...
  if (m_connection)
    m_connection->processIceTransportStateChange();
}
```

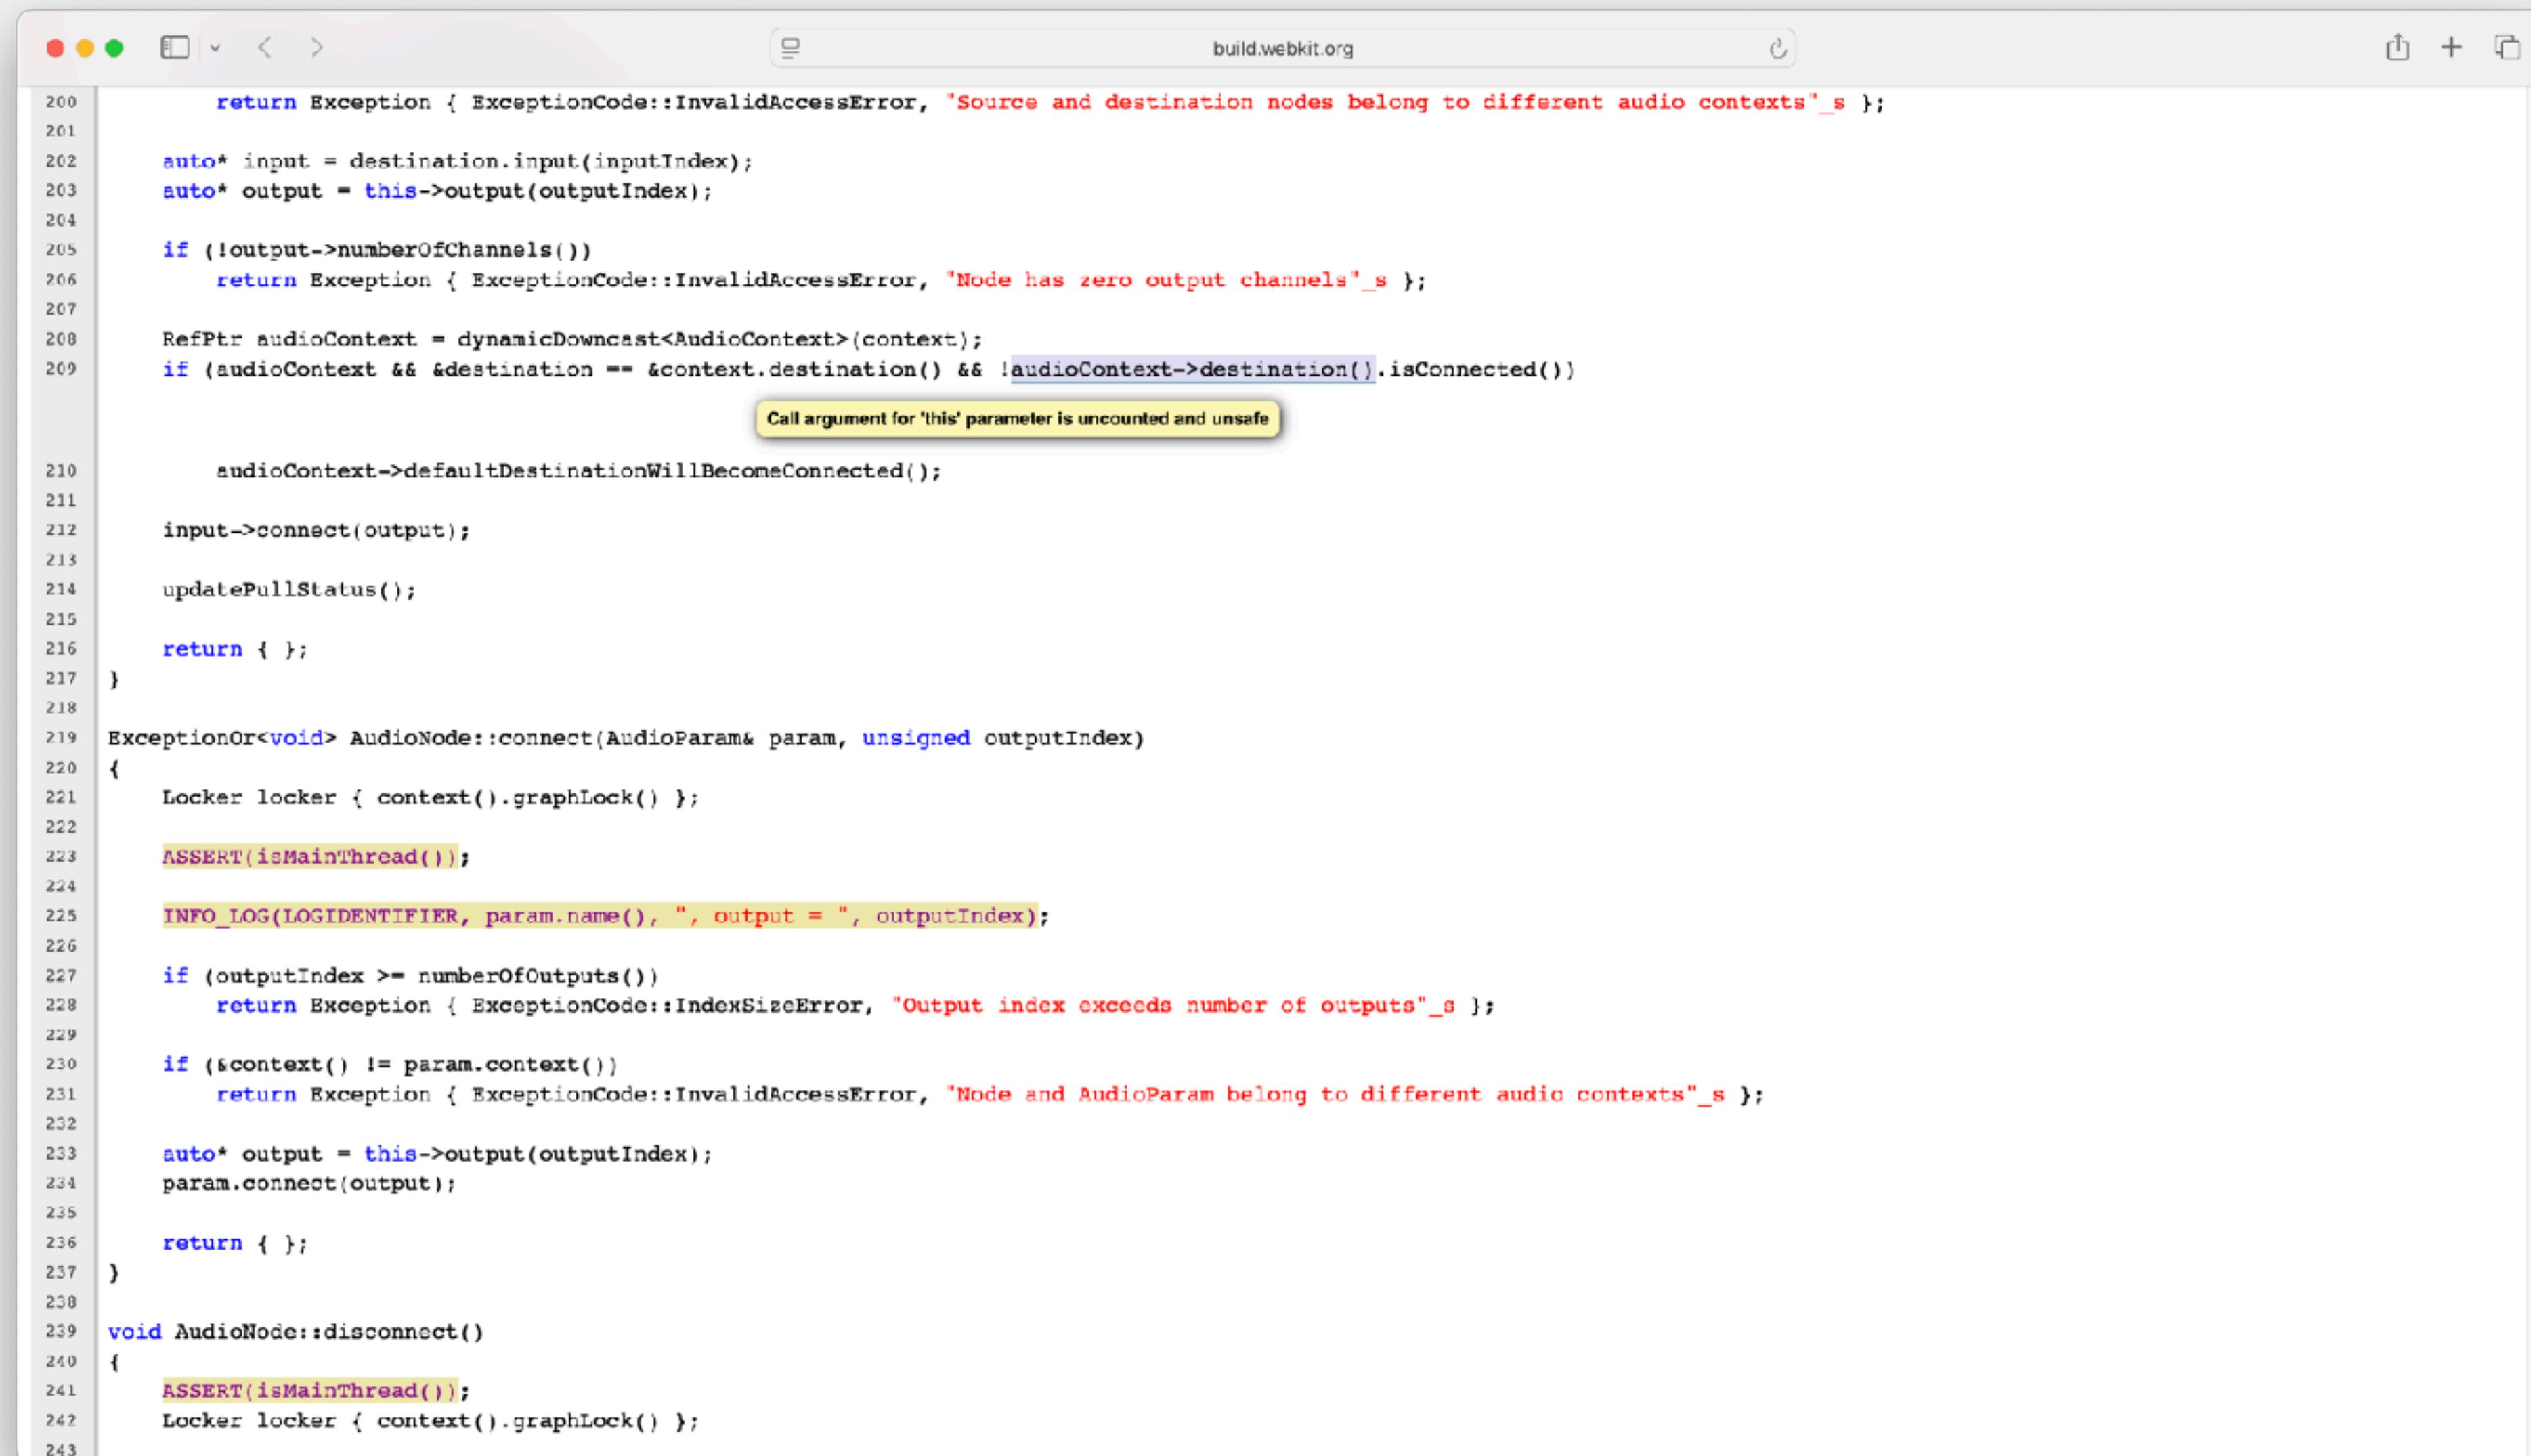
```
// Missing Reference Count – FIXED

class RTCPeerConnection
  : public RefCounted<RTCPeerConnection>
  , public CanMakeWeakPtr<RTCPeerConnection> { ... };

class RTCIceTransport {
  ...
  WeakPtr<RTCPeerConnection> m_connection;
};

void RTCIceTransport::onStateChanged()
{
  ...
  if (RefPtr connection = m_connection.get())
    connection->processIceTransportStateChange();
}
```

# Enforcement



The screenshot shows a terminal window on a Mac OS X system. The title bar reads "build.webkit.org". The window contains a block of C++ code with line numbers from 200 to 243. Several lines of code are highlighted in red, indicating errors or warnings. A yellow callout box points to line 209, which contains the code "if (audioContext && &destination == &context.destination() && !audioContext->destination().isConnected())". The callout box contains the text "Call argument for 'this' parameter is uncounted and unsafe". The code is part of the `AudioNode::connect` method, which handles audio node connections.

```
200     return Exception { ExceptionCode::InvalidAccessError, "Source and destination nodes belong to different audio contexts"_s };
201
202     auto* input = destination.input(inputIndex);
203     auto* output = this->output(outputIndex);
204
205     if (!output->numberOfChannels())
206         return Exception { ExceptionCode::InvalidAccessError, "Node has zero output channels"_s };
207
208     RefPtr<AudioContext> = dynamicDowncast<AudioContext>(context);
209     if (audioContext && &destination == &context.destination() && !audioContext->destination().isConnected())
210         audioContext->defaultDestinationWillBecomeConnected();
211
212     input->connect(output);
213
214     updatePullStatus();
215
216     return {};
217 }
218
219 ExceptionOr<void> AudioNode::connect(AudioParam& param, unsigned outputIndex)
220 {
221     Locker locker { context().graphLock() };
222
223     ASSERT(isMainThread());
224
225     INFO_LOG(LOGIDENTIFIER, param.name(), ", output = ", outputIndex);
226
227     if (outputIndex >= numberOfOutputs())
228         return Exception { ExceptionCode::IndexSizeError, "Output index exceeds number of outputs"_s };
229
230     if (&context() != param.context())
231         return Exception { ExceptionCode::InvalidAccessError, "Node and AudioParam belong to different audio contexts"_s };
232
233     auto* output = this->output(outputIndex);
234     param.connect(output);
235
236     return {};
237 }
238
239 void AudioNode::disconnect()
240 {
241     ASSERT(isMainThread());
242     Locker locker { context().graphLock() };
243 }
```

```
// Reference Count Overlooking Scope Rule

class Object : public RefCounted<Object> {
public:
    void usePointers(Object*);

private:
    const RefPtr<Object> constMember;
    RefPtr<Object> member;
    Object* rawMember; // ERROR
};

Object* getObject();

extern void use(Object*);
inline void useInline(Object*) { }

void Object::usePointers(Object* argument)
{
    use(argument);
    use(constMember.get());
    use(RefPtr { member }.get());
    use(member.get()); // ERROR
    useInline(member.get());
    use(getObject()); // ERROR
}
```

```
// Reference Count Overlooking Scope Rule

class Object : public RefCounted<Object> {
public:
    void usePointers(Object*);

private:
    const RefPtr<Object> constMember;
    RefPtr<Object> member;
    Object* rawMember; // ERROR
};

Object* getObject();

extern void use(Object*);
inline void useInline(Object*) { }

void Object::usePointers(Object* argument)
{
    RefPtr localMember = member.get();
    if (Object* rawMember = localMember.get())
        use(rawMember);

    localMember = nullptr; // ERROR
}
```

# Adoption (1 Year)

90%

# Performance Cost

0%

In Benchmarks

End-to-End

After Optimization

# Lifetime (Not Reference Counted)



```
// View type use-after-free  
  
return stringBuilder.stringView();  
  
          ^^^ string view  
          ^^^ string builder local
```

```
// View type use-after-free - FIXED

class StringBuilder {
    ...
    StringView stringView() const [[clang::lifetimebound]];
    ...
};
```

```
// View type use-after-free - FIXED

return stringBuilder.stringView();
      ^^^

error: address of stack memory associated with local variable
'stringBuilder' returned [-Werror, -Wreturn-stack-address]
```

# Enforcement

- -Wdangling (enabled by default)
- -Werror
- [[clang::lifetimebound]]

# Future Direction

- Require `[ [clang::lifetimebound] ]`
- Expand `-Wdangling`
- Smart pointer for mutation safety

Adoption (1 Year)

TBD%

# Performance Cost

0%

# Undefined Behavior



# The Friendly Neighborhood Ghost

```
RELEASE_ASSERT(destination.size() >= source.size());
```



```
// Infinite loop

uint8_t load(std::span<uint8_t> input, size_t index)
{
    if (index >= input.size()) {
        while (1) { // 🤪
    }
    }
    return input[index]; // UH OH
}
```

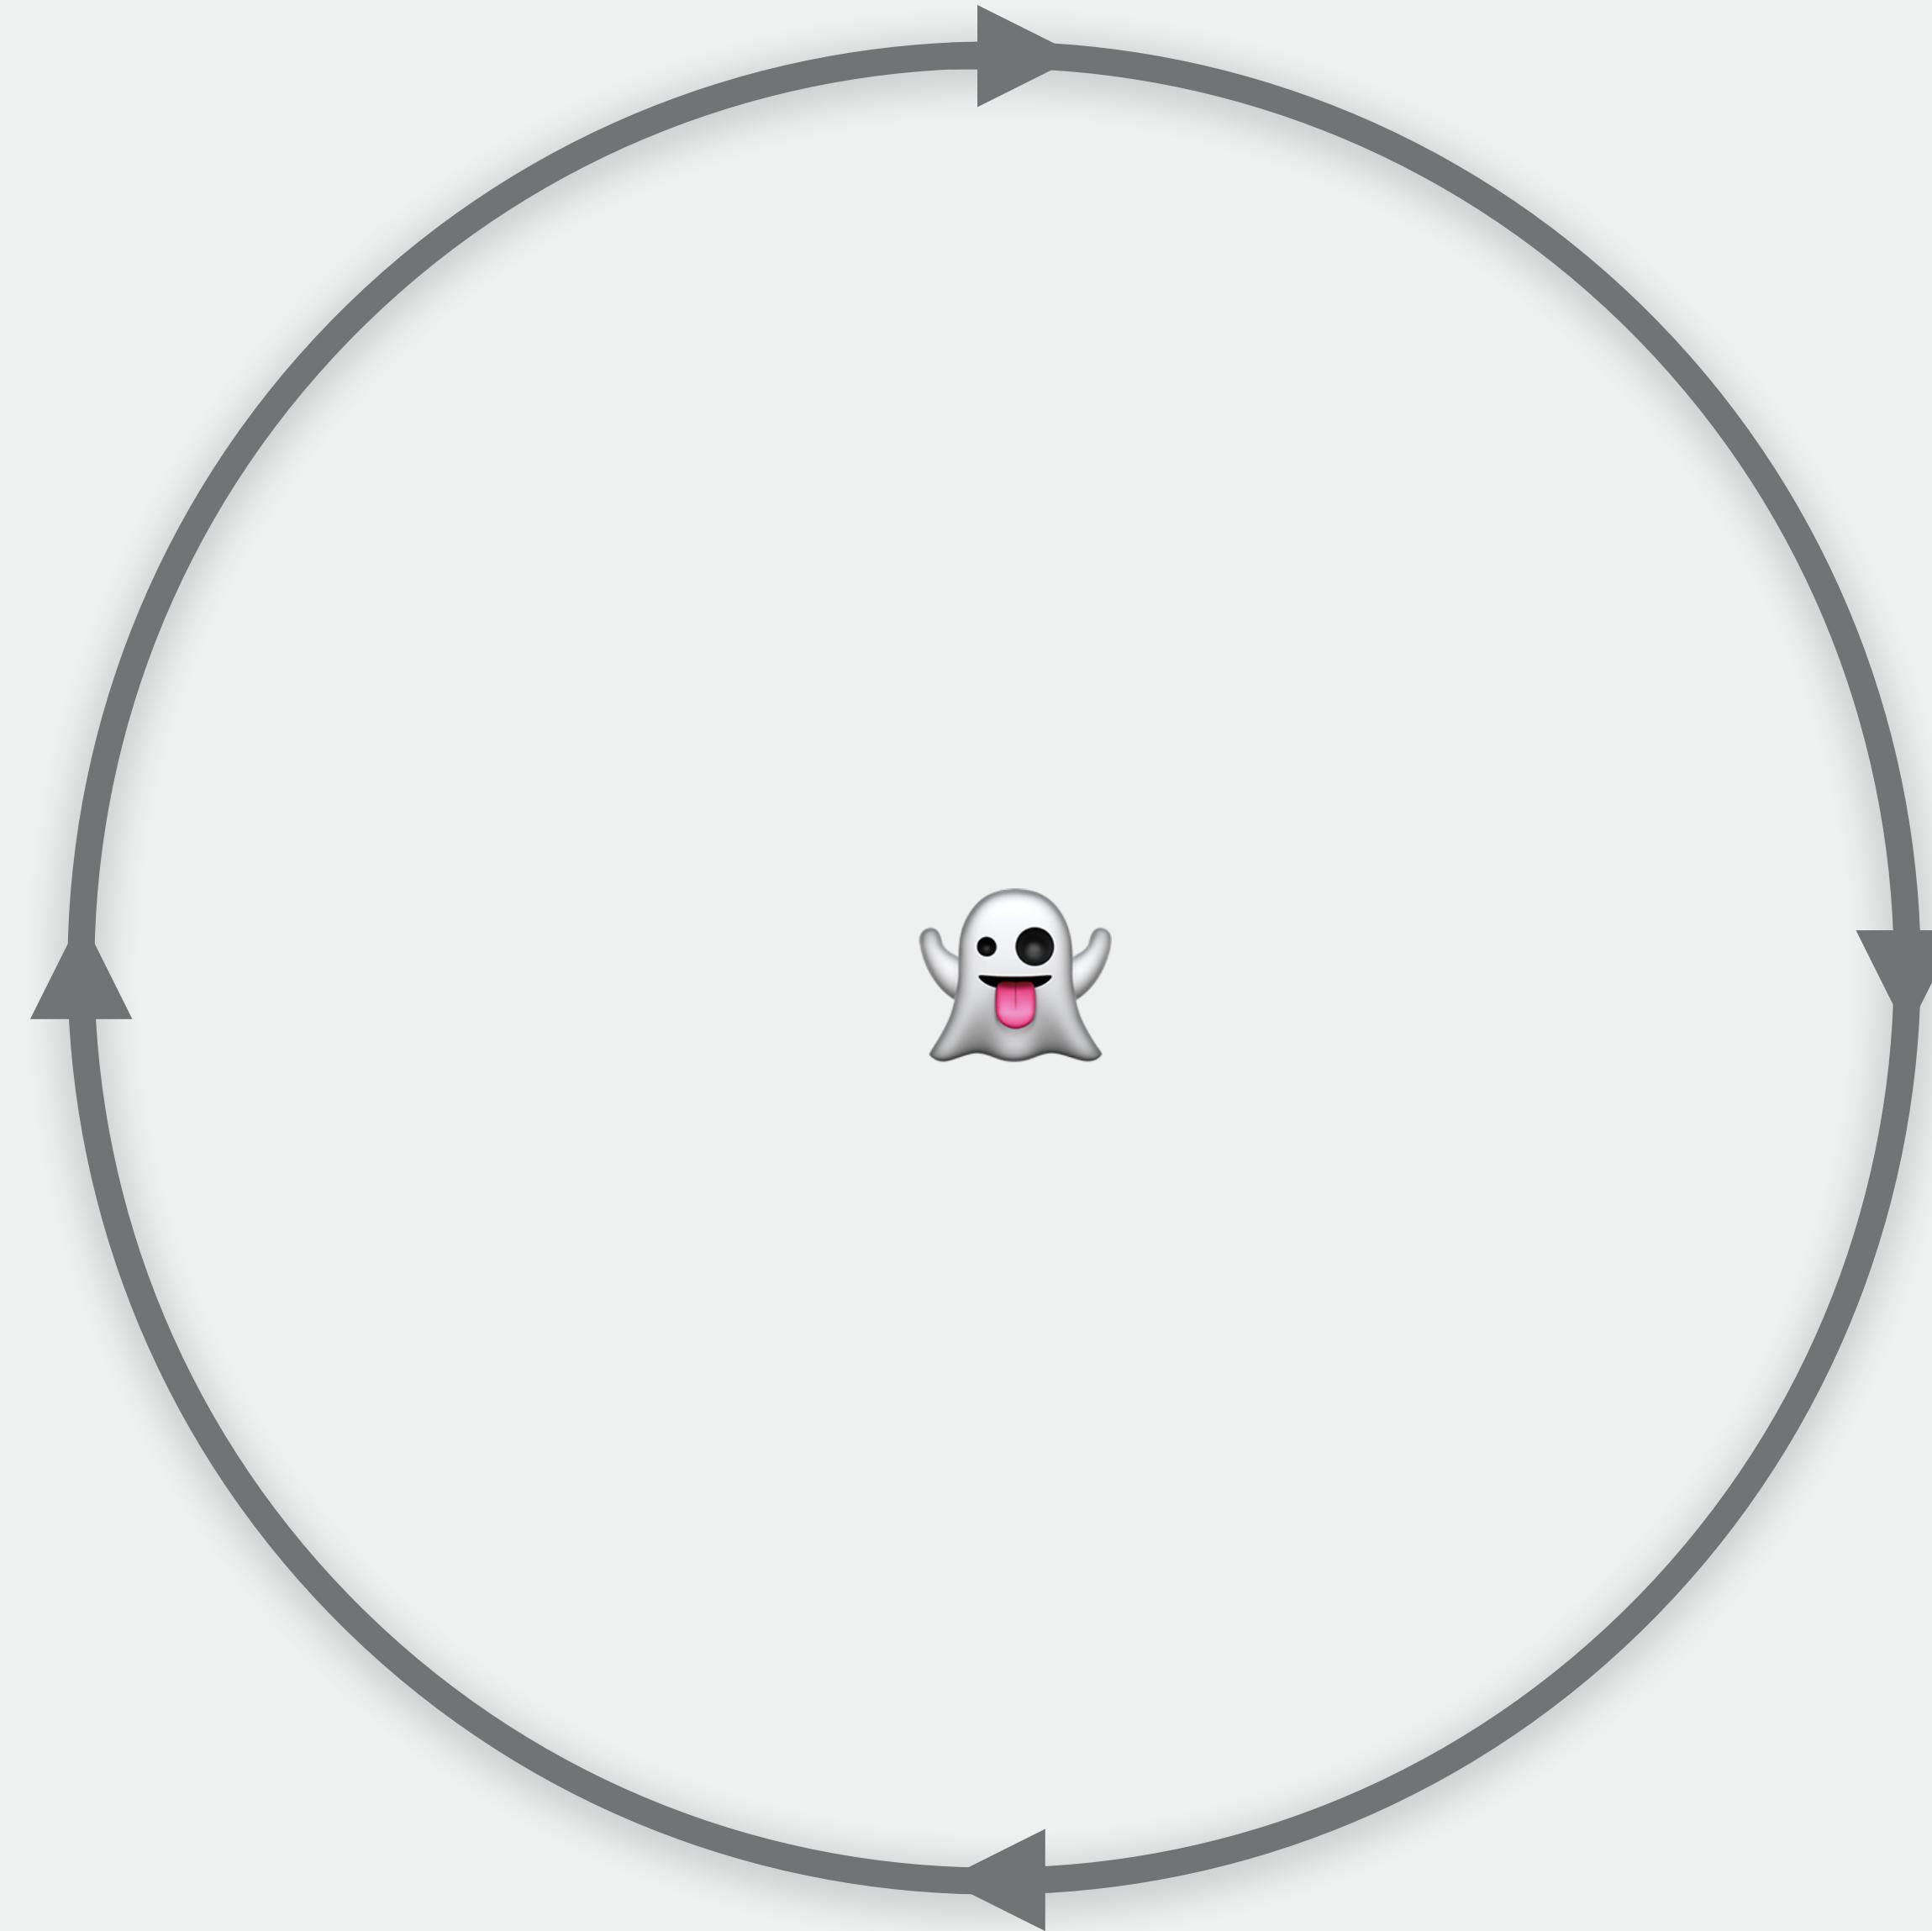
```
// Null pointer dereference

WeakPtr passwordPopup = getPasswordPopup();
if (shouldShowPasswordPopup())
    RefPtr { passwordPopup.get() /* 🧐 */ }->show(); // UH OH
```

*{ Redacted }*

Memory Un-Safety

Static Analysis



# Swift

- Direct C++ integration
- Shared concepts
- Strict Memory Safety Checking

# Shared Concepts

Safety Issue	Concept	Annotation
Bounds	Bounds-Checked Collection and View Types	N/A
Type	Checked Pointer Casts	N/A
Lifetime (Reference Counted)	Overlooking Reference Count	SWIFT_SHARED_REFERENCE
Lifetime (Not Reference Counted)	Lifetime Dependency	SWIFT_NONESCAPABLE [[clang::lifetimebound]]
Undefined Behavior	"If it's not safe, it shouldn't compile"	N/A

```
// Lifetime (Reference Counted)

// Objective-C++

template<typename T> class RefCounted { ... } SWIFT_SHARED_REFERENCE(ref, deref);

class Texture : public RefCounted<Texture> {
    ...
    id<MTLTexture> texture() const { return m_texture; }
};

// Swift

let texture = fromAPI(attachment.view)
    ^^^ C++ reference count
    ...

let mtlTexture = texture.texture()
    ^^^ Objective-C reference count
```

```
// Lifetime (Reference Counted)

// Objective-C++

template<typename T> class RefCounted { ... };

class Texture : public RefCounted<Texture> {
    ...
    id<MTLTexture> texture() const { return m_texture; }
};

// Swift

let texture = fromAPI(attachment.view)
    ^^^ error: expression uses unsafe constructs but is not marked with 'unsafe'
...
let mtlTexture = texture.texture()
    ^^^ Objective-C reference count
```

```
// Lifetime (Not Reference Counted)

// C++

class SWIFT_NONNULL StringView { ... };

class StringBuilder {
    ...
    StringView stringView() const [[clang::lifetimebound]];
    ...
};

// Swift

return stringBuilder.stringView();
          ^^^ error: lifetime-dependent value escapes local scope
```

```
// Lifetime (Not Reference Counted)

// C++

class SWIFT_NONNULL StringView { ... };

class StringBuilder {
    ...
    StringView stringView() const;
    ...
};

// Swift

return stringBuilder.stringView();
               ^^^ error: lifetime-dependent value is missing a lifetime annotation
```

```
// Lifetime (Not Reference Counted)

// C++

class StringView { ... };

class StringBuilder {
    ...
    StringView stringView() const;
    ...
};

// Swift

return stringBuilder.stringView();
^/^ error: expression uses unsafe constructs but is not marked with 'unsafe'
```

# Summary

# Summary

- Incremental progress toward memory safety can start now, everywhere
- The end-to-end performance cost in WebKit has been zero
- Undefined behavior haunts C++ 
- You can help!
- A memory safe language can be a “Yes! And...”