

C++ now

2025

C++ as a Microscope Into Hardware

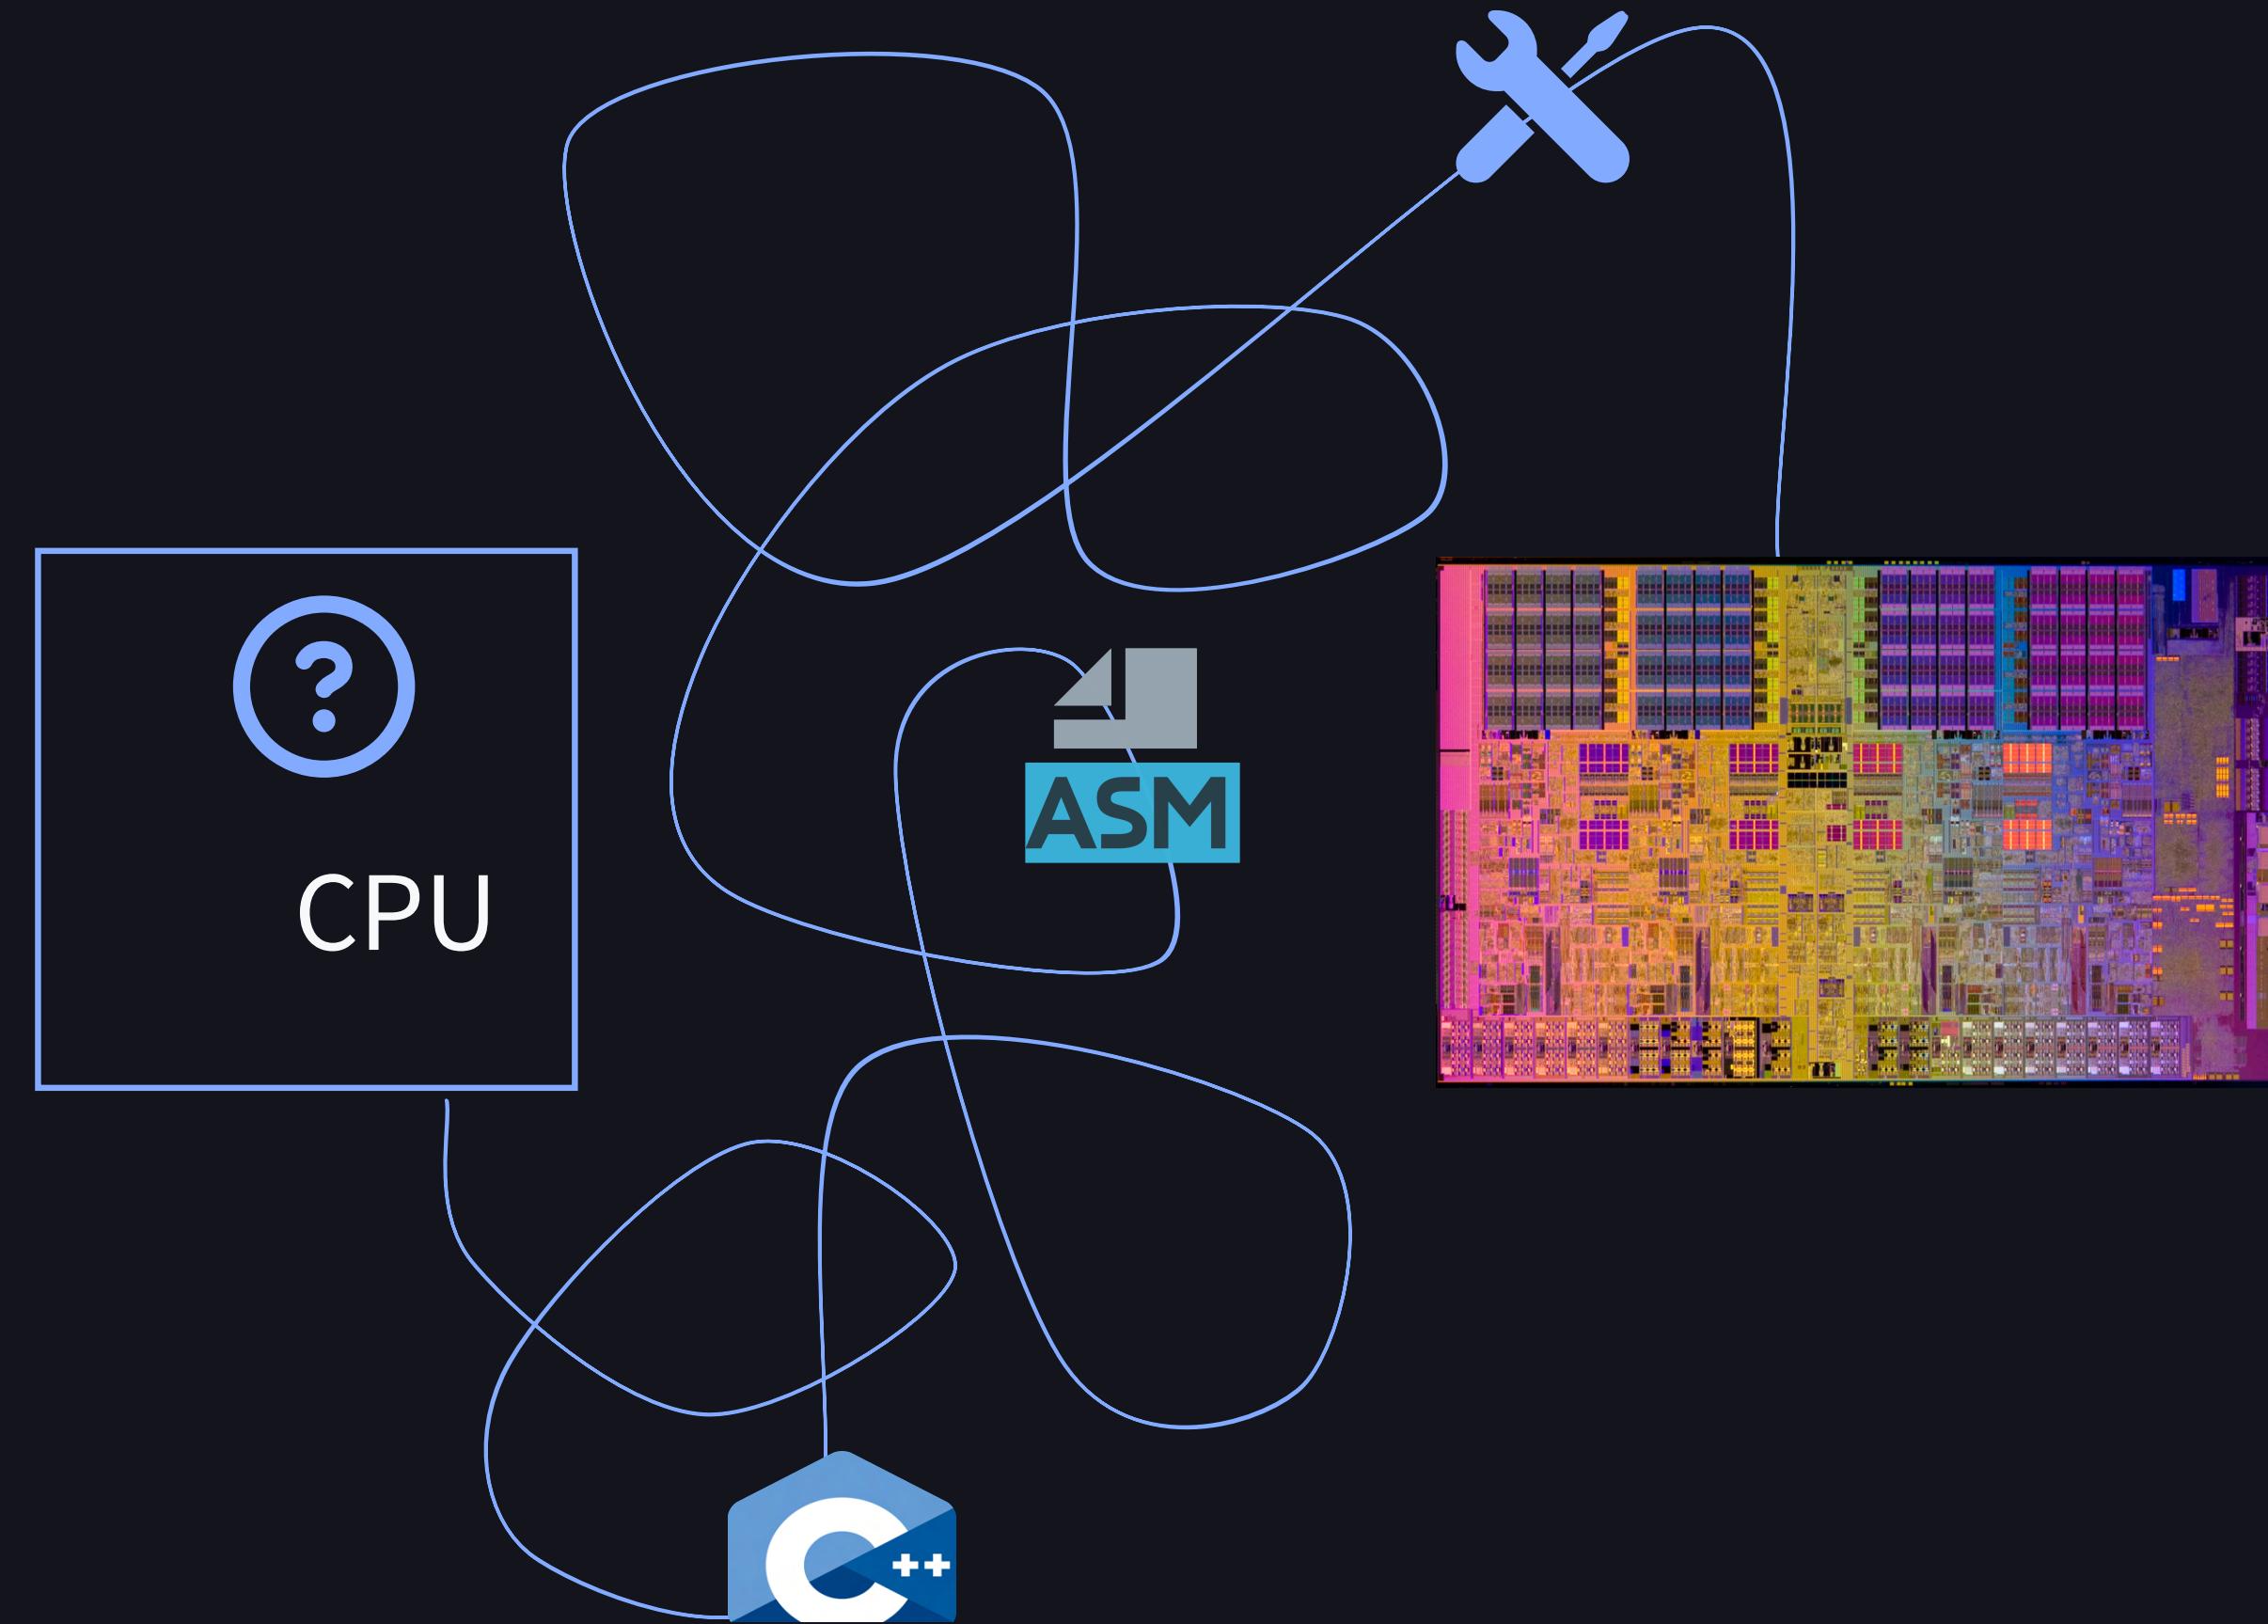
Linus Boehm



car ↔ \$

ENGINEERS	GATE
-----------	------

+ x86-64



```
int fun() {  
    return 0;  
}
```

```
g++ -g -c -O1 return_zero.cpp # -g:debug info, -c:no main, -O1:optimize  
less return_zero.o
```

```
int fun() {  
    return 0;  
}
```

```
g++ -g -c -O1 return_zero.cpp # -g:debug info, -c:no main, -O1:optimize >_  
xxd -b return_zero.o          # -b:binary dump
```

```
>> 00000000: 01111111 01000101 01001100 01000110 00000010 00000001 .ELF..  
00000006: 00000001 00000000 00000000 00000000 00000000 00000000 .....  
0000000c: 00000000 00000000 00000000 00000000 00000001 00000000 .....  
00000012: 00111110 00000000 00000001 00000000 00000000 00000000 >.....  
00000018: 00000000 00000000 00000000 00000000 00000000 00000000 .....  
0000001e: 00000000 00000000 00000000 00000000 00000000 00000000 .....  
00000024: 00000000 00000000 00000000 00000000 10000000 00000101 .....  
0000002a: 00000000 00000000 00000000 00000000 00000000 00000000 .....  
00000030: 00000000 00000000 00000000 00000000 01000000 00000000 ....@.  
00000036: 00000000 00000000 00000000 00000000 01000000 00000000 @
```

fetch → decode → execute

```
int fun() {  
    return 0;  
}
```

```
g++ -g -c -O1 return_zero.cpp # -g:debug info, -c:no main, -O1:optimize  
objdump -d -C -Mintel return_zero.o # -d:disassemble,-Mintel:intel syntax >-
```

```
» return_zero.o:      file format elf64-x86-64  
  
Disassembly of section .text:  
  
0000000000000000 <fun():>  
 0:  b8 00 00 00 00      mov     eax,0x0  
 5:  c3                  ret
```

fetch → decode → execute

Registers / Instructions

```
int many_args(int, int, int, int,
              int, int, int num7) {
    return num7;
}
```



```
/////////-01
<many_args(int, int, int, int, int, int, int)>:
8b 44 24 08      mov    eax,DWORD PTR [rsp+0x8]
c3                ret
```



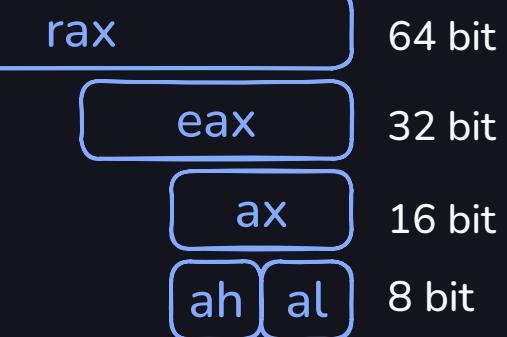
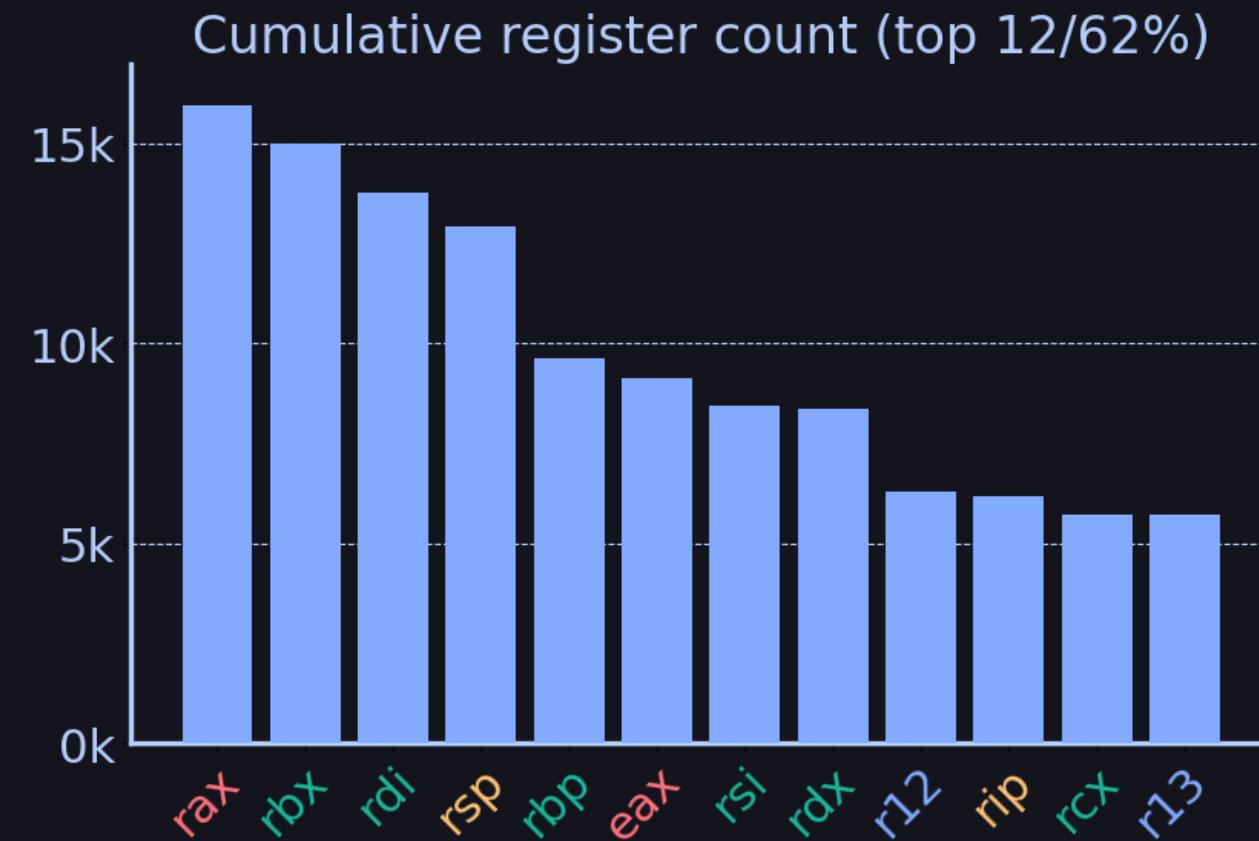
- Variable instruction length ([Combined Intel 64 Manuals](#))
- `eax / rax`: return register (32/64 bit)
- Volatile/caller owned
- `rdi , rsi , ...` : function arguments
- Calling convention: [System V AMD64 ABI](#)
- Memory operands `mov eax,DWORD PTR [rsp+0x8]`

- Questions
 - How many instructions/registers are there?
 - What are common ones?

/usr/bin/gcc registers

```
# disassemble gcc
objdump -d /usr/bin/gcc
```

» total registers: 186k



- rax Return
- rbp Base pointer
- rip Instruction pointer
- rsp Stack pointer
- rflags Flags register (Carry, sign, ...)

Function calling convention



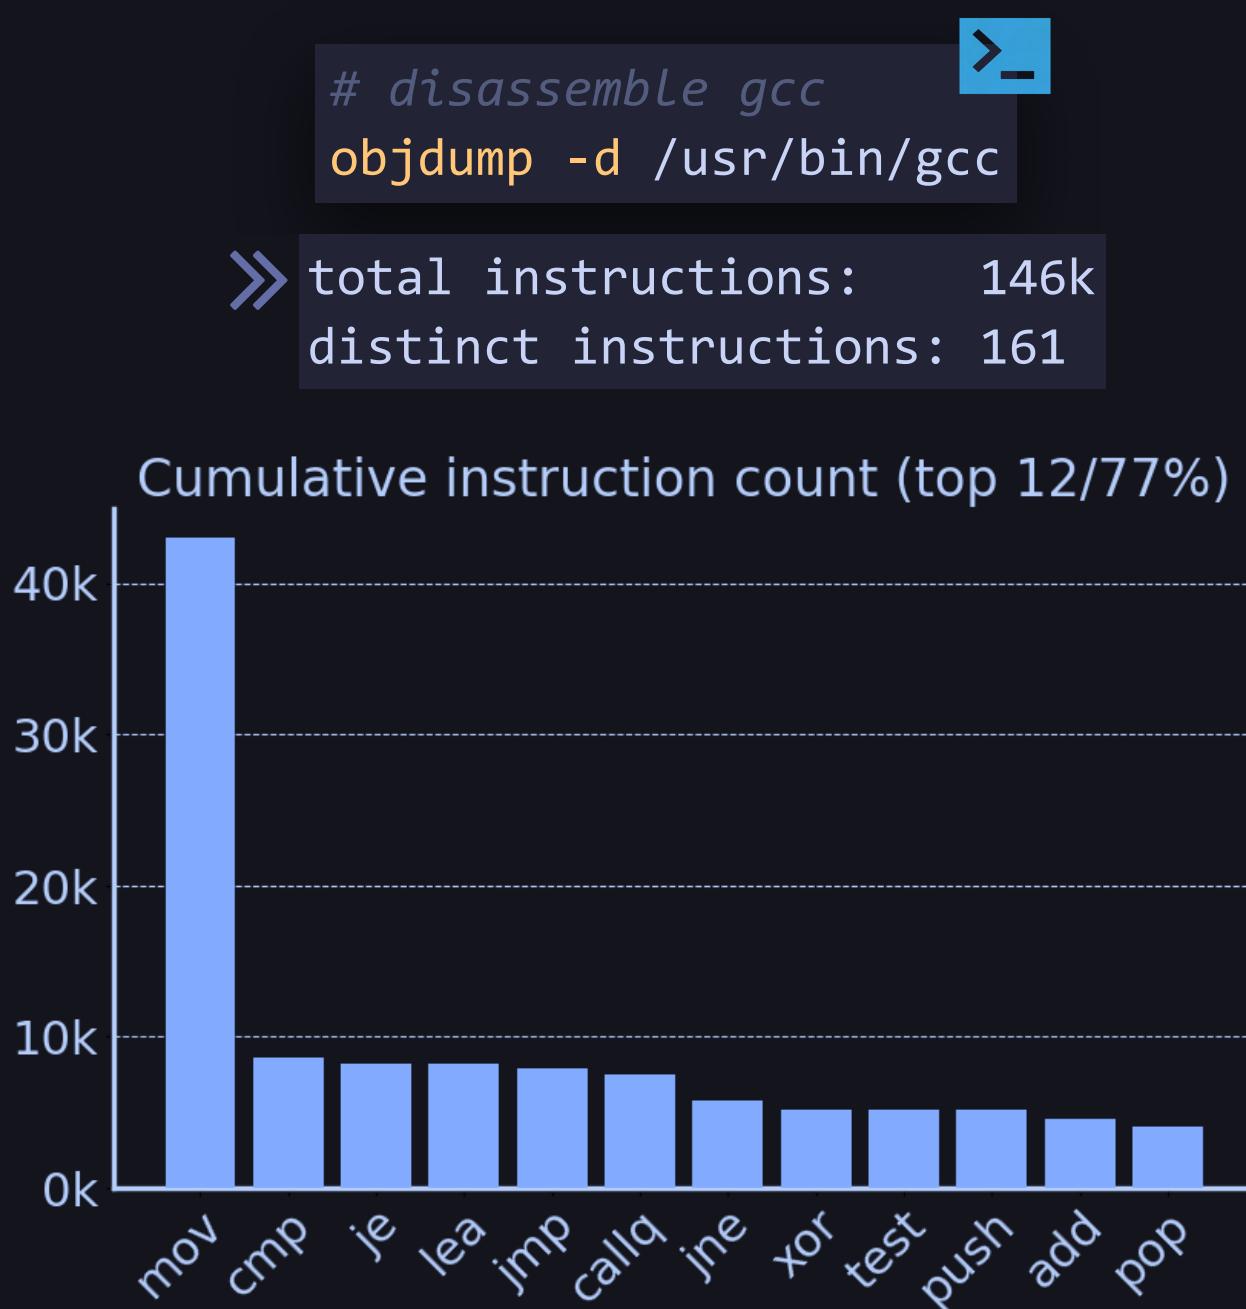
Scratch/volatile/callee owned



Caller owned (restored by callee)



/usr/bin/gcc instructions



- `mov` : copy data (e.g. into registers for processing)
- `cmp` : compare numbers (if statements, loops)
- `je` : jump if previous comp values were equal (loops)
- `lea` : fancy way of doing arithmetic (adding)
- `callq` : function call
- `jmpq` : unconditional jump
- `jne` : jump not equal
- `xor` : xor (e.g. to null out registers)
- `test` : ANDs operands and updates flags
- `push` : pushes value on stack, stack management
- `add` : adds two values
- `pop` : pops value from stack, stack management

the x86 instruction set is [...] overcomplicated, and **redundantly redundant**.

[...] it remains Turing-complete when reduced to just **one instruction**.

— Stephen Dolan, University of Cambridge, 2013 — [mov is Turing-complete](#)

movfuscator

find_primes.c:

```
1 #include <stdio.h>
2
3 #define size 8190
4 #define sizepl 8191
5 char flags[sizepl];
6
7 int main() {
8     int i, prime, k, count, iter;
9     for (iter = 1; iter <= 100; iter++) {
10         count = 0;
11         for (i = 0; i <= size; i++)
12             flags[i] = 1;
13         for (i = 0; i <= size; i++) {
14             if (flags[i]) {
15                 prime = i + i + 3;
16                 k = i + prime;
17                 while (k <= size) {
18                     flags[k] = 0;
19                     k += prime;
20                 }
21                 count = count + 1;
22             }
23         }
24     }
25     printf("found %d primes", count);
26 }
```

Overview

The M/o/Vfuscator (short 'o', sounds like 'move') is a compiler that generates assembly code using only mov instructions. Arithmetic, comparisons, and other operations are all performed through mov operations; there is no non-mov cheating.

Instructions, and only "mov" needs are all performed in, and no other form of

<https://github.com/xoreaxeaxeax/movfuscator>

The basic effects of the process can be seen in [overview](#), which illustrates compiling a simple prime number function with gcc and the M/o/Vfuscator.

movfuscator

```
movcc ./find_primes.c -o find_primes_mov >_
objdump -Mintel -d ./find_primes_mov
```

```
000 mov    eax, DWORD PTR [ecx*4+0x83f8160]
001 mov    edx, DWORD PTR ds:0x85f8194
002 mov    DWORD PTR [eax], edx
003 mov    DWORD PTR ds:0x83f8164,0x804f048
004 mov    eax, DWORD PTR [ecx*4+0x83f8160]
005 mov    edx, DWORD PTR ds:0x85f8198
006 mov    DWORD PTR [eax], edx
007 mov    DWORD PTR ds:0x83f8164,0x804f050
008 mov    eax, DWORD PTR [ecx*4+0x83f8160]
009 mov    edx, DWORD PTR ds:0x85f819c
010 mov    DWORD PTR [eax], edx
011 mov    edx, DWORD PTR ds:0x85f81a0
012 mov    DWORD PTR [eax+0x4], edx
013 mov    DWORD PTR ds:0x83f8164,0x804f058
014 mov    eax, DWORD PTR [ecx*4+0x83f8160]
015 mov    edx, DWORD PTR ds:0x85f81a4
```

ASM

```
gcc ./find_primes.c -o find_primes_gcc >_
objdump -Mintel -d ./find_primes_gcc
```

```
1 .LC0:
2     .string "found %d primes"
3 main:
4         sub    rsp, 8
5         mov    r9d, 100
6         mov    edi, OFFSET FLAT:flags+8191
7         jmp    .L2
8 .L5:
9         add    r8d, 1
10 .L4:
11        add    rsi, 1
12        add    rcx, 3
13        add    rdx, 2
14        cmp    rcx, 24576
15        je     .L13
16 .L7:
```

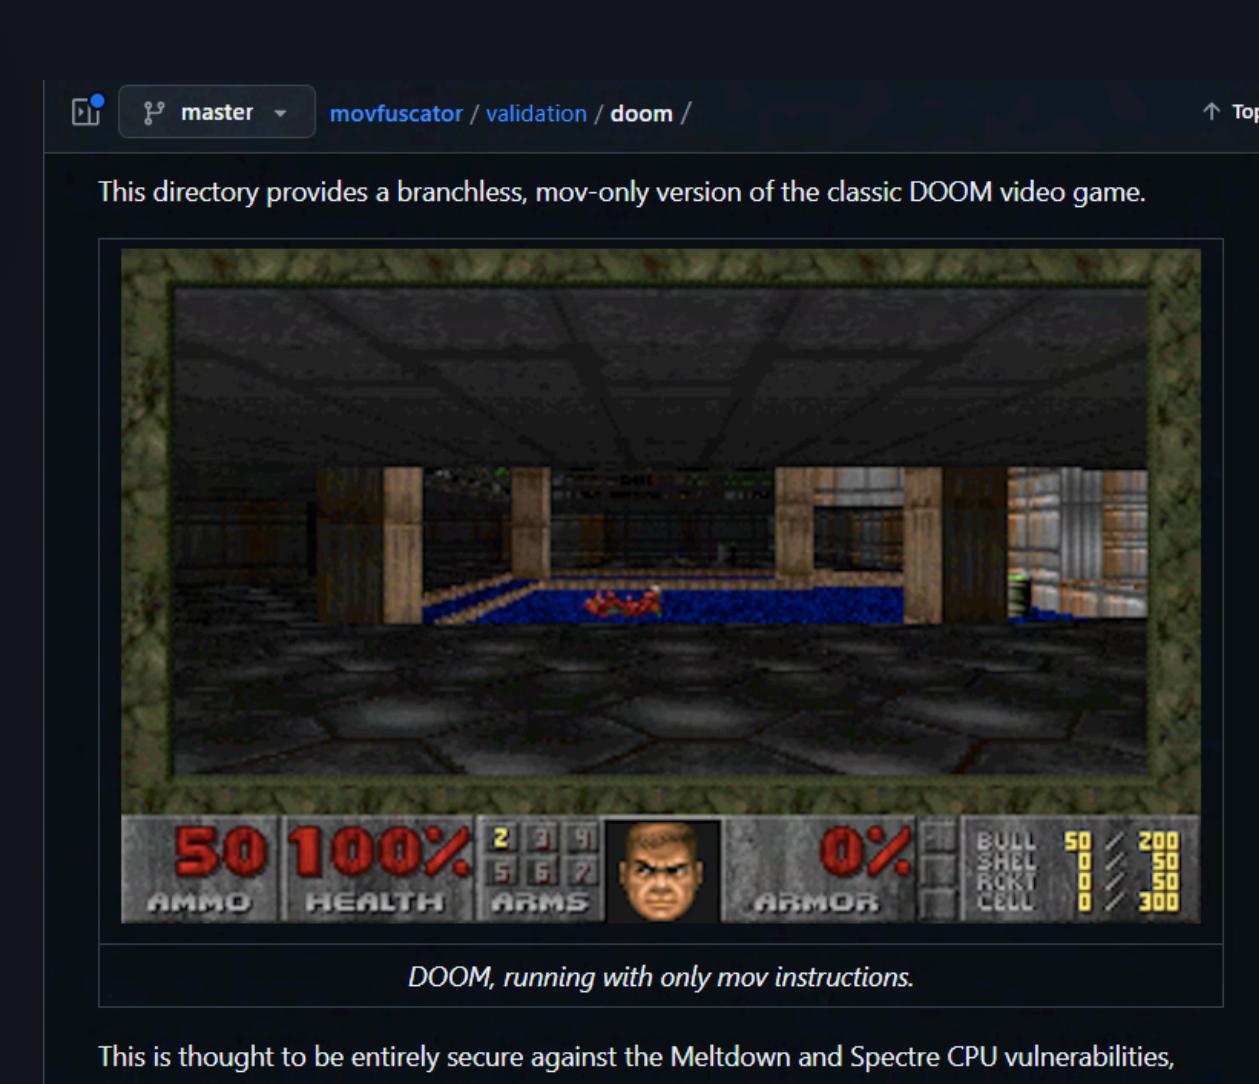
ASM

movfuscator

```
movcc ./find_primes.c -o find_primes_mov >_
objdump -Mintel -d ./find_primes_mov
```

ASM

```
000 mov    eax,DWORD PTR [ecx*4+0x83f8160]
001 mov    edx,DWORD PTR ds:0x85f8194
002 mov    DWORD PTR [eax],edx
003 mov    DWORD PTR ds:0x83f8164,0x804f048
004 mov    eax,DWORD PTR [ecx*4+0x83f8160]
005 mov    edx,DWORD PTR ds:0x85f8198
006 mov    DWORD PTR [eax],edx
007 mov    DWORD PTR ds:0x83f8164,0x804f050
008 mov    eax,DWORD PTR [ecx*4+0x83f8160]
009 mov    edx,DWORD PTR ds:0x85f819c
010 mov    DWORD PTR [eax],edx
011 mov    edx,DWORD PTR ds:0x85f81a0
012 mov    DWORD PTR [eax+0x4],edx
013 mov    DWORD PTR ds:0x83f8164,0x804f058
014 mov    eax,DWORD PTR [ecx*4+0x83f8160]
015 mov    edx,DWORD PTR ds:0x85f81a4
```



fetch → decode → execute

Toolbox Digression #1

- `objdump`: disassemble binary into assembly
- `perf` (`perf list` for available counters)

```
perf stat -e duration_time,cpu-migrations:u,instructions,cycles:u ls >-
```

```
>> Performance counter stats for 'ls':
```

1490843 ns	duration_time:u
0	cpu-migrations:u
1031098	instructions:u
971675	cycles:u

1.06 insn per cycle

```
0.001490843 seconds time elapsed
```

```
0.000000000 seconds user
```

```
0.001473000 seconds sys
```

movfuscator – perf

```
$ perf stat -e instructions,duration_time ./find_primes_mov >-
```

```
Performance counter stats for './find_primes_mov':
```

```
 8,896,386,311      instructions:u  
 6,819,887,233 ns  duration_time:u
```

F.A.Q.

- Q: Why did you make this? A: I thought it would be funny.
- `find_primes.c`: `movcc` version uses 800x instructions; 2,500x slower than `gcc` version.
- The mov-only DOOM renders approximately one frame every 7 hours.

Recap – Binary

- ELF file spec
- Instruction encoding
- Calling convention
- Tools: `objdump`, `perf`
- Instructions
- Registers
 - `mov eax, DWORD PTR [rsp+8]` (relative to stack pointer)
 - `mov eax, DWORD PTR [rsi]` (any memory address)



Memory Layout

```

1 #include <algorithm>
2 #include <string>
3 #include <vector>
4
5 int global = 10;
6 void foo() {}
7
8 int main(int argc, char** argv) {
9     int stack1;
10    int stack2;
11    int* heap1 = new int();
12    int* heap2 = new int();
13
14    std::vector<std::pair<std::string, void*>> items;
15    items.emplace_back("globals", reinterpret_cast<void*>(&global));
16    items.emplace_back("functions", reinterpret_cast<void*>(&foo));
17
18    items.emplace_back("stack1", reinterpret_cast<void*>(&stack1));
19    items.emplace_back("stack2", reinterpret_cast<void*>(&stack2));
20    items.emplace_back("heap1", reinterpret_cast<void*>(heap1));
21    items.emplace_back("heap2", reinterpret_cast<void*>(heap2));
22
23    items.emplace_back("command line args", reinterpret_cast<void*>(argv));
24    items.emplace_back("environment variables", getenv("HOME"));
25
26    std::ranges::sort(items, std::greater<>{},
27                      []<const auto& p> { return p.second; });
28    for (const auto& [name, ptr] : items) {
29        printf("%-25s %p\n", name.c_str(), ptr);
30    }
31 }
```



»	environment variables	0x7ffc56785fe8
	command line args	0x7ffc567844b8
	stack1	0x7ffc5678434c
	stack2	0x7ffc56784348
	heap2	0x18f5e2d0
	heap1	0x18f5e2b0
	globals	0x406058
	functions	0x4011a6

Questions

- What if two programs use the same memory address?
- Where do heap and stack meet?

Memory – addresses

```
1 #include <stdio.h>
2 #include <atomic>
3 #include <chrono>
4 #include <thread>
5
6 std::atomic<int> num = 40;
7
8 int main() {
9     ++num;
10    std::this_thread::sleep_for(
11        std::chrono::seconds(1));
12    printf("addr: %p, val: %d\n",
13          &num, num);
14 }
```



```
./globals_loc & ./globals_loc >_
```

```
>> addr: 0x602034, val: 41
addr: 0x602034, val: 41
```

- Memory addresses aren't real
- OS maintains table for address translation
- TLB (translation lookaside buffer) caches recent translations

Memory – Stack

```
1 #include <iostream>
2
3 static constexpr auto ONE_MiB = (1U << 20U);
4 struct OneMiB { char data[ONE_MiB]; };
5
6 void measureStackSize(OneMiB* first_addr) {
7     OneMiB local_var;
8     const auto diff = first_addr - &local_var;
9     std::cout << "stack size: " << diff
10    << " MiB" << std::endl;
11    measureStackSize(first_addr); // recurse
12 }
13
14 int main() {
15     OneMiB first;
16     measureStackSize(&first);
17     return 0;
18 }
```



```
» stack size: 1 MiB
stack size: 2 MiB
stack size: 3 MiB
stack size: 4 MiB
stack size: 5 MiB
stack size: 6 MiB
Segmentation fault (core dumped)
```

Memory – Heap

```

1 #include <cassert>
2 #include <iostream>
3
4 static constexpr auto ONE_GiB = (1U << 30U);
5 struct OneGiB { char data[ONE_GiB]; };
6
7 int main() {
8     OneGiB *curr_addr;
9
10    for (int cnt = 0;; ++cnt) {
11        try {
12            curr_addr = new OneGiB;
13        } catch (const std::bad_alloc &e) {
14            std::cout << "malloc failed after: "
15                         << cnt << " GiB" << std::endl;
16            break;
17        }
18    }
19 }
```



» malloc failed after: 131070 GiB

htop:

CPU	VIRT	CPU%	MEM%	TIME+	Command
3	127T	0.0	0.4	0:00.98	./heap_limit

Question

- What happens when we use the memory?

Memory – Mapping time + perf stat

```

1 static constexpr auto ONE_GiB = (1U << 30U);
2 struct OneGiB { char data[ONE_GiB]; };
3
4 int main() {
5     OneGiB* curr_addr;
6
7     for (int cnt = 0; cnt < 30; ++cnt) {
8         curr_addr = new OneGiB;
9     }
10 }
```



```

1 static constexpr auto ONE_GiB = (1U << 30U);
2 struct OneGiB { char data[ONE_GiB]; };
3
4 int main() {
5     OneGiB* curr_addr;
6
7     for (int cnt = 0; cnt < 30; ++cnt) {
8         curr_addr = new OneGiB{}; // value init
9     }
10 }
```



```

g++ -O3 heap_no_write.cpp -o heap_no_write
sudo perf stat -e duration_time ./heap_no_write
```

» 0.001664263 seconds time elapsed

```

sudo perf stat -e page-faults ./heap_no_write
```

» 126 page-faults

```

g++ -O3 heap_write.cpp -o heap_write
sudo perf stat -e duration_time ./heap_write
```

» 19.227826577 seconds time elapsed

```

sudo perf stat -e page-faults ./heap_write
```

» 7.864.447 page-faults

Toolbox Digression #2

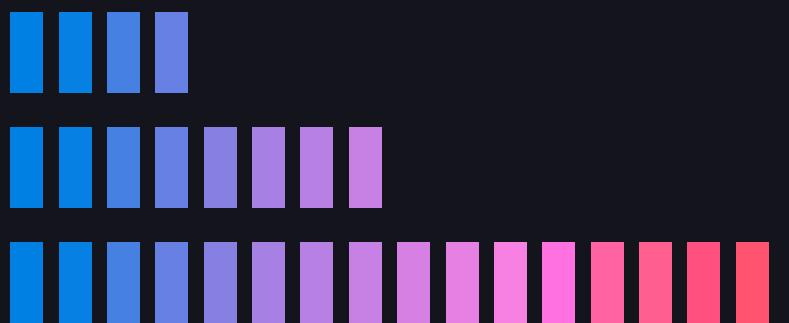
- **tsc** – time stamp counter

```
1 struct ProfileScope {  
2     ProfileScope() {  
3         tsc_start_ = __rdtsc();  
4     }  
5     ~ProfileScope() {  
6         const uint64_t elapsed = __rdtsc() - tsc_start_;  
7         std::cout << elapssed << std::endl;  
8     }  
9     uint64_t tsc_start_;  
10};
```

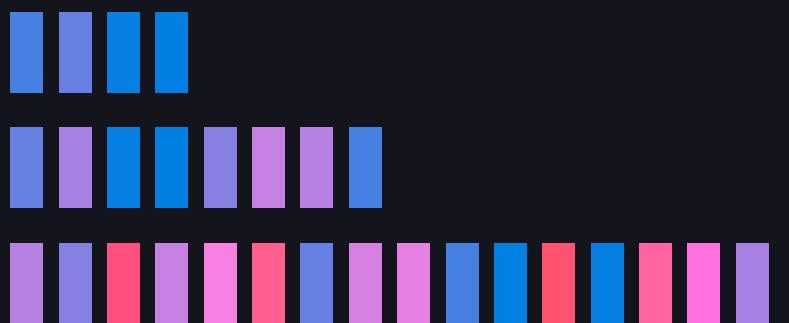
[Memory – Demo]

Memory – measurement patterns

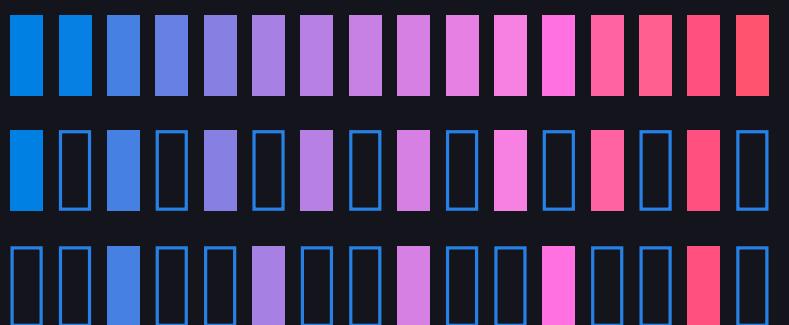
Sequential, $f(\text{range})$:



Random, $f(\text{range})$:



Stride, $f(\text{stride})$:



Toolbox Digression #3

- Google Benchmark
 - Benchmarking library
 - Statistical analysis

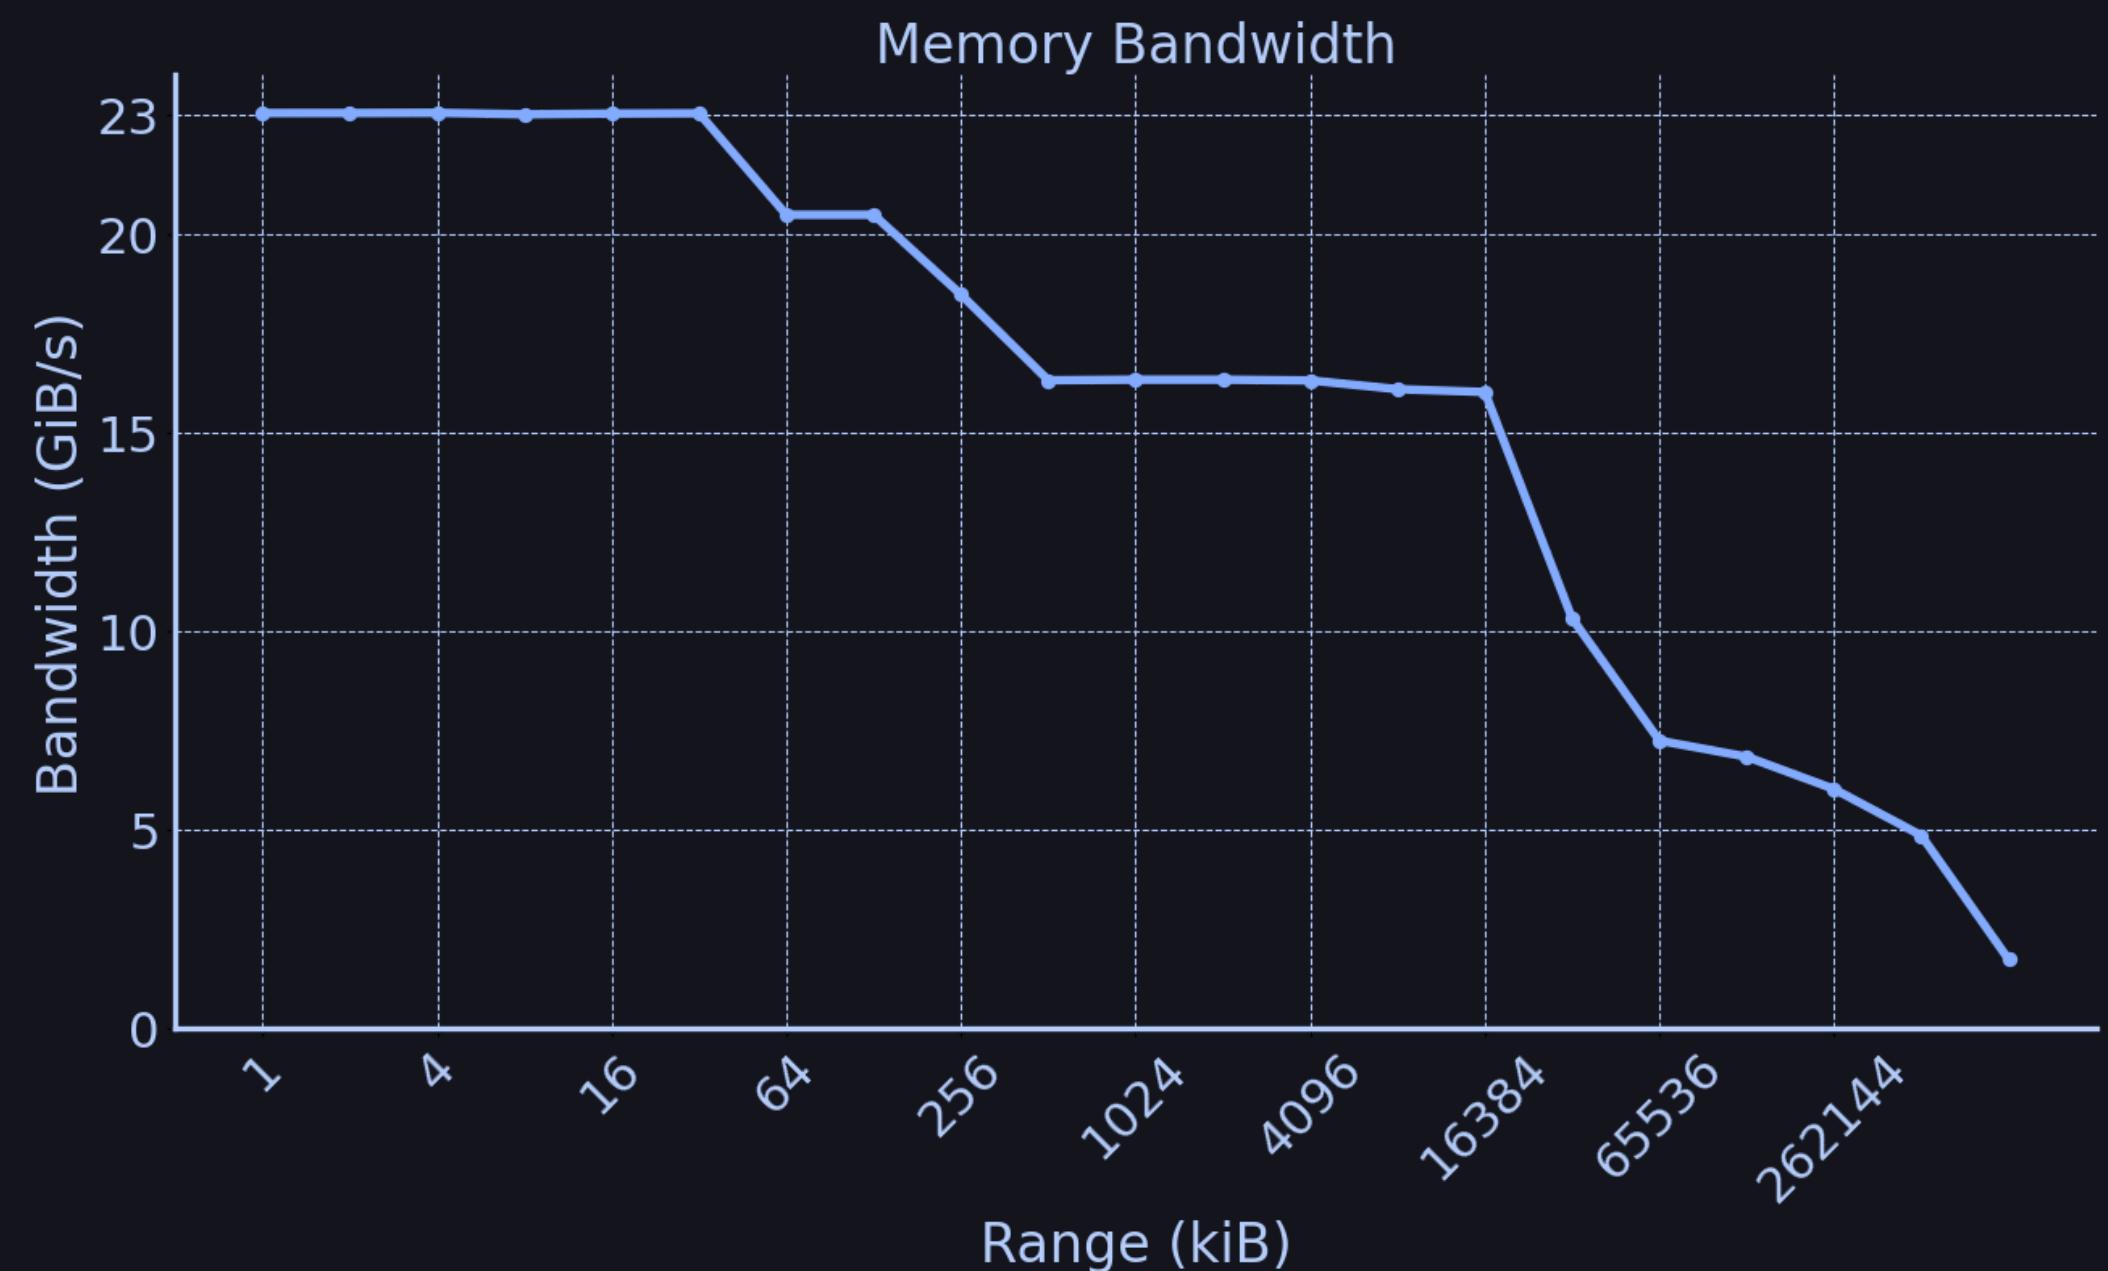
```
#include <benchmark/benchmark.h>

static void BM_SomeFunction(benchmark::State& state) {
    for (auto _ : state) {
        // test code
    }
}
BENCHMARK(BM_SomeFunction); // register test
BENCHMARK_MAIN();
```

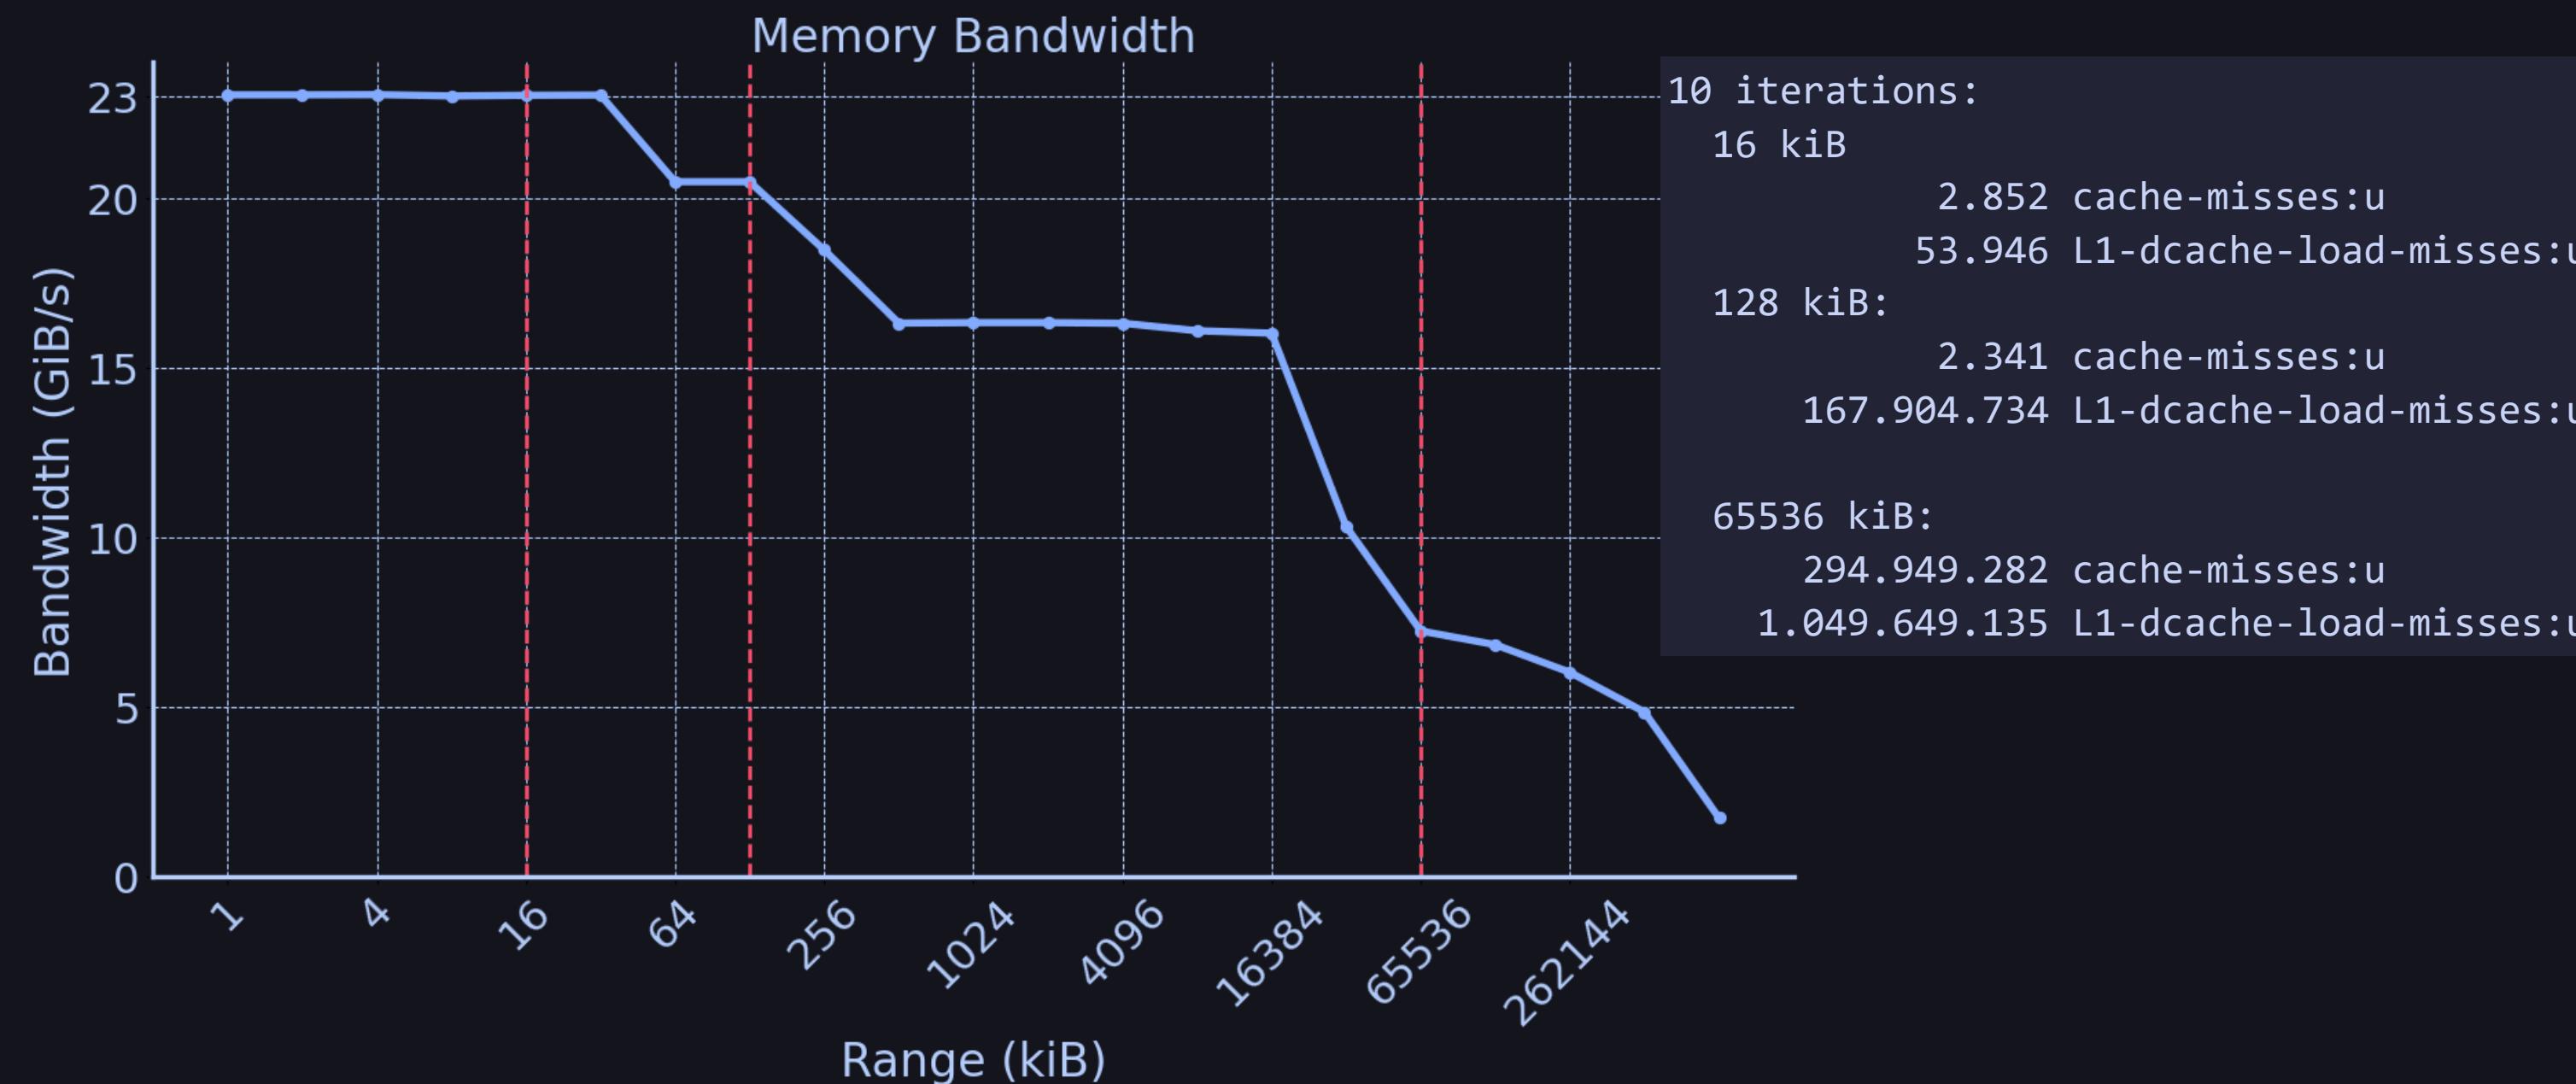


Memory – sequential write [DEMO]

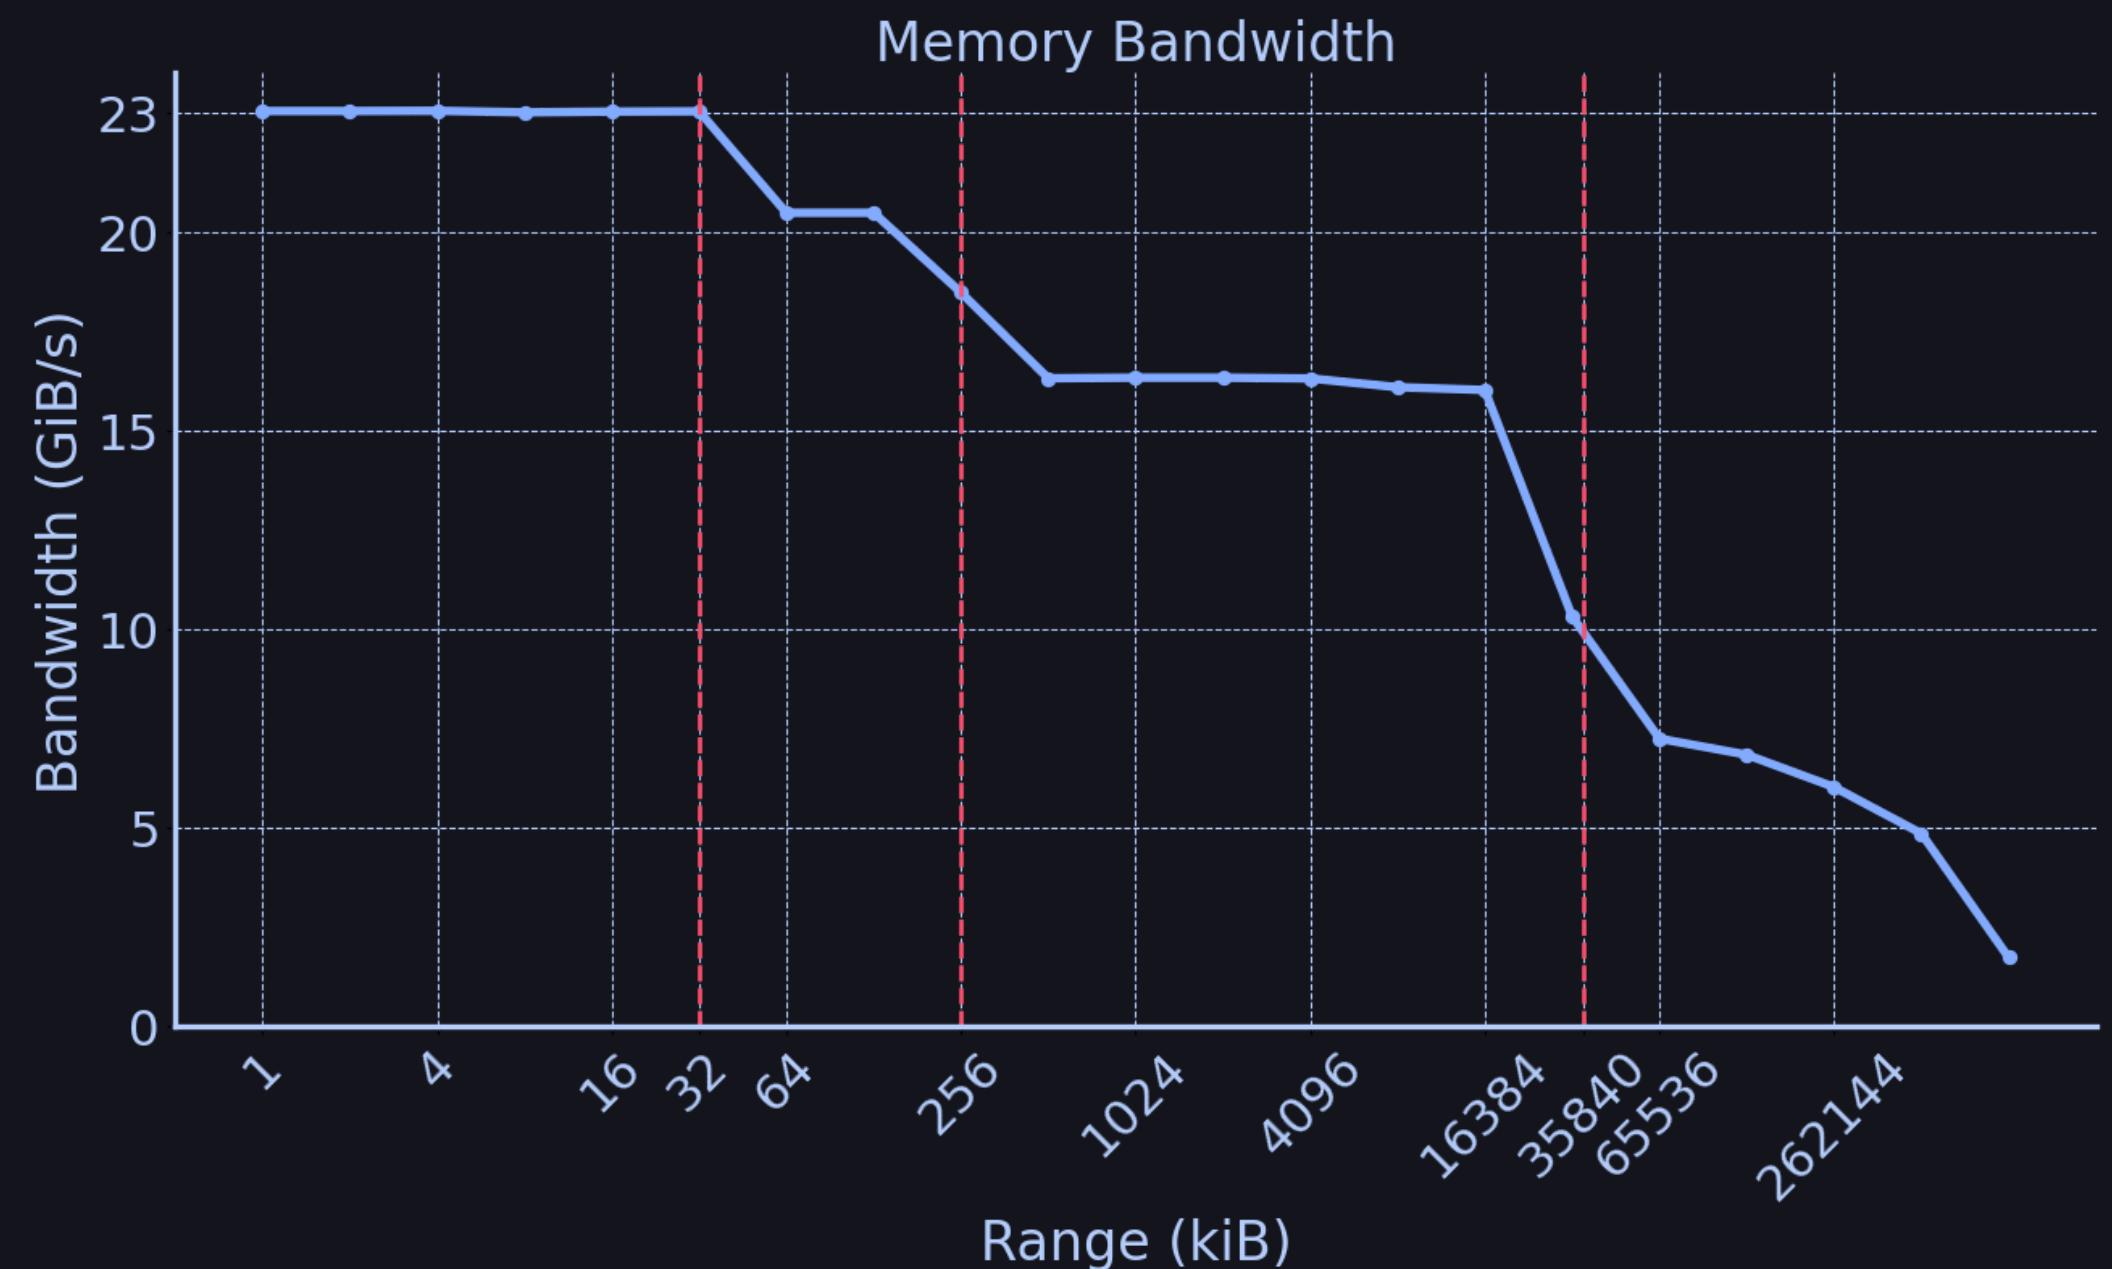
Memory – sequential write



Memory – sequential write



Memory – sequential write

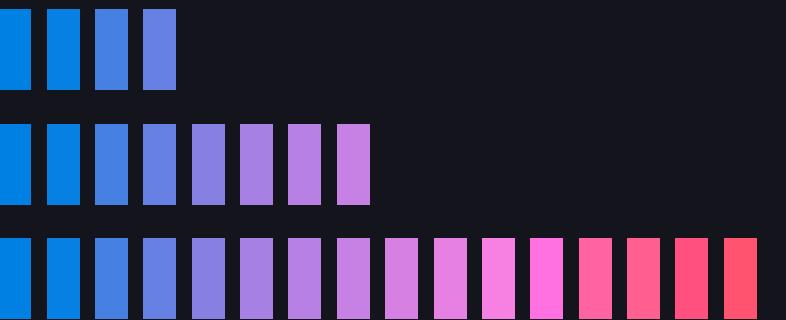


Memory – random write

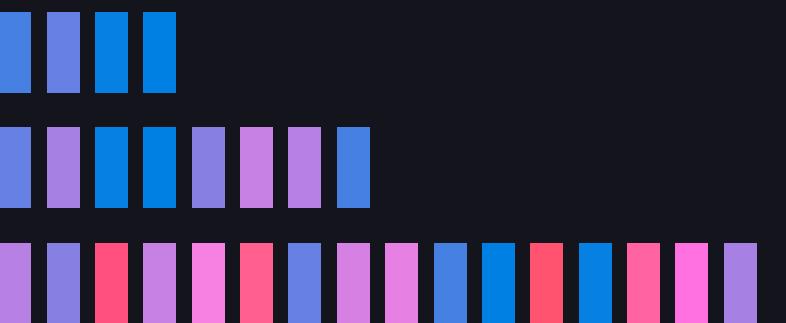


Cache – strided copy

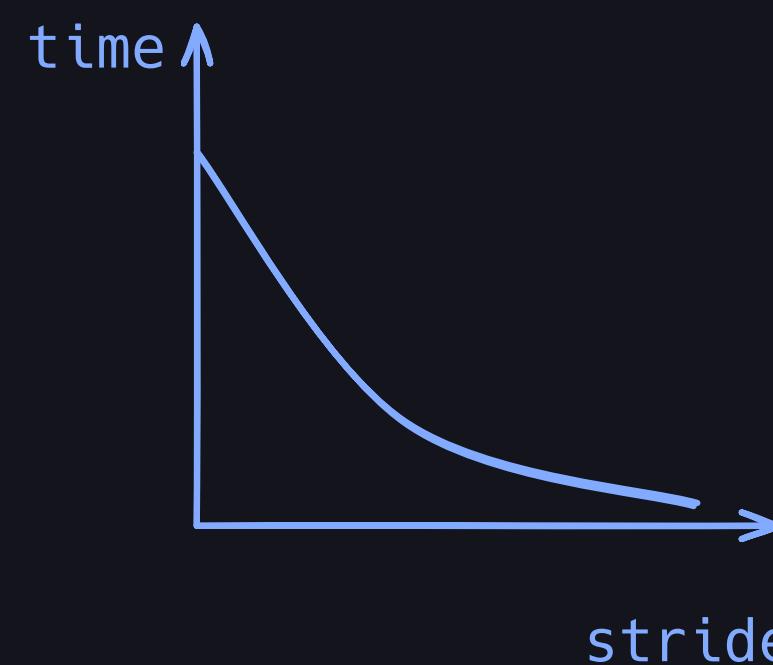
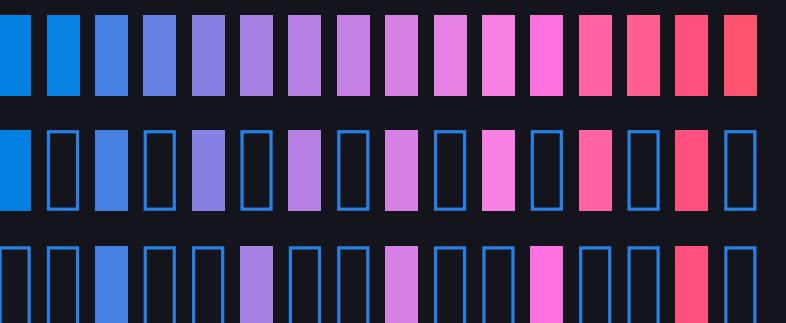
Sequential, $f(\text{range})$:



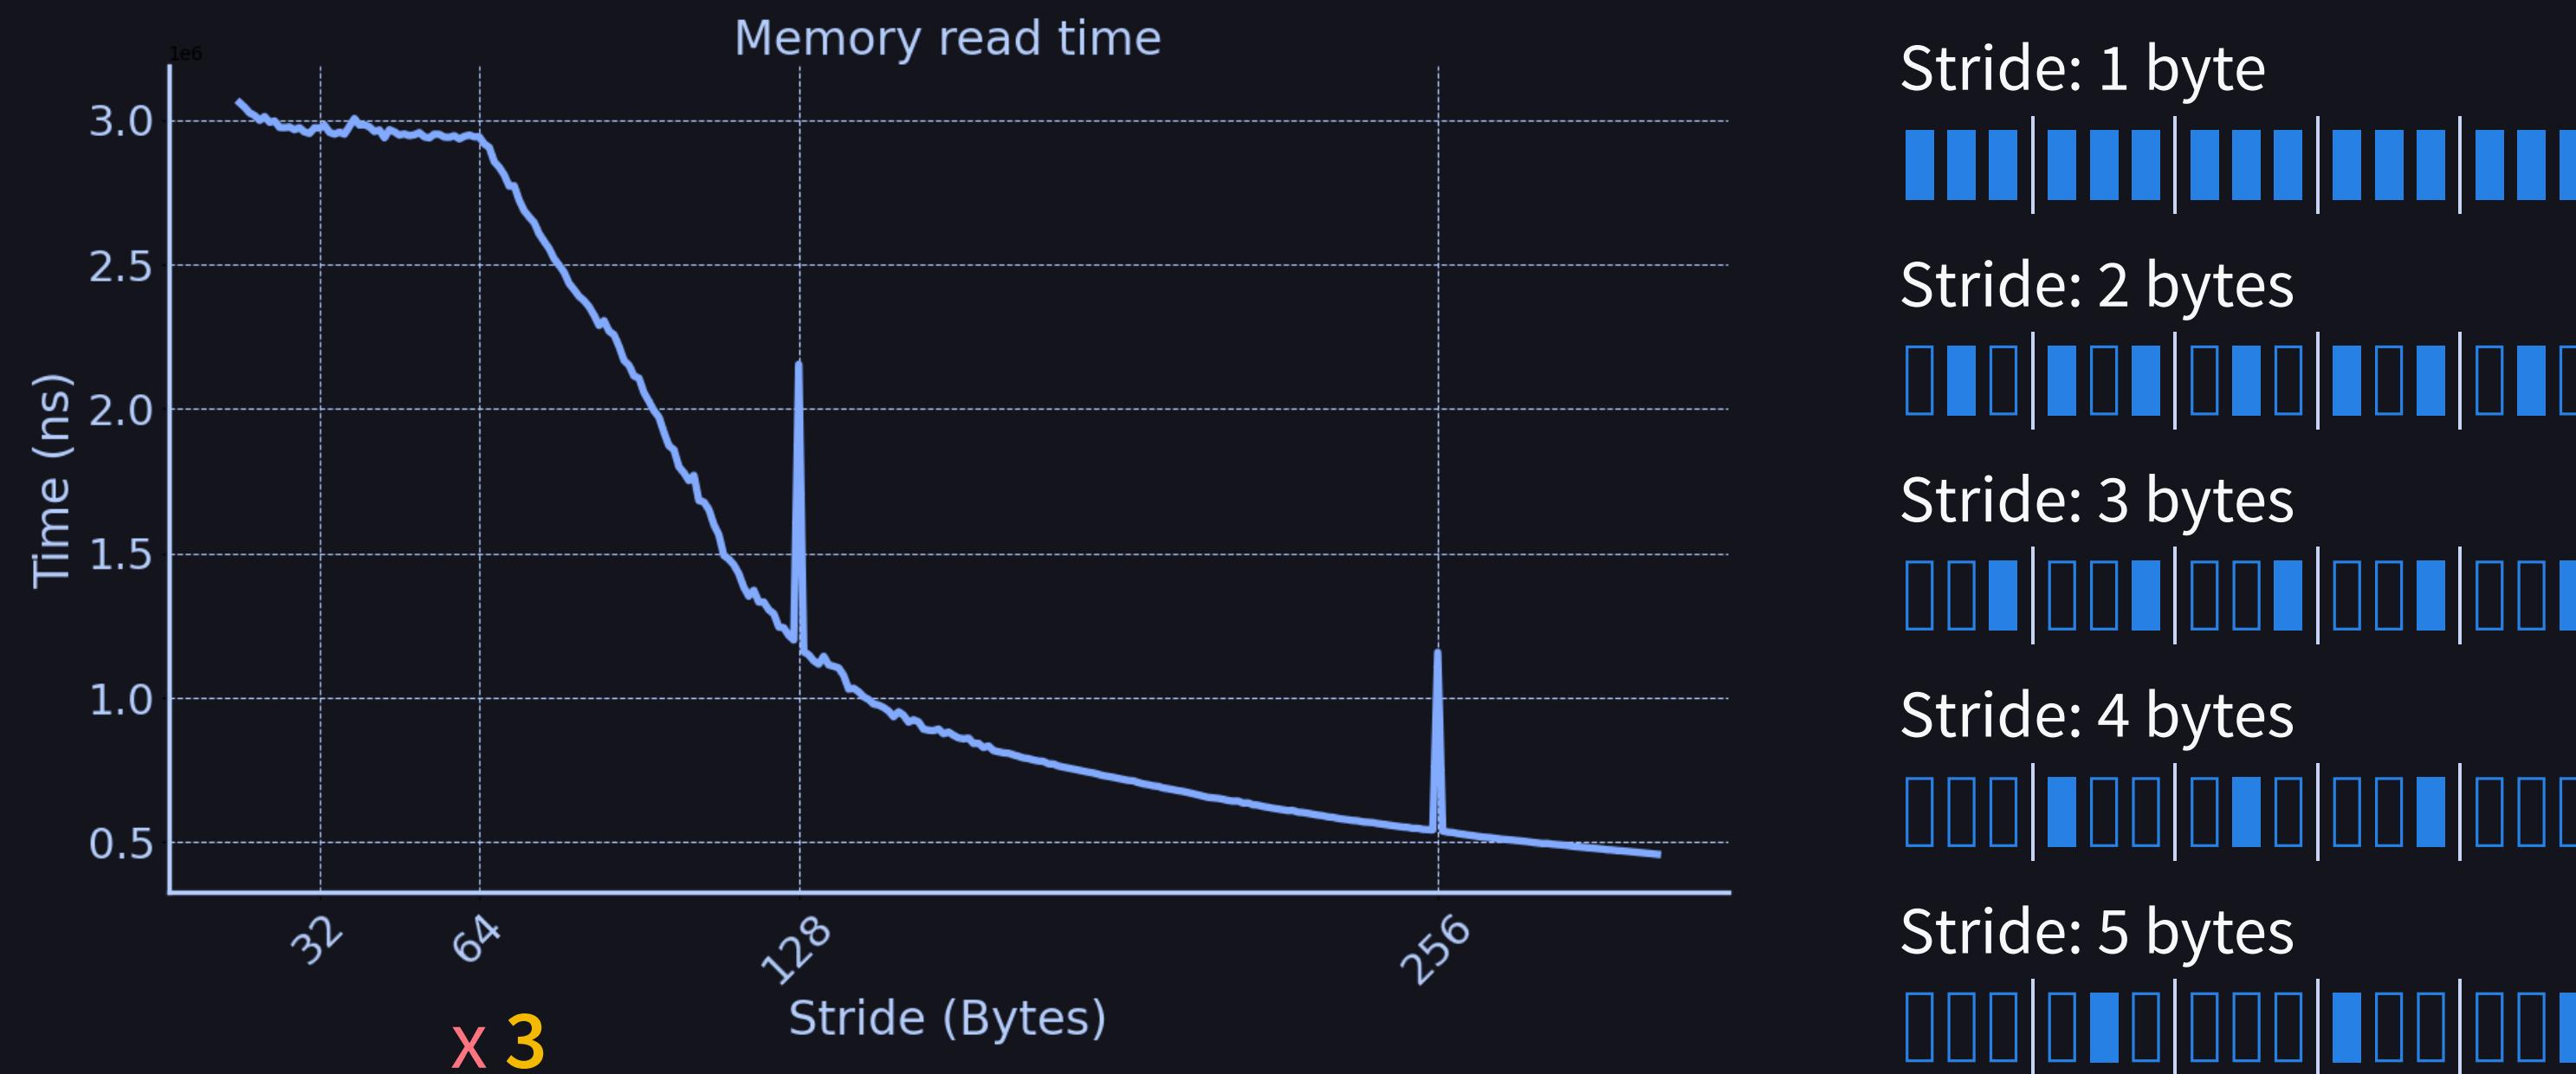
Random, $f(\text{range})$:



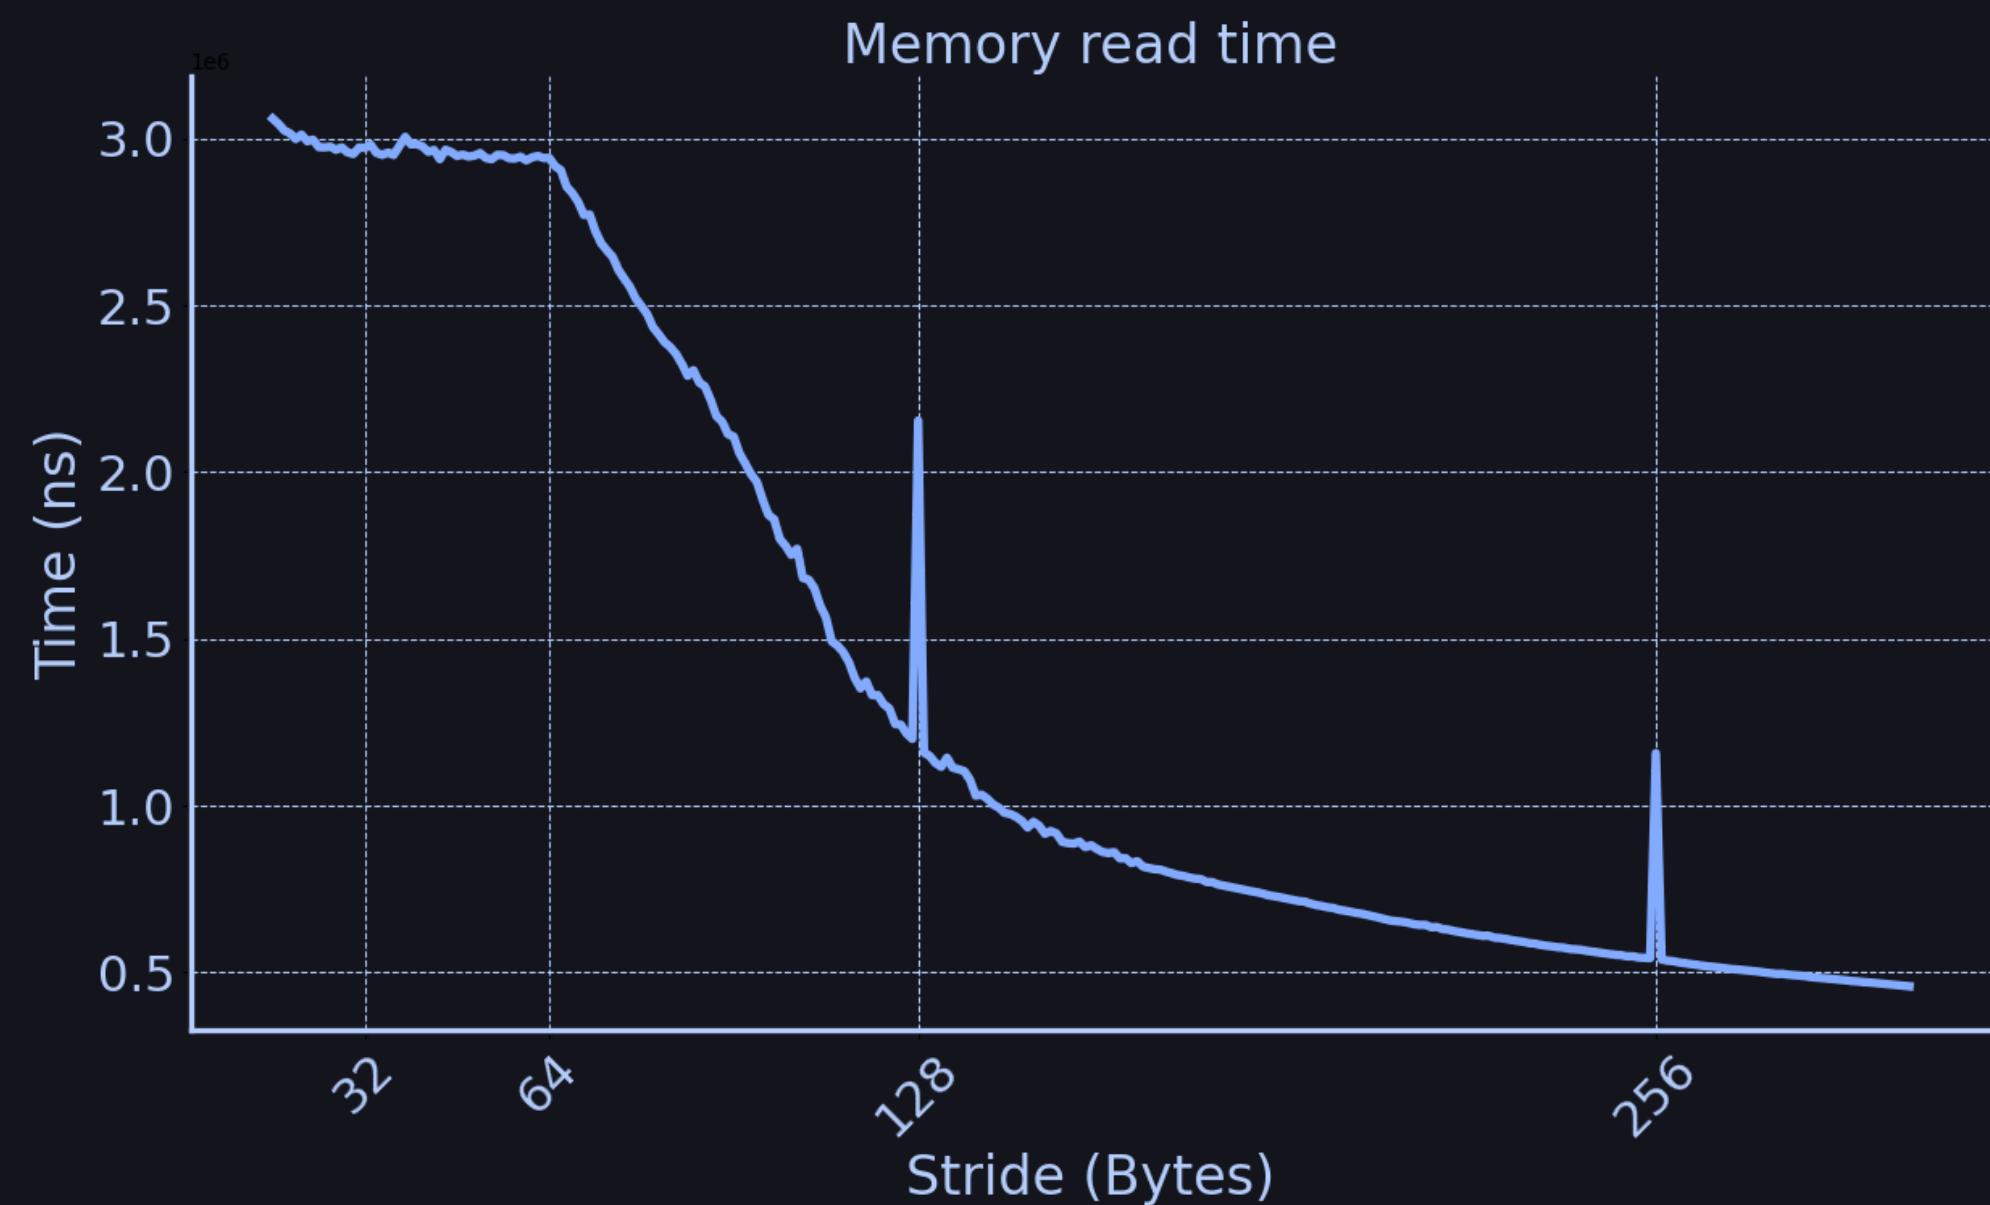
Stride, $f(\text{stride})$:



Memory – strided copy



Memory – strided copy



$$\text{cache_sz} = \text{cache_line_sz} \cdot \#\text{ways} \cdot \#\text{sets}$$

#ways = associativity

set index bits

Address: `xxxxxxxxx xxxxxxxxx 00111010 01101010 00001110 01000110 01110100 01xxxxxx`

16 bits: ignored

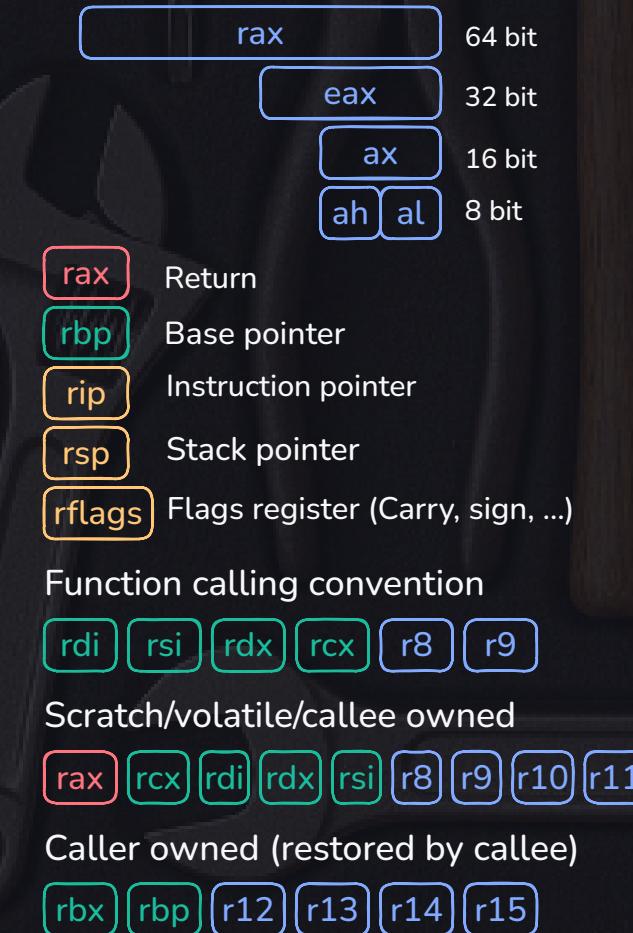
6 bits: cache line

Recap – Memory

- Memory layout & heap/stack
- Virtual memory (paging, TLB)
- Data Caching
 - Cache levels (L1, L2, L3)
 - Cache lines
 - Prefetching
 - Cache associativity
- Tools: `htop`, `tsc`, `google benchmark`

Toolbox Digression #4

- **nasm** (assembler)
 - Call assembly from C++
 - Exact control over executed code



```
enable_language(ASM_NASM)
set(CMAKE_ASM_NASM_FLAGS "-f elf64")
add_executable(addExample)
target_sources(addExample PRIVATE
               add_example.cpp add_example.asm)
```



```
global add
section .text
```



```
add:
    mov rax, rdi ; copy first argument into rax
    add rax, rsi ; add second argument
    ret
```



```
#include <cstdint>
#include <iostream>

extern "C" int64_t add(int64_t a, int64_t b);

int main() {
    int64_t result = add(10, 20);
    std::cout << "Result: " << result << std::endl;
    return 0;
}
```

CPU – frontend / backend

add	add loop	nop	nop loop	9 byte nop	9 byte nop loop
<pre>ASM_add: %rep 100000 add rax, 9 %endrep ret</pre>	<pre>ASM_add_loop: .loop: add rax, 9 dec rdi jnz .loop ret</pre>	<pre>ASM_nop: %rep 100000 nop %endrep ret</pre>	<pre>ASM_nop_loop: .loop: nop dec rdi jnz .loop ret</pre>	<pre>ASM_nop_long: %rep 100000 nop_9_byte %endrep ret</pre>	<pre>ASM_nop_long_loop: .loop: nop_9_byte dec rdi jnz .loop ret</pre>

	add	add_loop	nop	nop_loop	nop_9B	nop_9B_loop
time (μs)	32	32	8	32	51	32
instructions (k)	101	301	101	301	101	301
uops_issued.any (k)	101	201	101	201	101	201
uops_executed.core (k)	101	201	1	101	1	101
icache.misses	55	8	1	7	10.897	7

fetch → decode → ~~execute~~ execute

[Branch Prediction – Demo]

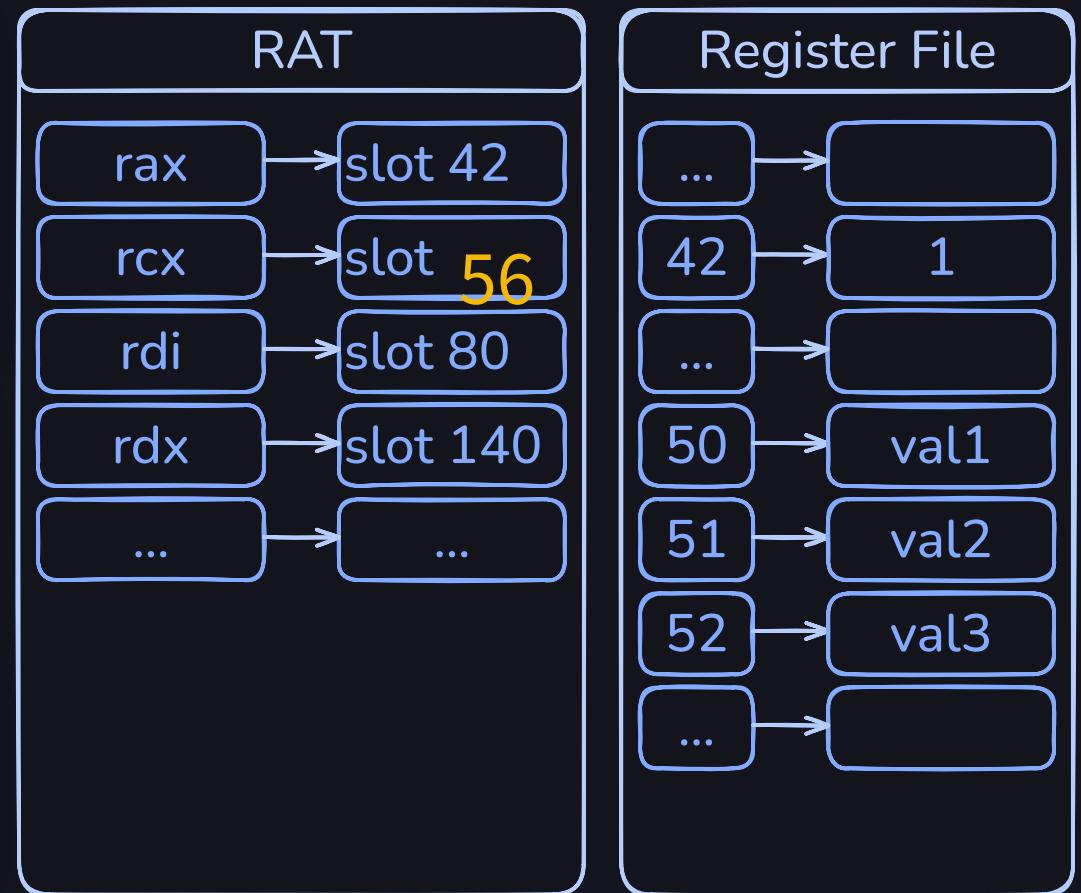
[Registers – Demo]

Registers

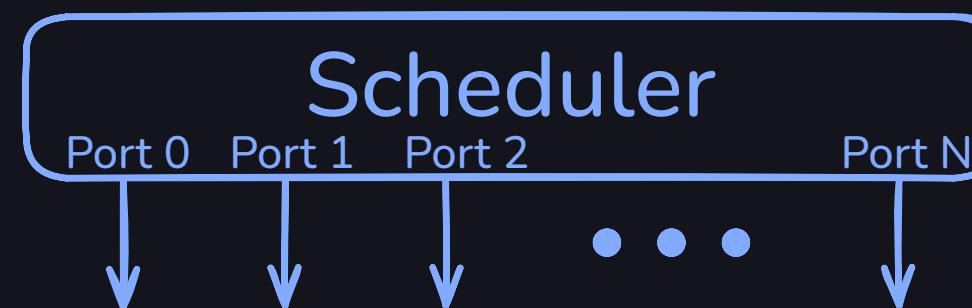
```

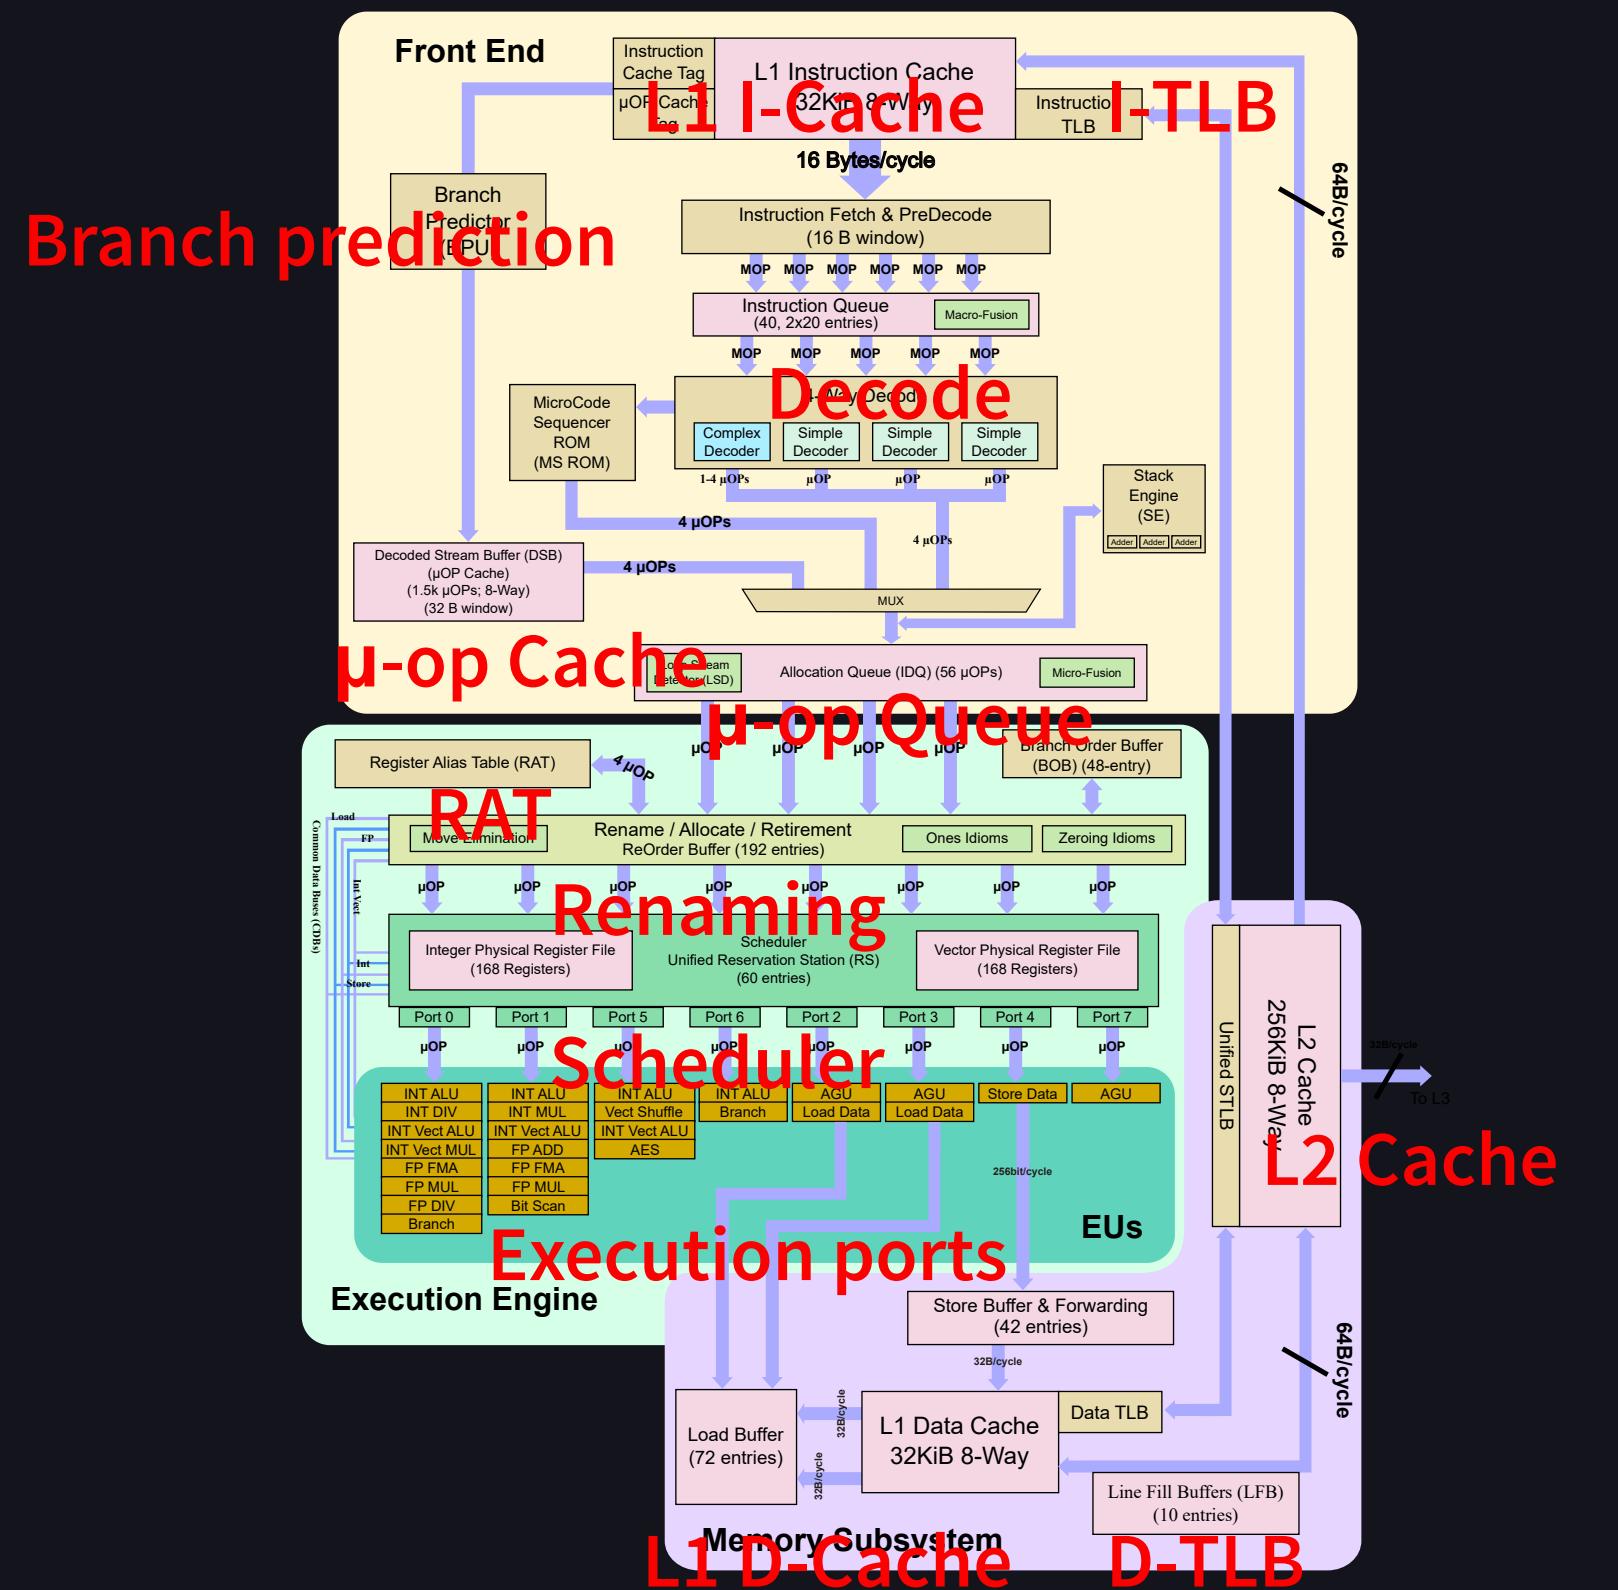
ASM_add:
    %rep 4000
        add rcx, 1      ; rcx += 1
    %endrep
    ret

ASM_mov_add:
    mov rax, 1
    %rep 4000
        mov rcx, rax ; rcx = rax
        add rcx, 1      ; rcx += 1
    %endrep
    ret
  
```



instruction	dest	op1	op2
...	-	-	-
mov rcx, rax	s51	s42	-
add rcx, 1	s52	s51	1
mov rcx, rax	s53	s42	-
add rcx, 1	s54	s53	1
mov rcx, rax	s55	s42	-
add rcx, 1	s56	s55	1
...	-	-	-





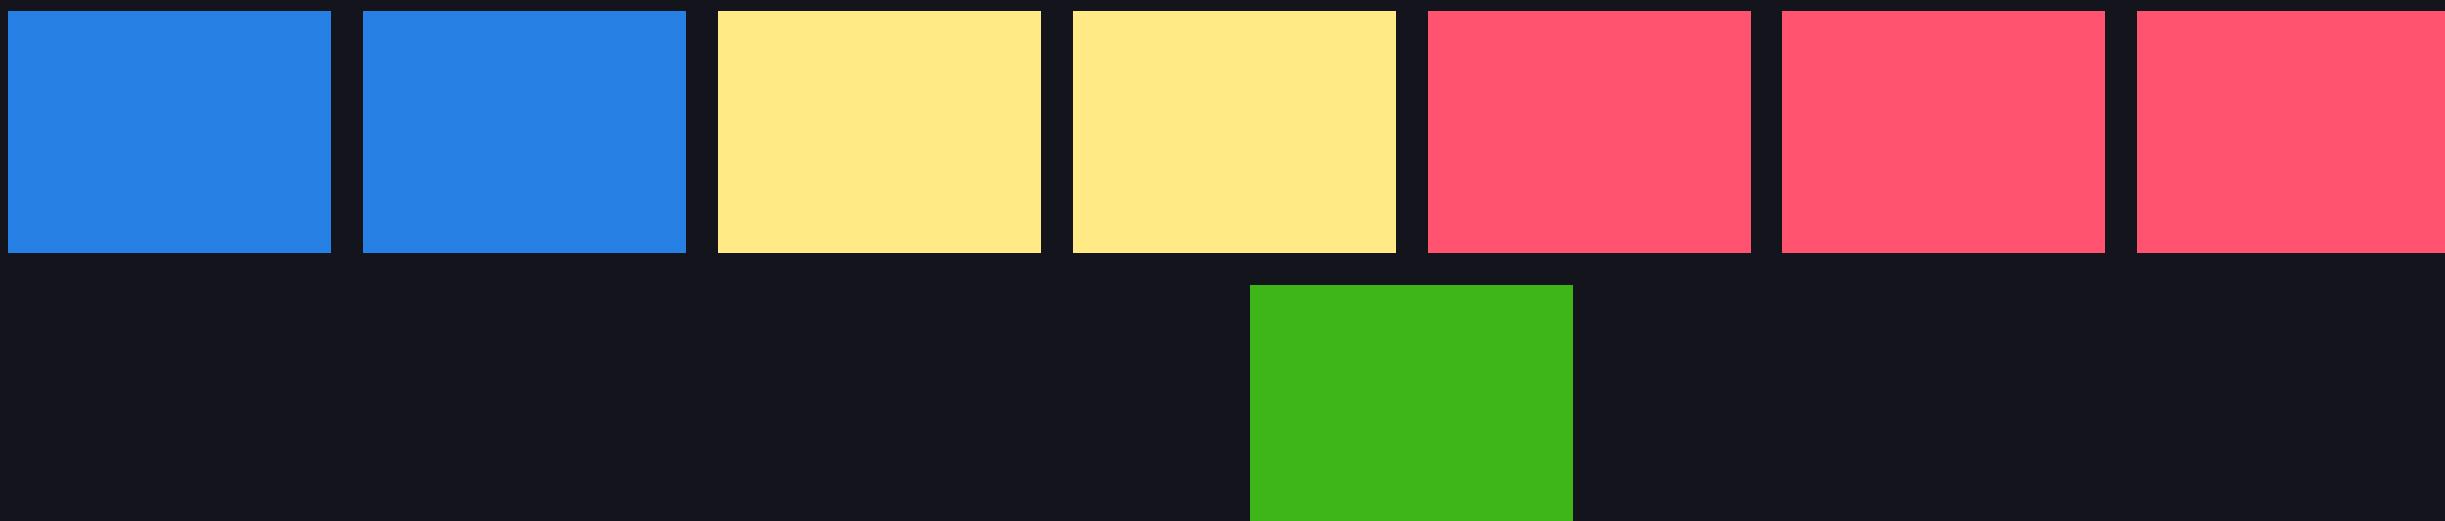
Sorted vector

- maintain sorted `vector<T>`, support insert/remove/update element

```
sorted_values_.push_back(new_val);  
std::sort(sorted_values_.begin(), sorted_values_.end());
```



- Test: 1'000'000 elements, 1'000 inserts: 46s



Sorted vector

- maintain sorted `vector<T>`, support insert/remove/update element

```
sorted_values_.push_back(new_val);  
std::sort(sorted_values_.begin(), sorted_values_.end());
```



- Test: 1'000'000 elements, 1'000 inserts: 46s



```
auto it = std::lower_bound(sorted_values_.begin(), sorted_values_.end(), new_val);  
sorted_values_.insert(it, new_val);
```



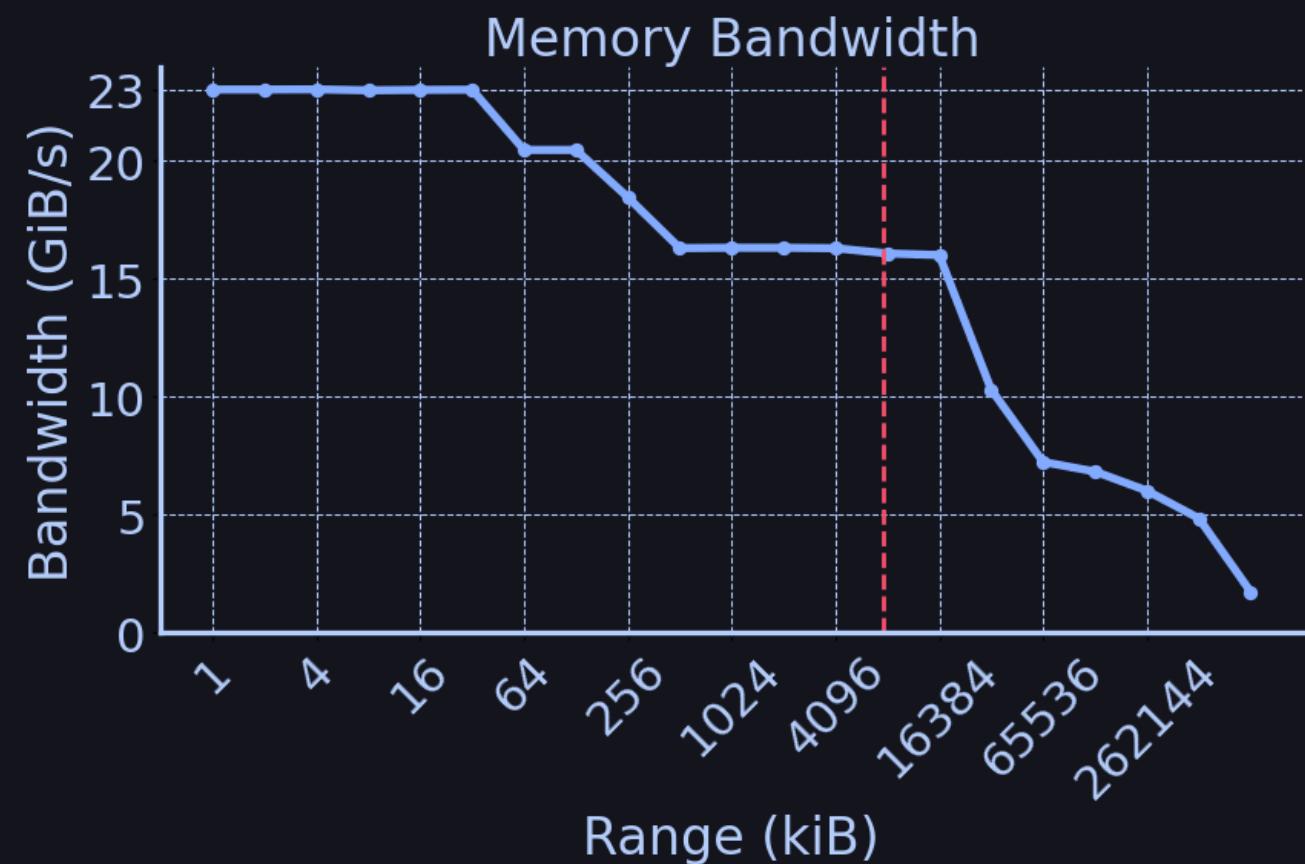
Sorted vector

- maintain sorted `vector<T>`, support insert/remove/update element

```
sorted_values_.push_back(new_val);
std::sort(sorted_values_.begin(), sorted_values_.end());
```



- Test: 1'000'000 elements, 1'000 inserts: **46s**



```
nr_moves = 1000000 / 2 * 1000
bw_gbyte_s = 16
bw_int64_s = (bw_gbyte_s << 30) / 8
expected_ms = int(nr_moves / bw_int64_s * 1e3)

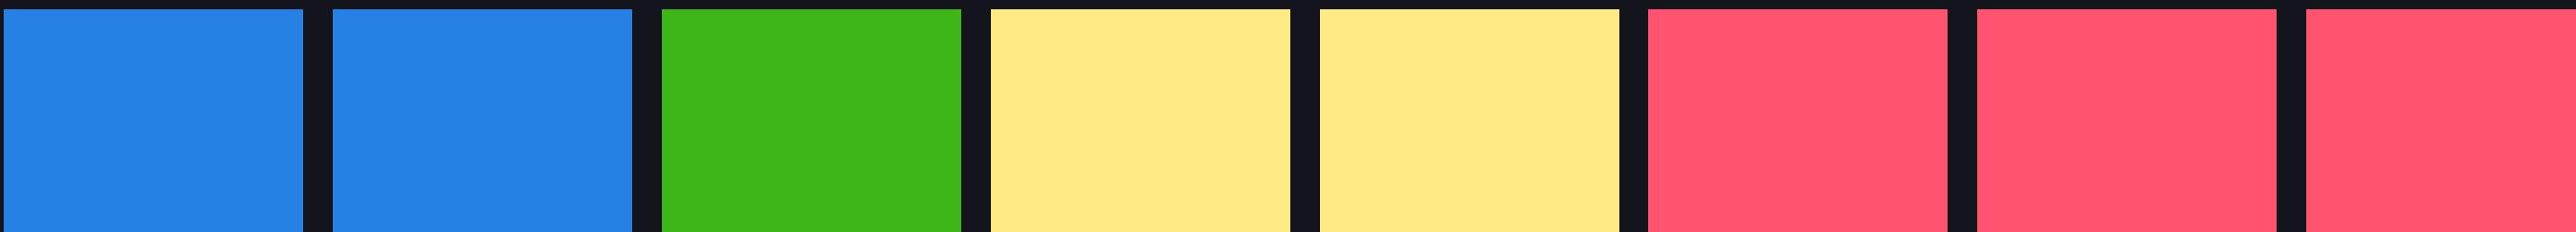
print(f"expected duration: {expected_ms}ms")
```



» expected duration: 232ms

Sorted vector

- maintain sorted `vector<T>`, support insert/remove/update element
- test with: `1'000'000` elements, `1'000` events
- `Insert (std::lower_bound + insert)`: **204ms** (calculated: **232ms**)
- `Remove (std::lower_bound + erase)`: **208ms**
- `Update (remove + insert)`: **427ms**



Sorted vector

- maintain sorted `vector<T>`, support insert/remove/update element
- test with: `1'000'000` elements, `1'000` events
- `Insert (std::lower_bound + insert)`: **204ms** (calculated: **232ms**)
- `Remove (std::lower_bound + erase)`: **208ms**
- `Update (remove + insert)`: **427ms**



```

auto curr_loc = std::lower_bound(sorted_values_.begin(), sorted_values_.end(), existing);
auto new_loc = std::lower_bound(sorted_values_.begin(), sorted_values_.end(), new_val);
*curr_loc = new_val; // update value

if ((curr_loc == new_loc) || ((curr_loc + 1) == new_loc)) { continue; } // same location

auto start = std::min(curr_loc, new_loc);
auto end = std::max(curr_loc + 1, new_loc);
auto mid = (new_loc < curr_loc) ? curr_loc : curr_loc + 1;
std::rotate(start, mid, end);

```



Sorted vector

- maintain sorted `vector<T>`, support insert/remove/update element
- test with: `1'000'000` elements, `1'000` events
- `Insert (std::lower_bound + insert)`: **204ms** (calculated: **232ms**)
- `Remove (std::lower_bound + erase)`: **208ms**
- `Update (remove + insert)`: **427ms**
- `Update (2x std::lower_bound + std::rotate)`: **141ms**

Memory Usage Python

- sum all quantities

```
things = get_things(5)
quantities = [thing.qty_ for thing in things]
print(sum(quantities))
```



→ mem usage: 9MB

```
things = get_things(10000000)
quantities = [thing.qty_ for thing in things]
print(sum(quantities))
```

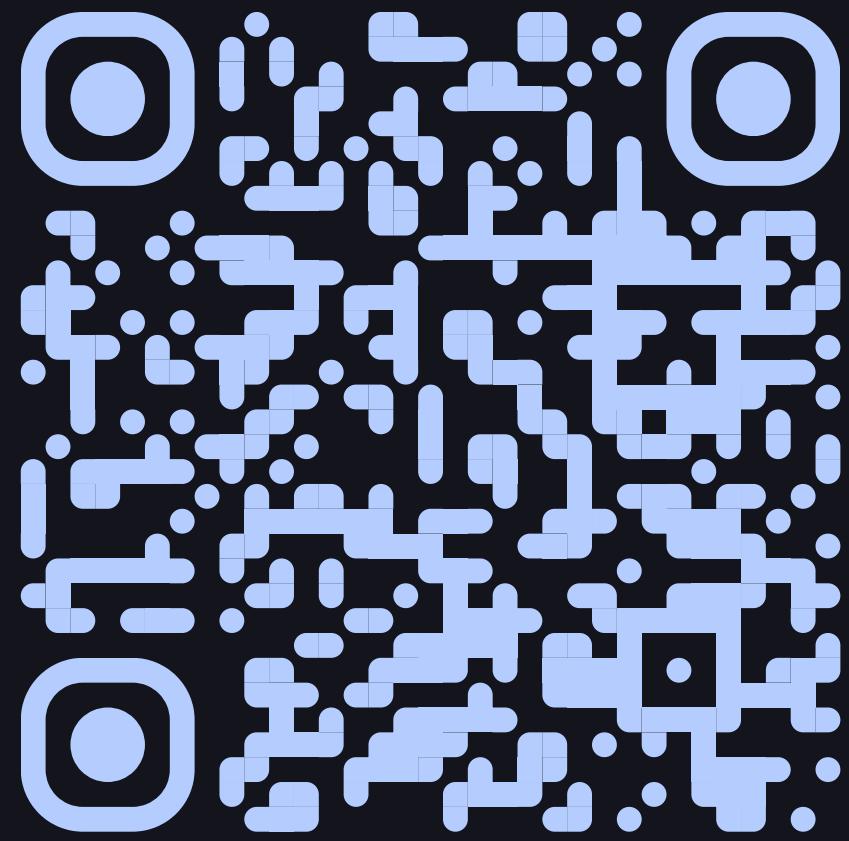
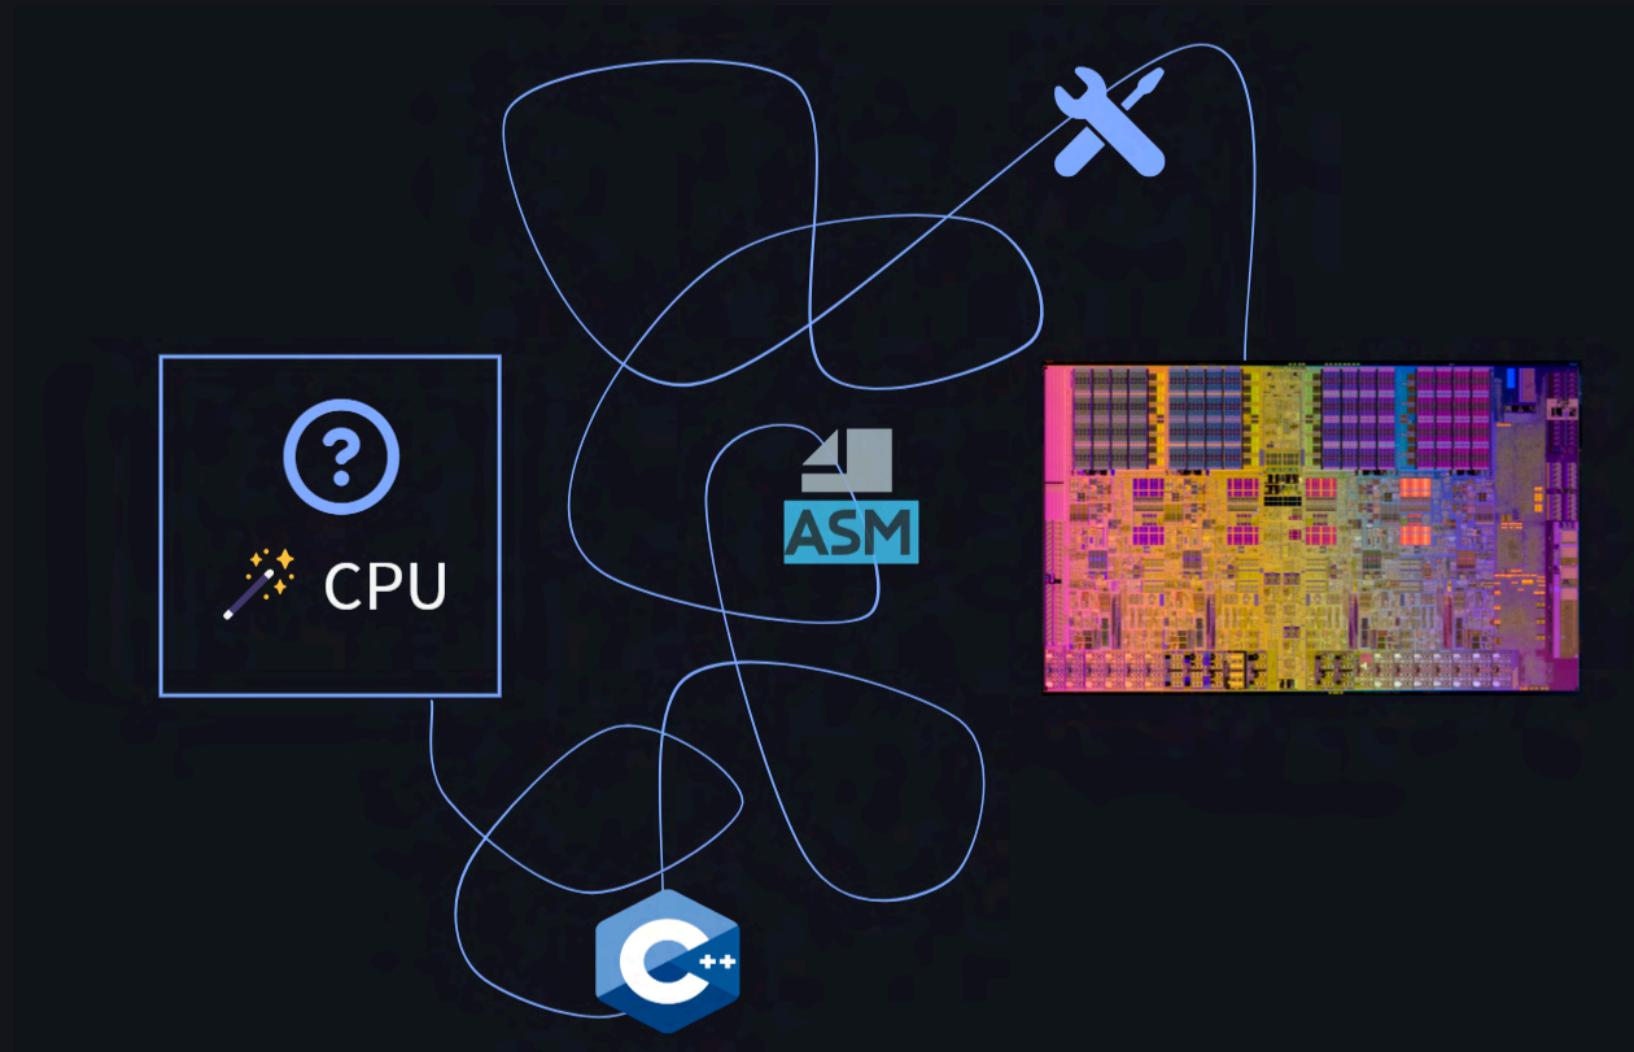


→ mem usage: 404MB

```
things = get_things(10000000)
quantities = (thing.qty_ for thing in things)
print(sum(quantities))
```



→ mem usage: 9MB



github.com/linusboehm/perf_stuff/

✖ : objdump , perf , htop , rdtsc , google benchmark , nasm , llvm-mca ,
compiler-explorer , beman exemplar