

Five Issues with `std::expected` and How to Fix Them

Vitaly Fanaskov

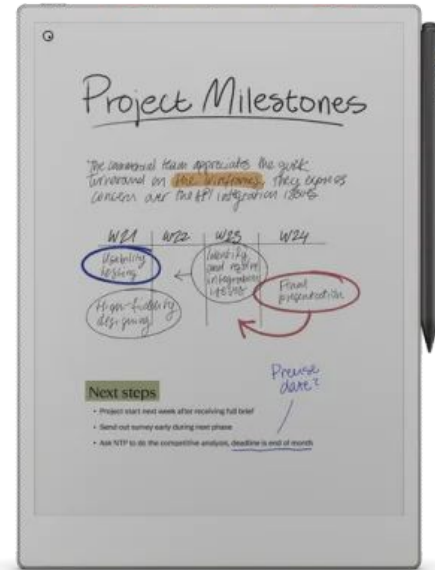
Five Issues with `std::expected` and How to Fix Them

About me

Vitaly Fanaskov

Principal software engineer at reMarkable

Work on C++ frameworks and libraries



Agenda

- Briefly about `std::expected`
- Some implementation details
- Monadic and regular interfaces
- Customization points
- `bind_back` and `bind_front`
- More monadic operations
- Conclusions

In this talk

- Almost no theory
- Simplified C++ code snippets
- Many change/improve suggestions

Examples and experimental implementation



<https://github.com/vt4a2h/fl>



[expected.hpp](#)

Briefly about `std::expected`

Definition

- Represents either of two values: *expected* (T) or *unexpected* (E)
- Union under the hood
- Has monadic operations

Basic example of using std::expected

```
#include <expected>
```

```
std::expected<int, Error> expectedBox;
```

```
expectedBox = 42;
```

```
fmt::println("The value is: {}", expectedBox.value());
```

std::expected as a return value

When an operation can fail and we need to know why:

```
std::expected<Widget, WidgetError> loadWidget()
{
    // If error
    return std::unexpected(WidgetError{ /* ... */ });

    // Actual result
    return Widget{ /* ... */ };
}
```

Process std::expected

```
void loadWidget()
{
    if (const auto widgetBox = getNewWidget(); widgetBox.has_value()) {
        const auto widget = widgetBox.value();
        // Do something with the widget ...
    } else {
        const auto error = widgetBox.error();
        // Handle the error ...
    }
}
```

That was not an entirely bad example...

```
void loadWidget()  
{  
    const auto widgetBox = getNewWidget();  
  
    if (widgetBox.has_value()) {  
        // Do something with the widget ...  
    } else {  
        const auto error = widgetBox.error();  
        log("Cannot get a new widget {}: {}.", widgetBox.value(), error);  
    }  
}
```

Monadic operations to the rescue!

```
getWidget()  
    .and_then([](const auto &widget) → std::expected<Widget, WidgetError> {  
        // Do something with the widget ...  
        return widget;  
    })  
    .transform([](const auto &widget) → ID { return widget.id(); })  
    .or_else([](const auto& error) → std::expected<ID, WidgetError> {  
        log(error);  
        // Possibly recover and/or cleanup ...  
        // Return value or error ...  
    })
```

Available monadic operations

- `std::expected<V, E>::transform(F &&f)`
 - $f: V \rightarrow T$
- `std::expected<V, E>::transform_error(F &&f)`
 - $f: E \rightarrow T$
- `std::expected<V, E>::and_then(F &&f)`
 - $f: V \rightarrow \text{std::expected}<T, E>$
- `std::expected<V, E>::or_else(F &&f)`
 - $f: E \rightarrow \text{std::expected}<V, T>$

For more information



[Monadic Operations
in Modern C++: A
Practical Approach -
Vitaly Fanaskov -
CppCon 2024](#)

The slide has a dark blue background. In the top right corner, there is a green graphic element consisting of a large plus sign and the number '24'. The title 'Monadic Operations in Modern C++:' is written in a large, bold, orange font, with 'A Practical Approach' in a smaller white font below it. The speaker's name 'VITALY FANASKOV' is in orange on the right side. At the bottom left is the Cppcon logo (a white circle with a stylized plus sign) and the text 'Cppcon The C++ Conference'. At the bottom right is the year '2024' in large white numbers, followed by a stylized mountain logo and the dates 'September 15 - 20'.

You will not be doing purely functional programming

- There will be side effects
- There will be object-oriented-style parts of the existing code base
- There will be integration with 3rd-party libraries
- ...

A quick look under the hood

Types

- `template <class E> class unexpected`
- `template <class E> class bad_expected_access`
- `template <class T, class E> class expected`

Data structure

```
union {  
    T val;  
    E unex; // NOTE: not unexpected<E>!  
};  
bool has_val;
```

Issue #1:

Value type and error type can be the same

This is perfectly fine code

```
using Expected = std::expected<std::string, std::string>;
```

```
Expected doSomething()
```

```
{
```

```
    return "Correct result"s;
```

```
    // Or
```

```
    return std::unexpected{"Something goes wrong"s};
```

```
};
```

How often do you
need this?

Consequences of the existing implementation

- Extra type `std::unexpected<E>`
- API users must explicitly specify this type
- Complex implementations with many different restrictions
- Harder to use `std::variant<T, E>` under the hood

What is wrong with std::unexpected<E>?

- Additional abstraction
- Not quite symmetrical

Examples from other languages:

```
enum Result<T, E> {  
    Ok(T),  
    Err(E),  
}
```

...

```
Result<i32, i32>  
Ok(42)  
Err(42)
```

```
sealed abstract class Either[+A, +B]  
extends Product with Serializable
```

...

```
Either[Int, Int]  
Left(42)  
Right(42)
```


Let's add some simple restrictions

```
!std::is_same_v<Value, Error>
```

```
!std::is_convertible_v<Value, Error>
```

```
!std::is_convertible_v<Error, Value>
```

Declare “new” expected

```
template <... Value, ... Error>
    requires (...)
struct expected : public std::variant<Value, Error>
{
    using base_t = std::variant<...>;

    using base_t::base_t;
};
```

- Use `monostate` when `Value` is void
- `Error` cannot be void

What do we get from this?

- Reuse the existing abstraction
- Simplify implementation (approx. 350 lines of code with some extra functionality, approx. 1k lines of tests)
- Use `std::get` and all variant-related utils
- Use `std::visit` (!)

A short story about `std::visit`

std::visit can be very convenient (regular use case)

```
using Shape = std::variant<Square, Rectangle, Circle>;
```

```
template<class... Ts>  
struct overloaded : Ts... { using Ts::operator()...; };
```

```
double calculateArea(const Shape &shape)  
{  
    return std::visit(overloaded{  
        [](const Square& square) { return std::pow(square.side, 2); },  
        [](const Rectangle& rect) { return rect.width * rect.height; },  
        [](const Circle& circle) { return std::numbers::pi * std::pow(circle.radius, 2); },  
    }, shape);  
}
```

```
std::println("The shape area is: {}", calculateArea(shape));
```

...and can be even better

```
template <IsVariant Variant, class... Matchers>
    requires(sizeof...(Matchers) ≥ 1)
decltype(auto) match(Variant&& variant, Matchers&&... matchers)
{
    return std::visit(
        overloaded{std::forward<Matchers>(matchers) ...},
        std::forward<Variant>(variant));
}
```

match can be more convenient

```
double calculateArea(const Shape &shape)
{
    return match(shape,
        [](const Square& square) { return std::pow(square.side, 2); },
        [](const Rectangle& rect) { return rect.width * rect.height; },
        [](const Circle& circle) { return std::numbers::pi * std::pow(circle.radius, 2); },
    );
}
```

Try it with expected

```
expected<Widget, WidgetError> widgetBox = getNewWidget();

match(widgetBox,
    [] (const Widget& widget) { /* Handle value ... */ },
    [] (const WidgetError& error) { /* Handle error ... */ }
);
```


We almost have
pattern matching!

Section conclusions

Pros

- Simpler implementation
- Drop redundant abstractions
- Use standard utilities for `std::variant`
- Almost pattern matching now

Cons

- Type restrictions
- Potentially negative impact on performance

Issue #2: Monadic and regular interfaces

API duplication

```
widgetBox.and_then(  
    [](const Widget& widget) → Expected {  
        // Do something with a widget;  
        return widget;  
    }).or_else(  
        [](const WidgetError &error) → Expected {  
            // Handle error  
            return std::unexpected{error};  
        });
```

```
if (widgetBox) {  
    const auto &widget = widgetBox.value();  
    // Do something with a widget;  
} else {  
    const auto &error = widgetBox.error();  
    // Handle error  
}
```

It seems that non-monadic interface is simpler and less verbose, right?

Benefits of monadic operations

- Safety
- Easy to propagate value or error
- Can be used in most of cases

Example: monadic operations in unit tests

```
std::ignore = calculateArea(shape)
    .and_then([expectedArea](double actualArea) → Expected {
        REQUIRE(actualArea == expectedArea);
        return actualArea;
    })
    .or_else([](const auto &error) → Expected {
        FAIL(fmt::to_string(error));
        return std::unexpected{error};
    });
```

Places where you still need the regular interface

- `main()`
- Libraries boundary
- “Integration” boundary

Example: “integration” boundary (QML interface)

```
class IWindowLayout
{
public:
    virtual core::Expected<> add(const window::Uri& windowUri) = 0;
};
```

```
class WindowLayout : public QObject
{
    Q_OBJECT
    // ...
public:
    Q_INVOKABLE bool WindowLayout::add(const window::Uri& windowUri)
private:
    std::shared_ptr<IWindowLayout> m_windowLayout;
};
```


Example: “Integration” boundary (expose into QML)

```
bool WindowLayout::add(const window::Uri& windowUri)
{
    const auto result = m_windowLayout->add(windowUri).or_else(&printError);
    return result.has_value();
    // Or return a value of another type, such as QString
}
```

Section conclusions

- Monadic API is safer than regular API
- Monadic API can be used in most common use cases
- You may still need regular API
- Implementation via `std::variant` can gracefully hide dangerous API

Issue #3:
No customization points

We have a problem!

```
constexpr const T& value() const &
```

Throws: `bad_expected_access(as_const(error()))` if `has_value()` is `false`

```
constexpr const T* operator→() const // Ditto for operator*
```

Preconditions: `has_value()` is `true`.

If `has_value()` is `false`, the behavior is undefined.

```
constexpr const E& error() const & noexcept
```

Preconditions: `has_value()` is `false`.

If `has_value()` is `true`, the behavior is undefined.

Issues with value/error APIs

- Not symmetrical
- Throws
- Undefined behavior
- Cannot customize

Different error handling techniques



[Error handling in
C++ - Vitaly
Fanaskov - NDC
Techtown 2022](#)



Provide a customization point

```
template<class Self>
    requires (...)
[[nodiscard]] constexpr auto&& value(this Self&& self) noexcept (...)
{
    if (self.has_error()) {
        if constexpr (CustomValueHandlerFound<Self>) {
            handle_bad_value(std::get<error_t>(std::forward<Self>(self)));
        } else {
            default_handle_bad_value<Self>();
        }
    }
    return std::get<value_t>(std::forward<Self>(self));
}
```

Default handler example

```
template <class>
constexpr void default_handle_bad_value()
{
    if consteval {
        throw "Expected object contains error";
    } else {
        std::terminate();
    }
}
```


Custom handler example

```
namespace gui
{

    void handle_bad_value(const WidgetError&) noexcept
    {
        // Handle, log, print stack trace...
        std::terminate();
    }

} // namespace gui
```

It's not that easy

- Can handlers throw?
- Terminate by default?
- Singleton-like get/set or template specialization or else
- Function signatures are unclear

Contracts in C++26?

- For `operator*()` and `operator->()`, not for `.value()`
- With customization points
- Implementation details of contracts are still questionable
- Still with potential UB

Section conclusions

- Customization points for `value()/error()` can be super useful
 - No exceptions
 - No undefined behavior
 - Set different behavior for libraries or test code
- Many questions to API design and use cases
- Contracts could help
- We don't need to think of these issues if we have `std::variant` + monadic API

Issue #4:
No built-in `bind_back`/`bind_front`

bind_back/front

```
auto add = [](int a, int b) { return a + b; };  
auto addOne = std::bind_back(add, 1);
```

```
std::println("{} ", addOne(2)); // prints 3
```

```
auto inc = [](int &a, int v) { a += v; };
```

```
int a{};
```

```
auto incA = std::bind_front(inc, std::ref(a));
```

```
incA(2);
```

```
std::println("{} ", a);
```

Use cases with monadic operations

- Lambda functions can potentially reduce readability
 - Too long
 - Too many
 - Too nested
- There are already functions to use
 - More than single parameter
 - Used in several places
 - Not easy to change
- A function is a class method

Some issues when using with monadic operations

- May not be available
 - `std::bind_front` – C++20
 - `std::bind_back` – C++23
- Preserves arguments
- Arguments must be move constructible
- Implementation can use extra intermediate structures (e.g. tuple)

Some implementation details

```
template <class _Fn, class... _Args>
requires
    is_constructible_v<decay_t<_Fn>, _Fn> &&
    is_move_constructible_v<decay_t<_Fn>> &&
    (is_constructible_v<decay_t<_Args>, _Args> && ...) &&
    (is_move_constructible_v<decay_t<_Args>> && ...)
constexpr auto bind_front(_Fn&& __f, _Args&&... __args)
{
    [...bound_args = std::forward<_Args>(__args), f = std::forward<_Fn>(__f)]
    <class Self, class... T>(this Self&&, T&&... call_args)
    noexcept(/* ... */)
    → /* ... */
    {
        return std::invoke(f, std::forward_like<Self>(bound_args) ..., std::forward<T>(call_args) ...);
    }
}
```

Make it built-in for monadic operations

```
template<class Self, class F, class ...Args>
    requires (/* ... */)
[[nodiscard]] constexpr auto and_then(this Self&& self, F &&f, Args &&...back_args)
    noexcept (/* ... */)
    → /* ... */
{
    if (self.has_value()) {
        return std::invoke(
            std::forward<F>(f),
            std::get<value_t>(std::forward<Self>(self)),
            std::forward<Args>(back_args) ... );
    } else {
        return std::get<error_t>(std::forward<Self>(self));
    }
}
```

Almost the same for bind_front

```
/* ... */  
[[nodiscard]] constexpr auto and_then(  
    this Self&& self, bind_front_t, F &&f, Args &&...front_args)  
/* ... */  
    return std::invoke(  
        std::forward<F>(f),  
        std::forward<Args>(front_args) ...  
        std::get<value_t>(std::forward<Self>(self)));  
/* ... */
```

Pros of the approach

- Required no extra functions
- Don't need to store arguments
- Requirements to arguments are relaxed
- Don't need to use `std::ref` or `std::cref`
- Monadic operations keep the same signature (except for `bind_front`)

Section conclusions

- It's relatively easy to have built-in `bind_back` and `bind_front`
- Using built-in `bind_back` and `bind_front` has some practical benefits
- May not be fully idiomatically correct

Issue #5:
Not enough monadic operations!

General information



[Monoids, Monads, and
Applicative Functors:
Repeated Software
Patterns - David Sankel
- CppCon 2020](#)



Applicative-like may be the most useful



[Applicative: The
Forgotten Functional
Pattern in C++ - Ben
Deane - CppNow
2023](#)



Let's make it a little
more practical

More canonical pattern quite far from regular use cases

```
template <class V, class E = Error>
using Expected = std::expected<V, E>;

template <class T, class U>
Expected<U> apply(Expected<std::function<U(T)>> f, Expected<T> v)
{
    if (f && v) {
        return std::invoke(*f, *v);
    } else {
        return std::unexpected{!f ? f.error() : v.error()};
    }
}

template <class T>
auto pure(T t) { return Expected<T>{t}; }
```

Let's say we use it as

```
const std::function<std::string(int)> toString =  
    [](int v) { return std::to_string(v); };
```

```
const Expected<int> number = getNumber();
```

```
const Expected<std::string> resultStr = apply(pure(toString), number);
```

```
// Do something with the result
```

However, most likely...

- A function is not wrapped into expected/optional/etc.
- A function parameter types are not expected/optional/etc.
- Some arguments will be wrapped into expected/optional/etc.

More real example

```
const auto toString =  
    [](int v) { return std::to_string(v); };  
  
const auto number = getNumber();  
  
const Expected<std::string> resultStr = apply(  
    toString, // ← just a function  
    number    // ← can be Expected<int> or int  
);
```

For expected type

- We can have an `ap`-like method (super close to `and_then`)
- Function to call is just a function
- Some arguments can be wrapped into `expected` with the same error type

Method interface

```
template <class Self, class F, class ...Args>
    requires (ApInvocable<F, error_t, value_t, Args...>)
[[nodiscard]] constexpr auto ap(this Self&& self, F &&f, Args &&...args)
    noexcept(
        std::is_nothrow_invocable_v<F, value_t, UnwrapOrForward<Args>...>
    )
    → ApInvocableResult<F, error_t, value_t, Args...>
{ /* ... */ };
```

Argument restrictions

```
template <class Result, class Error>
concept NotExpectedOrSameError =
    !is_expected<std::decay_t<Result>> ||
    SameError<Error, std::decay_t<Result>>;
```

```
template <class Error, class ...Args>
concept ValidApArgs = (... && NotExpectedOrSameError<Args, Error>);
```


Function restrictions

```
template <class F, class Error, class ...Args>
concept ApInvocable = requires {
    requires ValidApArgs<Error, Args...>;
    requires std::is_invocable_v<F, UnwrapOrForward<Args>...>;
    requires NotExpectedOrSameError<
        std::invoke_result_t<F, UnwrapOrForward<Args>...>,
        Error>;
};
```

// UnwrapOrForward is Arg or typename std::decay_t<Arg>::value_t

For the result type

```
template <class F, class Error, class ...Args>
using ApInvocableResult =
    typename WrapIntoExpectedOrForward<
        JustInvocableResult<F, Args...>, Error>::type;

// WrapIntoExpectedOrForward is expected<Arg, Error>
// or Arg (if already expected)
```

Handle arguments with errors

Handle errors in arguments

```
if (auto firstError = firstError<error_t>(self, args...)) {  
    return firstError.value();  
}  
else {  
    /* ... */  
}
```

Get first error implementation

```
template <class Error, class ...Args>
    requires (ValidApArgs<Error, Args...>)
constexpr std::optional<Error> firstError(const Args&...args)
{
    return (std::optional<Error>{} << ... << args);
}
```

Get first error implementation (operator <<)

```
template <class E, class Arg>
constexpr auto operator <<(std::optional<E> optErr, Arg &&arg) noexcept → std::optional<E>
{
    if (optErr) {
        return optErr;
    }

    if constexpr (is_expected<std::decay_t<Arg>>) {
        if (arg.has_error()) {
            return std::make_optional<E>(
                std::get<typename std::decay_t<Arg>::error_t>(std::forward<Arg>(arg)));
        }
    }

    return std::nullopt;
}
```

Merge errors

- Semigroup-like combine (associative binary operation)
- $\text{combine}(x, \text{combine}(y, z)) = \text{combine}(\text{combine}(x, y), z)$

Combine two errors example

```
struct Error
{
    std::string data;
};

constexpr auto combine(Error e1, Error e2)
{
    return Error{
        .data = fmt::format("{};{}", std::move(e1.data), std::move(e2.data))
    };
}
```


Combine optional example

```
template <class T>
    requires (is_optional<std::decay_t<T>>())
constexpr std::decay_t<T> combine(T &&a1, T &&a2)
{
    if (a1 && a2) { return combine(*std::forward<T>(a1), *std::forward<T>(a2)); }

    if (a1) { return std::forward<T>(a1); }

    if (a2) { return std::forward<T>(a2); }

    return std::nullopt;
}
```

Combine errors example with operator <<

```
template <class E, class Arg>
    requires (CanCombine<std::optional<std::decay_t<E>>> &&
              CanCombine<std::decay_t<E>>>)
constexpr auto operator <<(std::optional<E> optErr, Arg &&arg) noexcept → std::optional<E>
{
    if constexpr (is_expected<std::decay_t<Arg>>) {
        if (arg.has_error()) {
            return combine(std::move(optErr),
                          std::make_optional<E>(
                              std::get<typename std::decay_t<Arg>::error_t>(std::forward<Arg>(arg))));
        }
    }

    return optErr;
}
```

Continue with ap implementation

Invoke a function

```
if (auto firstError = /* ... */) {  
    /* ... */  
} else {  
    return std::invoke(  
        std::forward<F>(f),  
        std::get<value_t>(std::forward<Self>(self)),  
        unwrapOrForward(std::forward<Args>(args)) ...  
    );  
}
```

Use case

```
Expected openWindow(  
    const WindowNavigator& windowNavigator, const Uri& uri, const State& initialState);  
  
// ...  
  
const fl::expected<gui::Uri, gui::Error> expectedUri = gui::fromString(/* input */);  
const gui::State initialState = gui::getInitialState();  
  
const auto result = gui::getWindowNavigator()  
    .ap(&gui::openWindow, expectedUri, initialState);  
  
// Or  
  
/* ... */ = ap(gui::getWindowNavigator(), &gui::openWindow, expectedUri, initialState);
```

Future improvements

- Can be a part of `and_then` (with a certain tag)
- Can combine errors if supported or return the first one if not

Section conclusions

- Built-in “applicative” is valuable when
 - Some arguments are wrapped into expected
 - There are relatively many arguments
- Not general functionality
- May not belong to the expected interface

Can we add useful functionality without
changing expected class?

Many changes are possible

- Freestanding functions
- Operator |
- Don't use standard interface

Widget example with |

```
getWidget()
```

```
| and_then([](const auto &widget) → std::expected<Widget, WidgetError> {  
    // Do something with the widget ...  
    return widget;  
})  
| transform([](const auto &widget) → ID { return widget.id(); })  
| or_else([](const auto& error) → std::expected<ID, WidgetError> {  
    log(error);  
    // Possibly recover and/or cleanup ...  
    // Return value or error ...  
})
```

Create unwrap-like functionality for tests

- Used with expected
- Fails if an expected object contains an error
- Returns a value if expected object contains a value

Simplified implementation

```
namespace detail {

struct unwrap_t {
    template <IsExpected Expected>
    [[nodiscard]] friend constexpr decltype(auto) operator |(Expected &&e, detail::unwrap_t)
    {
        if (e.has_error()) {
            FAIL(fmt::format("Expected object contains an error {:?}.", e.error()));
        }

        return std::forward<Expected>(e).value();
    }
};

}

inline constexpr detail::unwrap_t unwrap;
```

Usage example of unwrap

```
TEST_CASE("...")
{
    const Shape shape = /* ... */;

    /* ... */

    const auto area = calculateArea(shape) | unwrap;

    REQUIRE(area == /* ... */);
}
```

Section conclusions

- All existing monadic operations can be implemented via operator |
- Many extra monadic operations can be implemented via operator |
- You don't need to re-create `std::expected` from scratch in this case

Final thoughts

Conclusions

- Current interface of `std::expected` has some issues
- Some of the issues result in error-prone code
- The interface lacks some practically useful functionality
- Most likely you may need to create your own "expected"
- Creating a pipe-like interface for `std::expected` can be a good compromise

Thank you!