

C++ now

2025

Getting the Lazy Task Done

Dietmar Kühl

Getting The Lazy Task Done

Engineering

Bloomberg

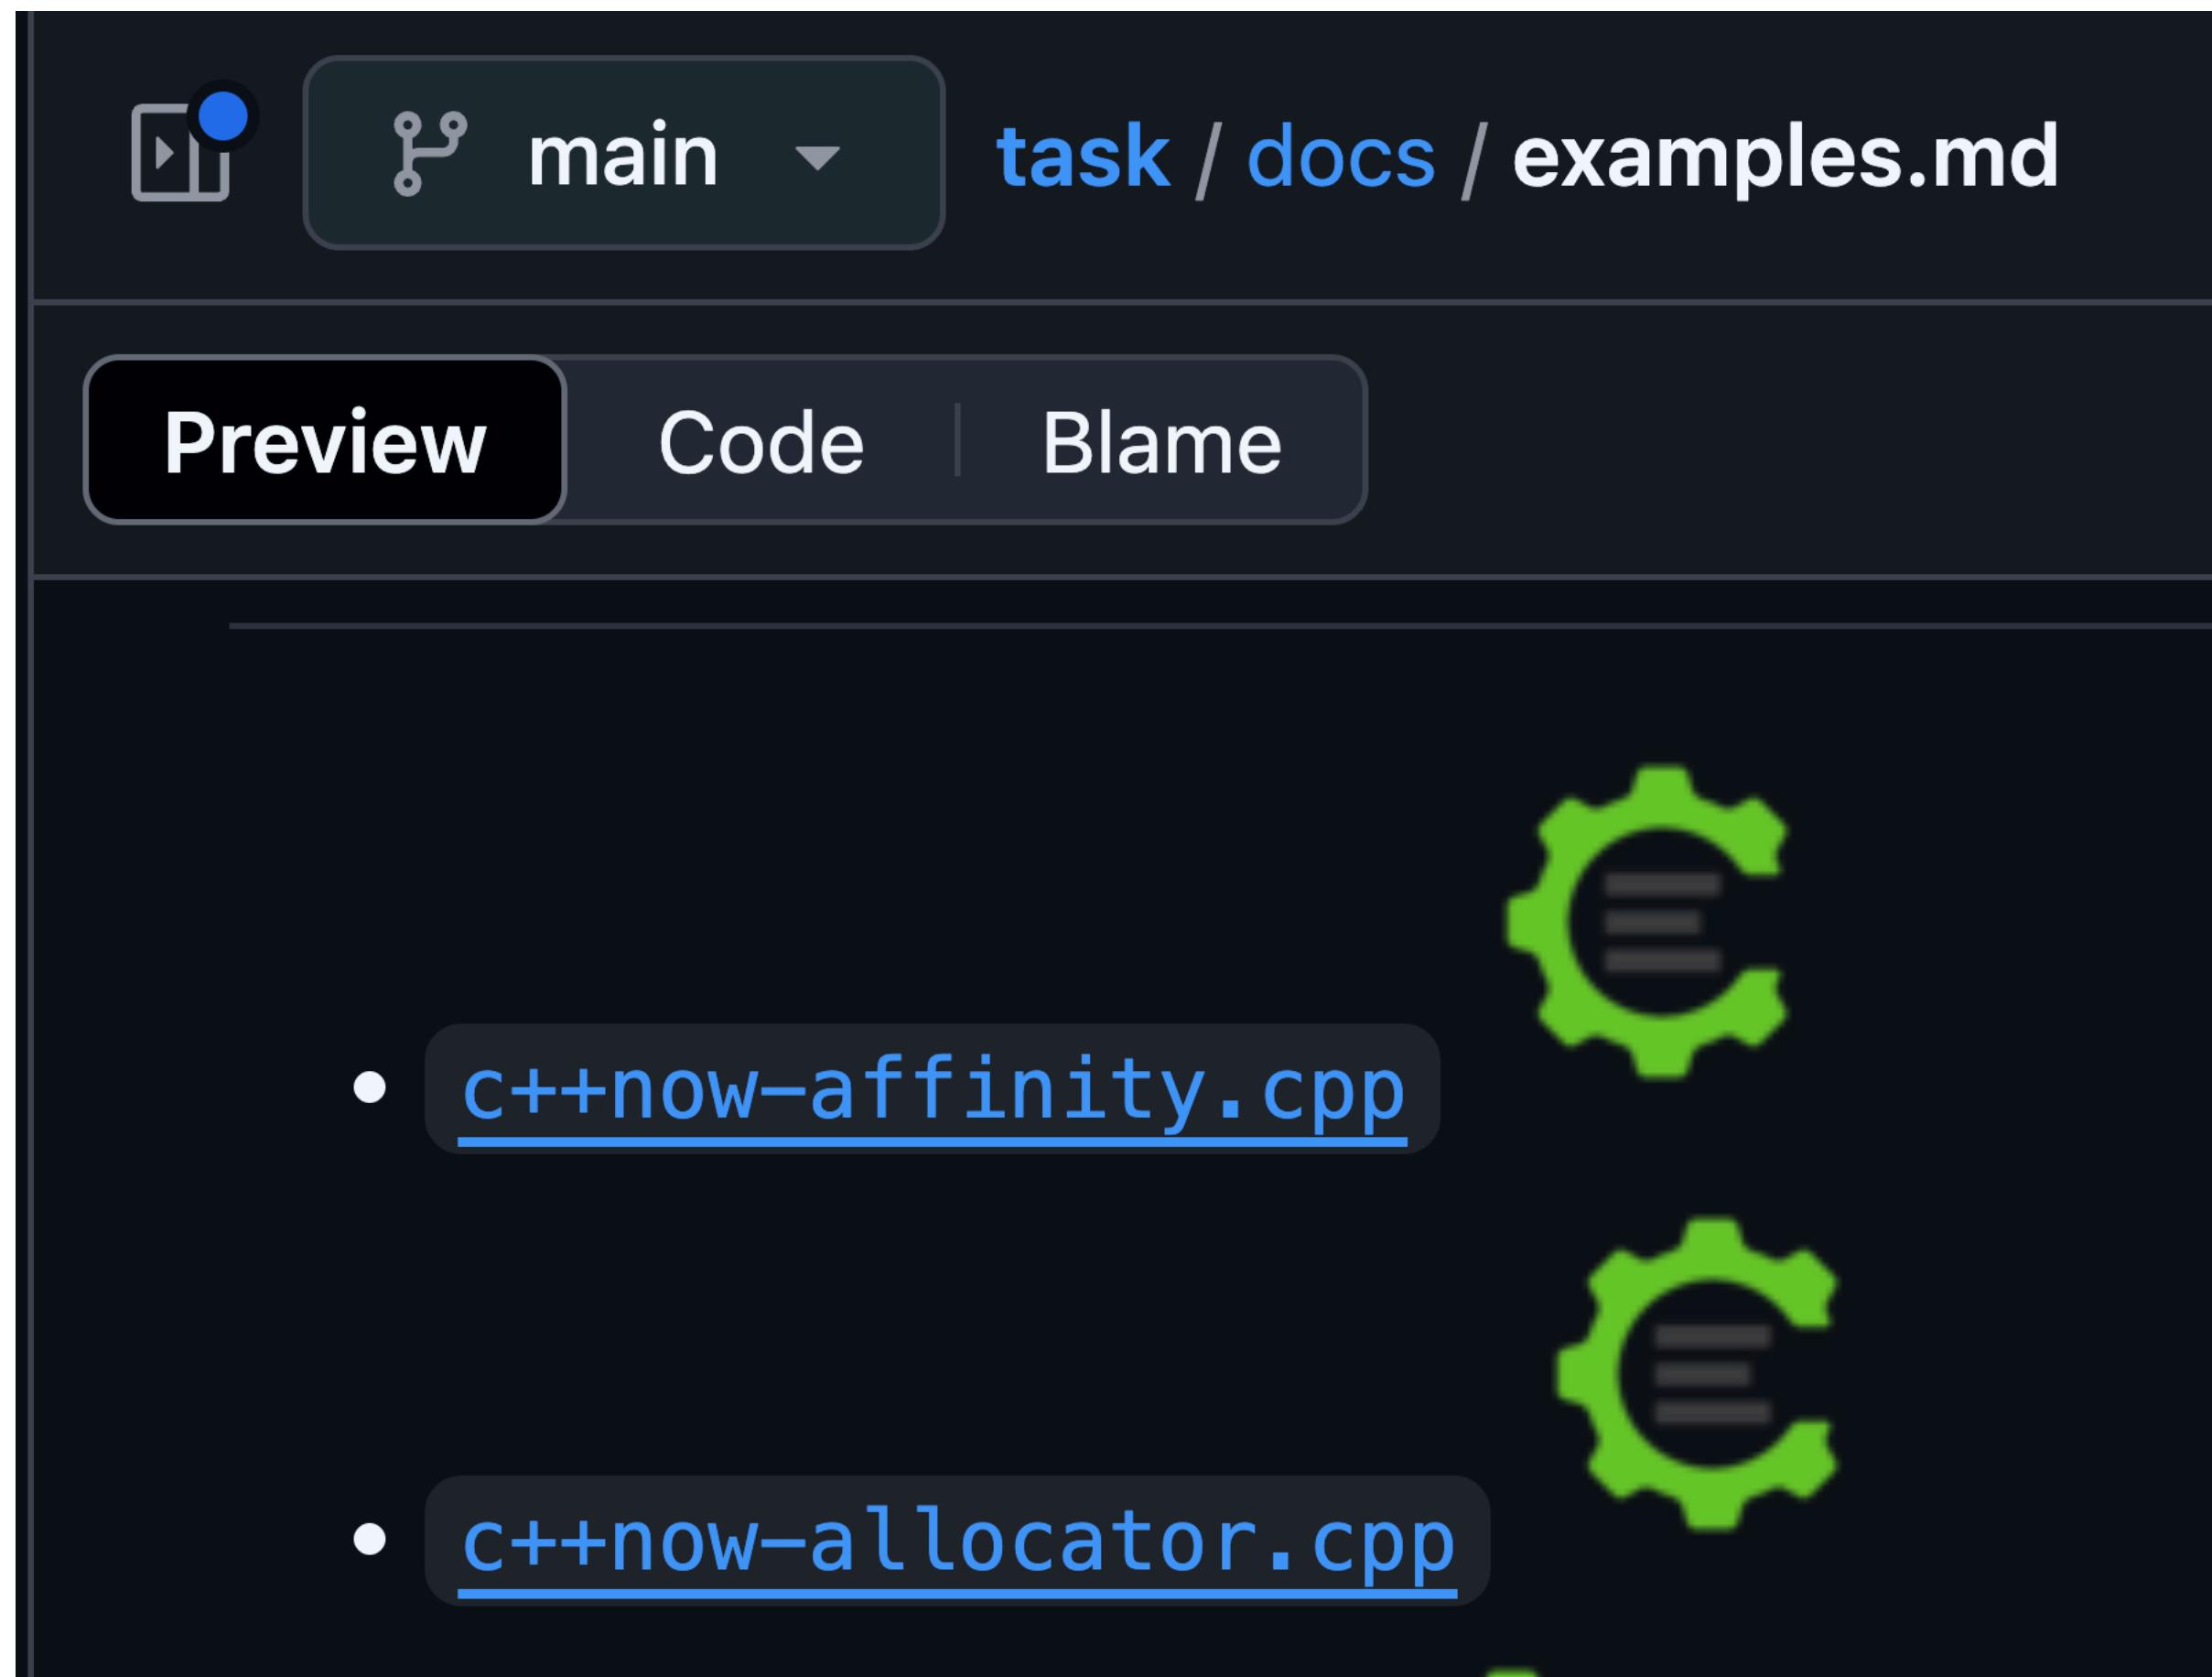
C++Now
2025-05-01

Dietmar Kühl

dkuhl@bloomberg.net

TechAtBloomberg.com

Live Code Examples



Objective

- `std::execution` is bound to become part of C++26
- users most likely want to use coroutines to create senders
- create a coroutine task type for the standard C++ library

Motivation

- Composing work from algorithms is possible but unfamiliar
 - Writing a function with some `co_*` keywords is more familiar
 - Writing a coroutine is often a lot easier
- Coroutines provide a type-erased interface

Sender/Receiver in a Nutshell

- sender: represent work
`S::comp_signs<set_value_t(v...), set_error(e), set_stopped>`
`s.connect(receiver)`
- receiver: consume results from a sender
`r.set_value(v...), r.set_error(e), r.set_stopped()`
`r.get_env()`
- operation state = `connect(sender, receiver)`
`o.start()`

Basics

```
#include <execution>
```

```
std::execution::task<> basic() {
    co_await std::suspend_never{};
}
```

```
int main() { std::execution::sync_wait(basic()); }
```

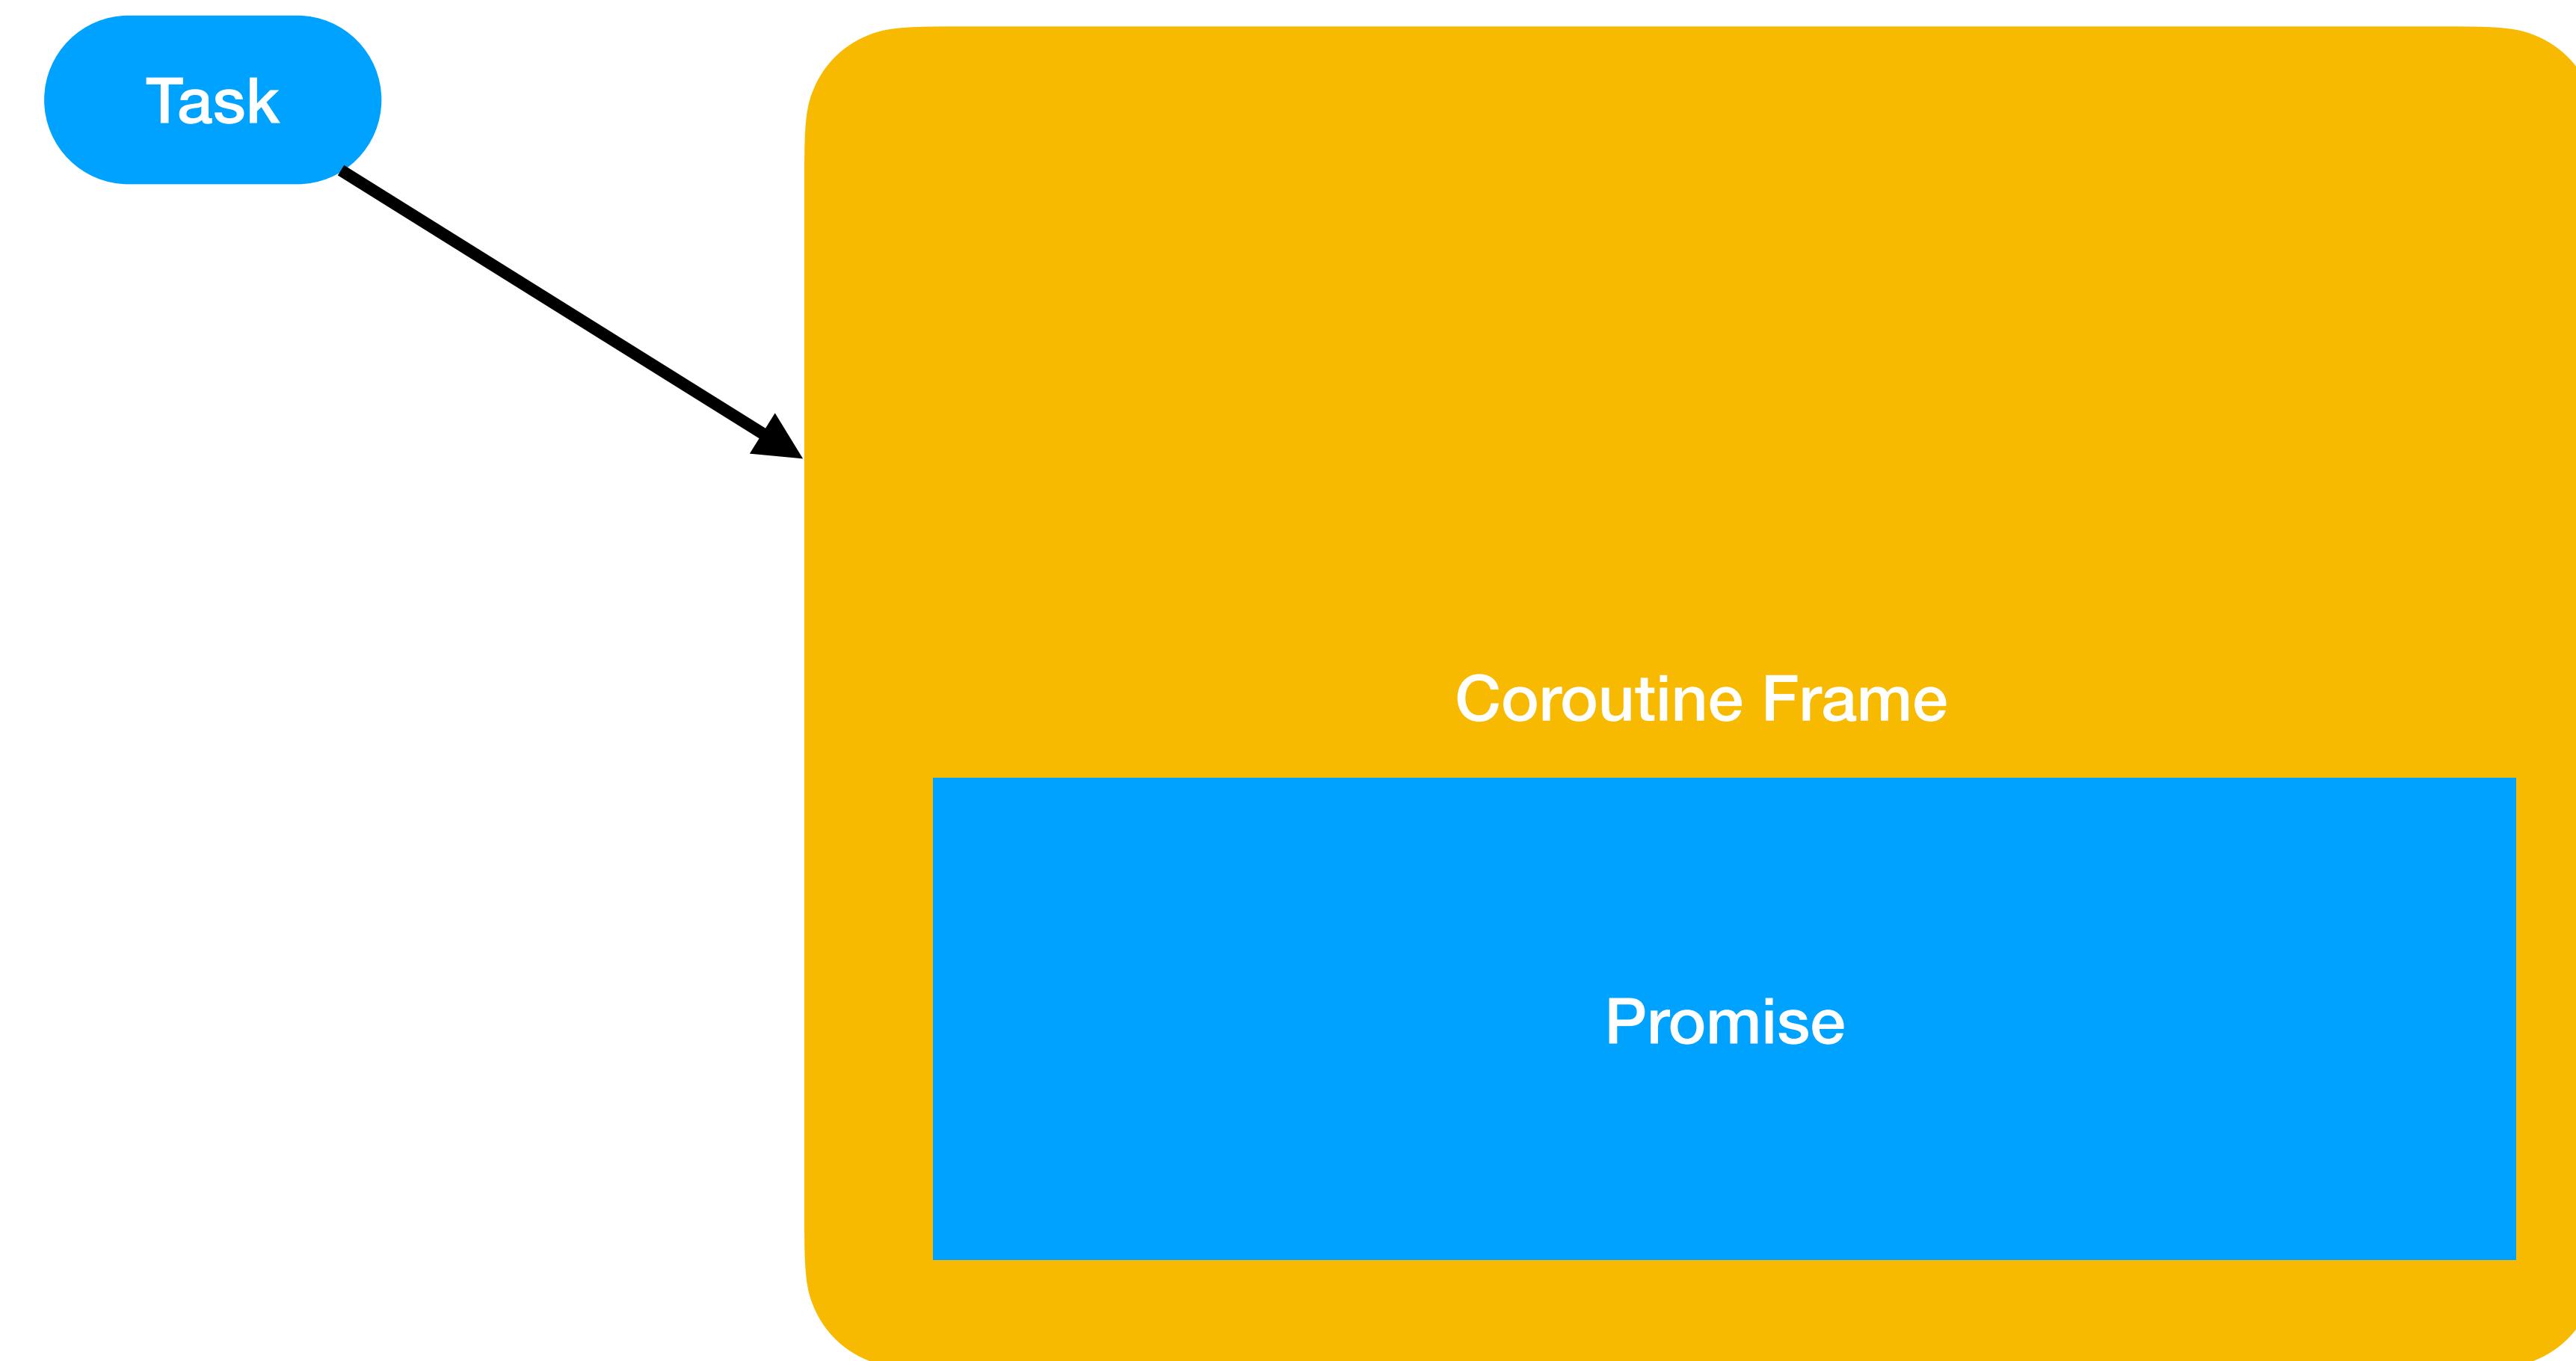
Basics

```
#include <execution>

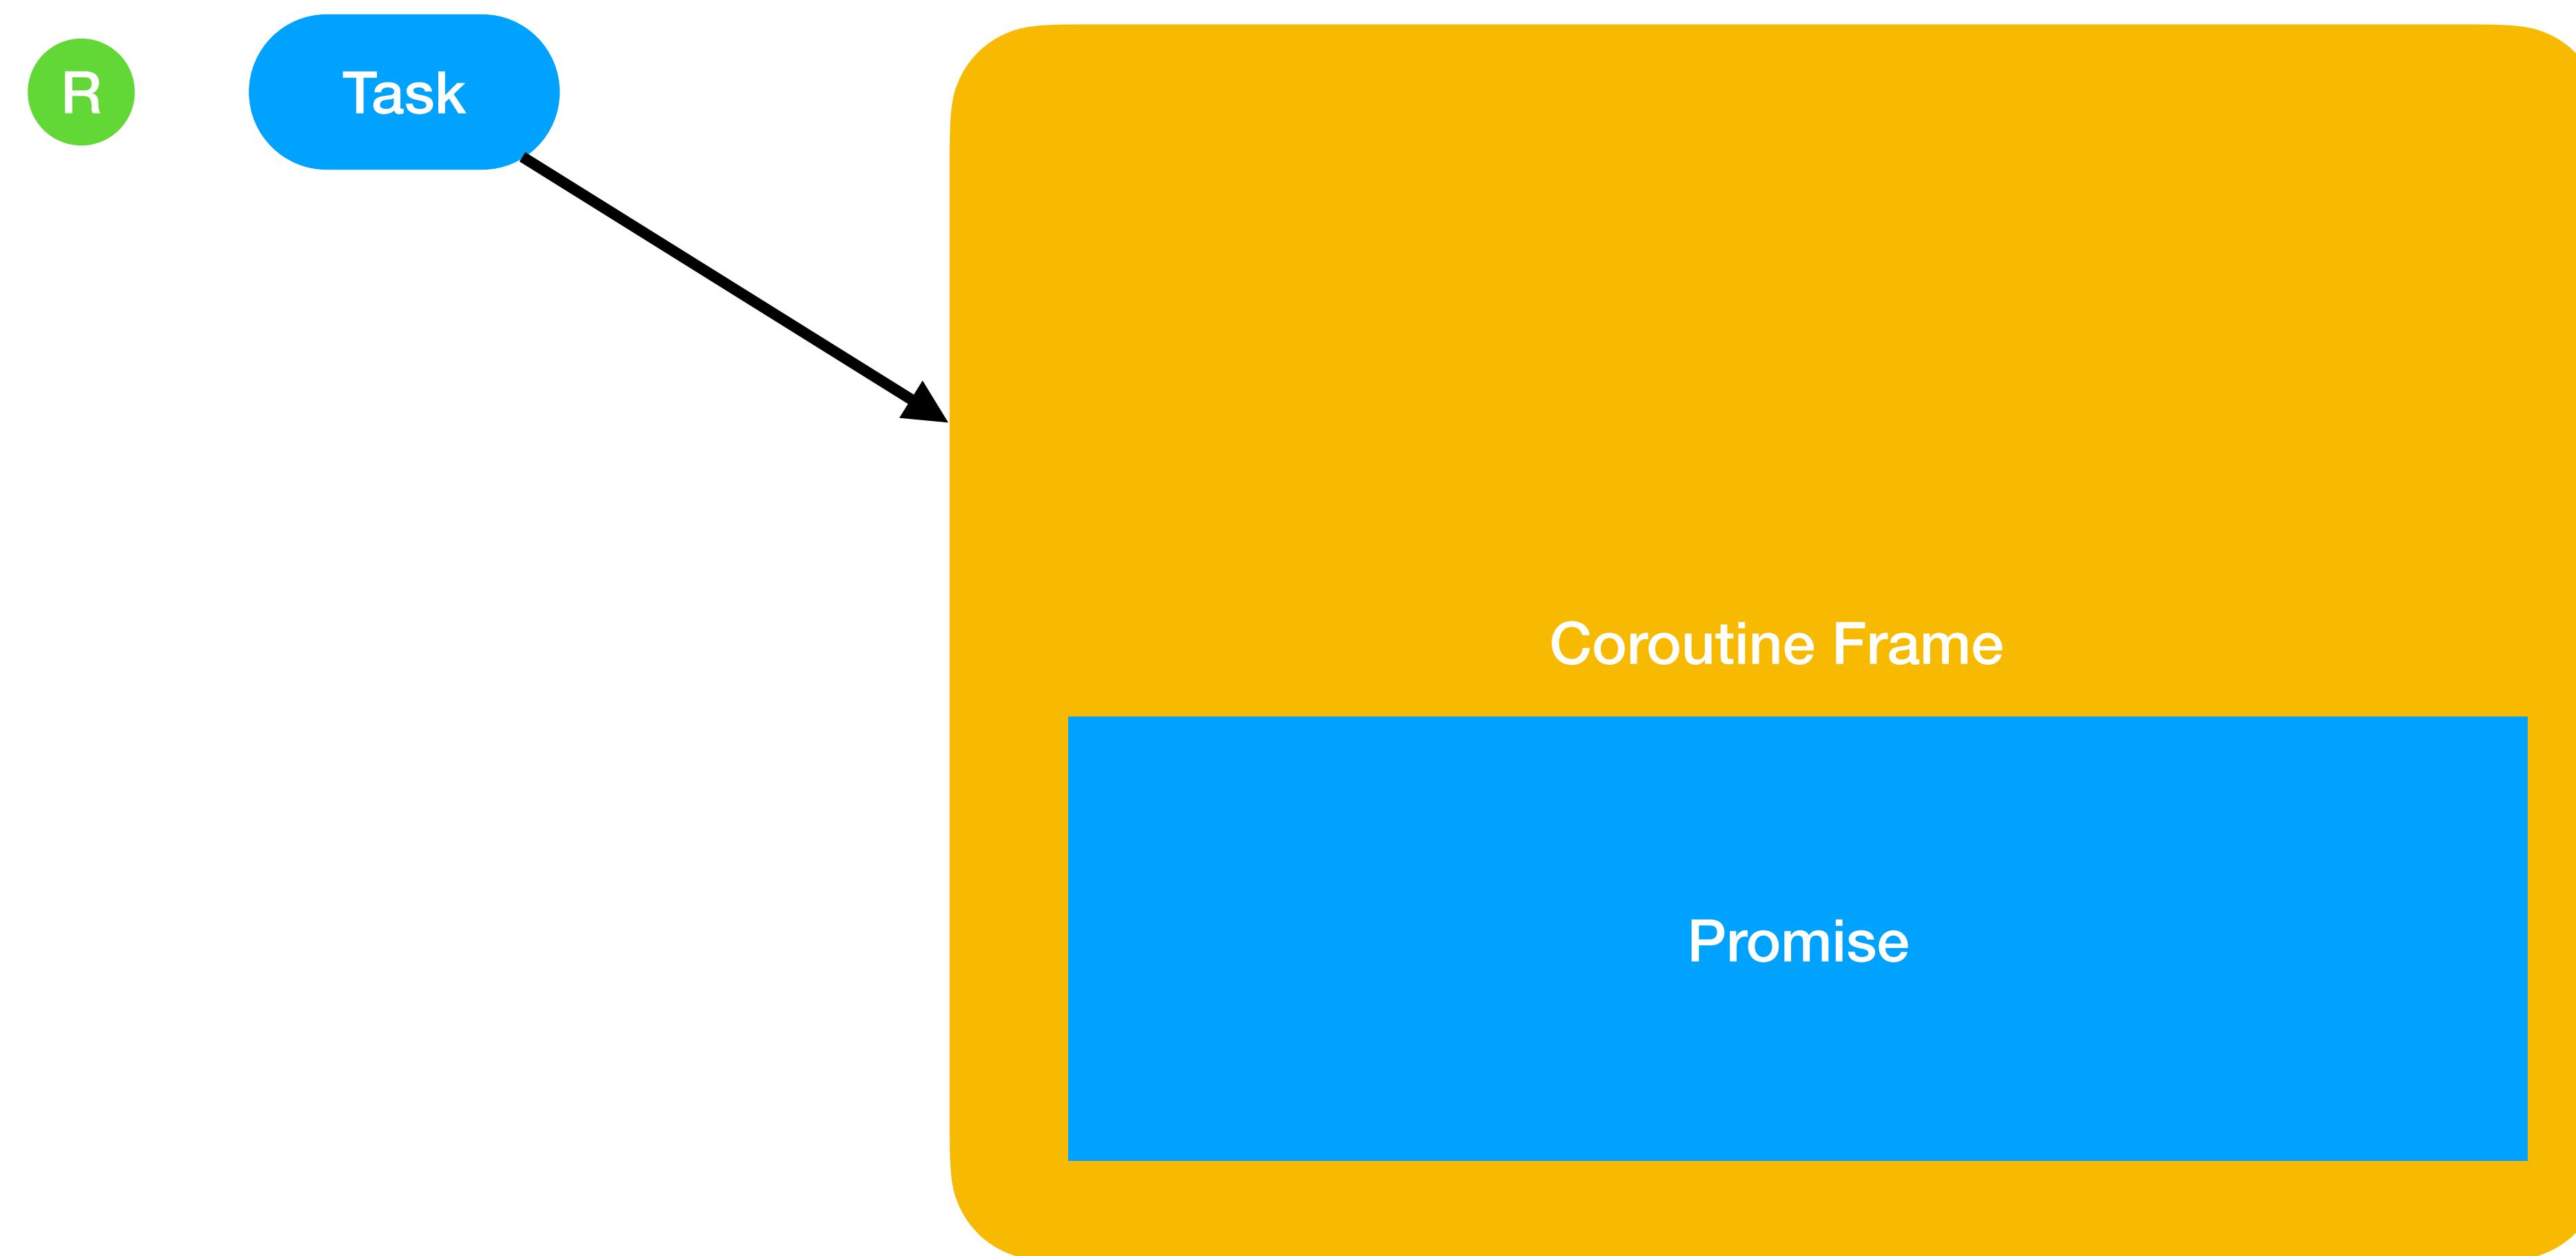
std::execution::task<> basic() {
    co_await std::execution::just();
}

int main() { std::execution::sync_wait(basic()); }
```

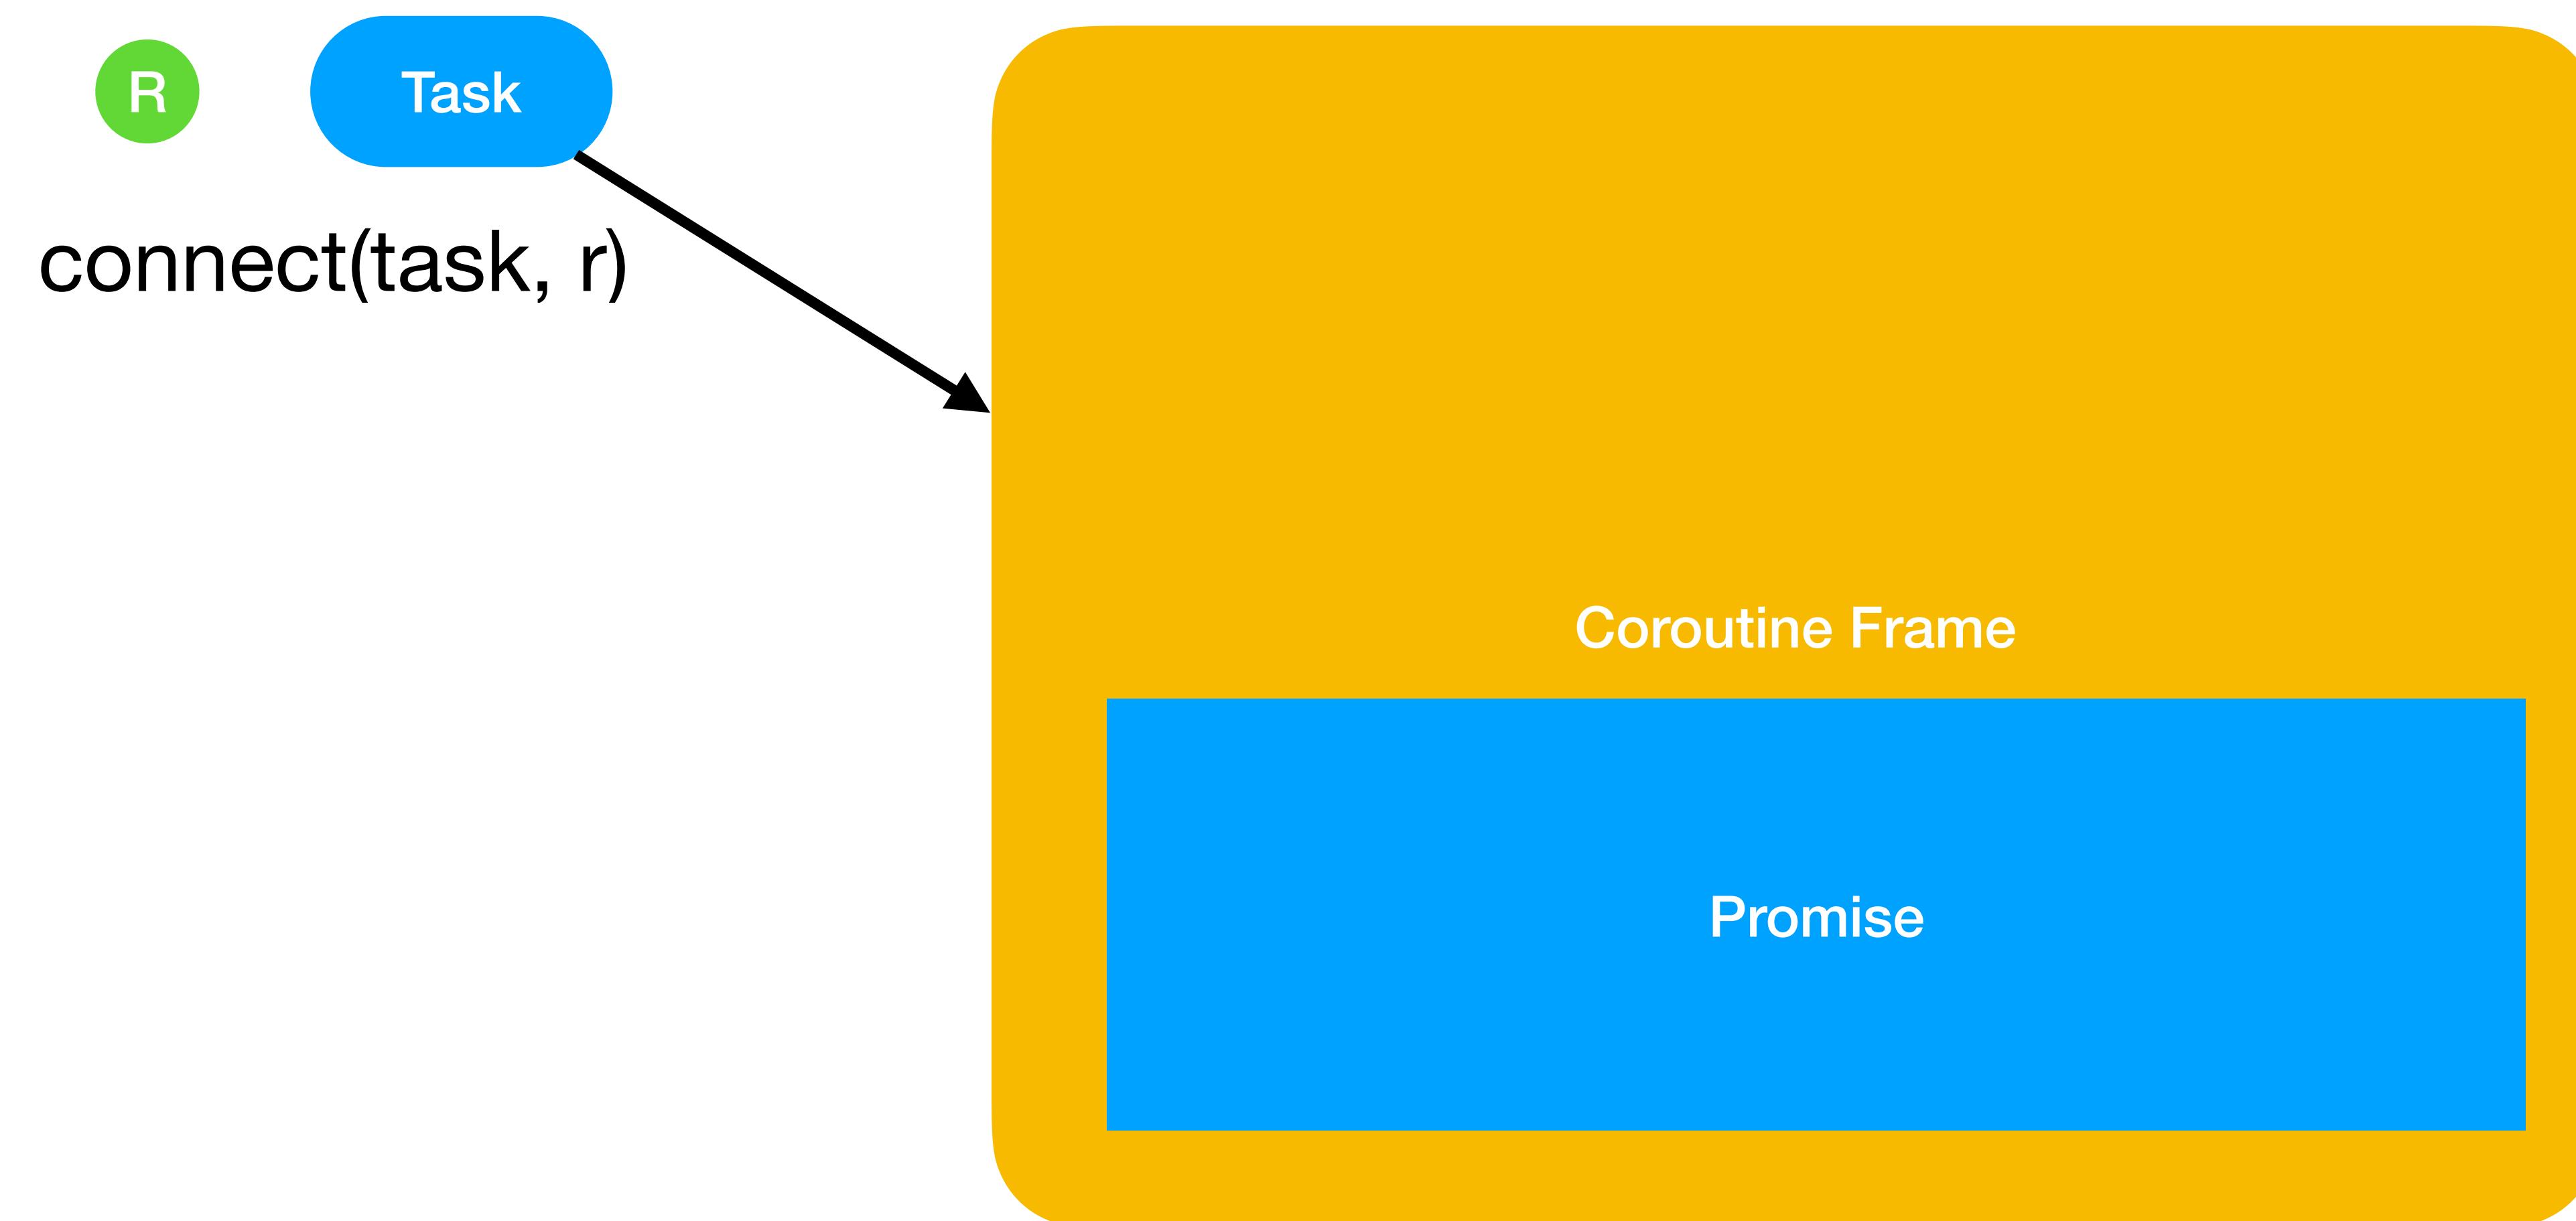
Task is a Sender



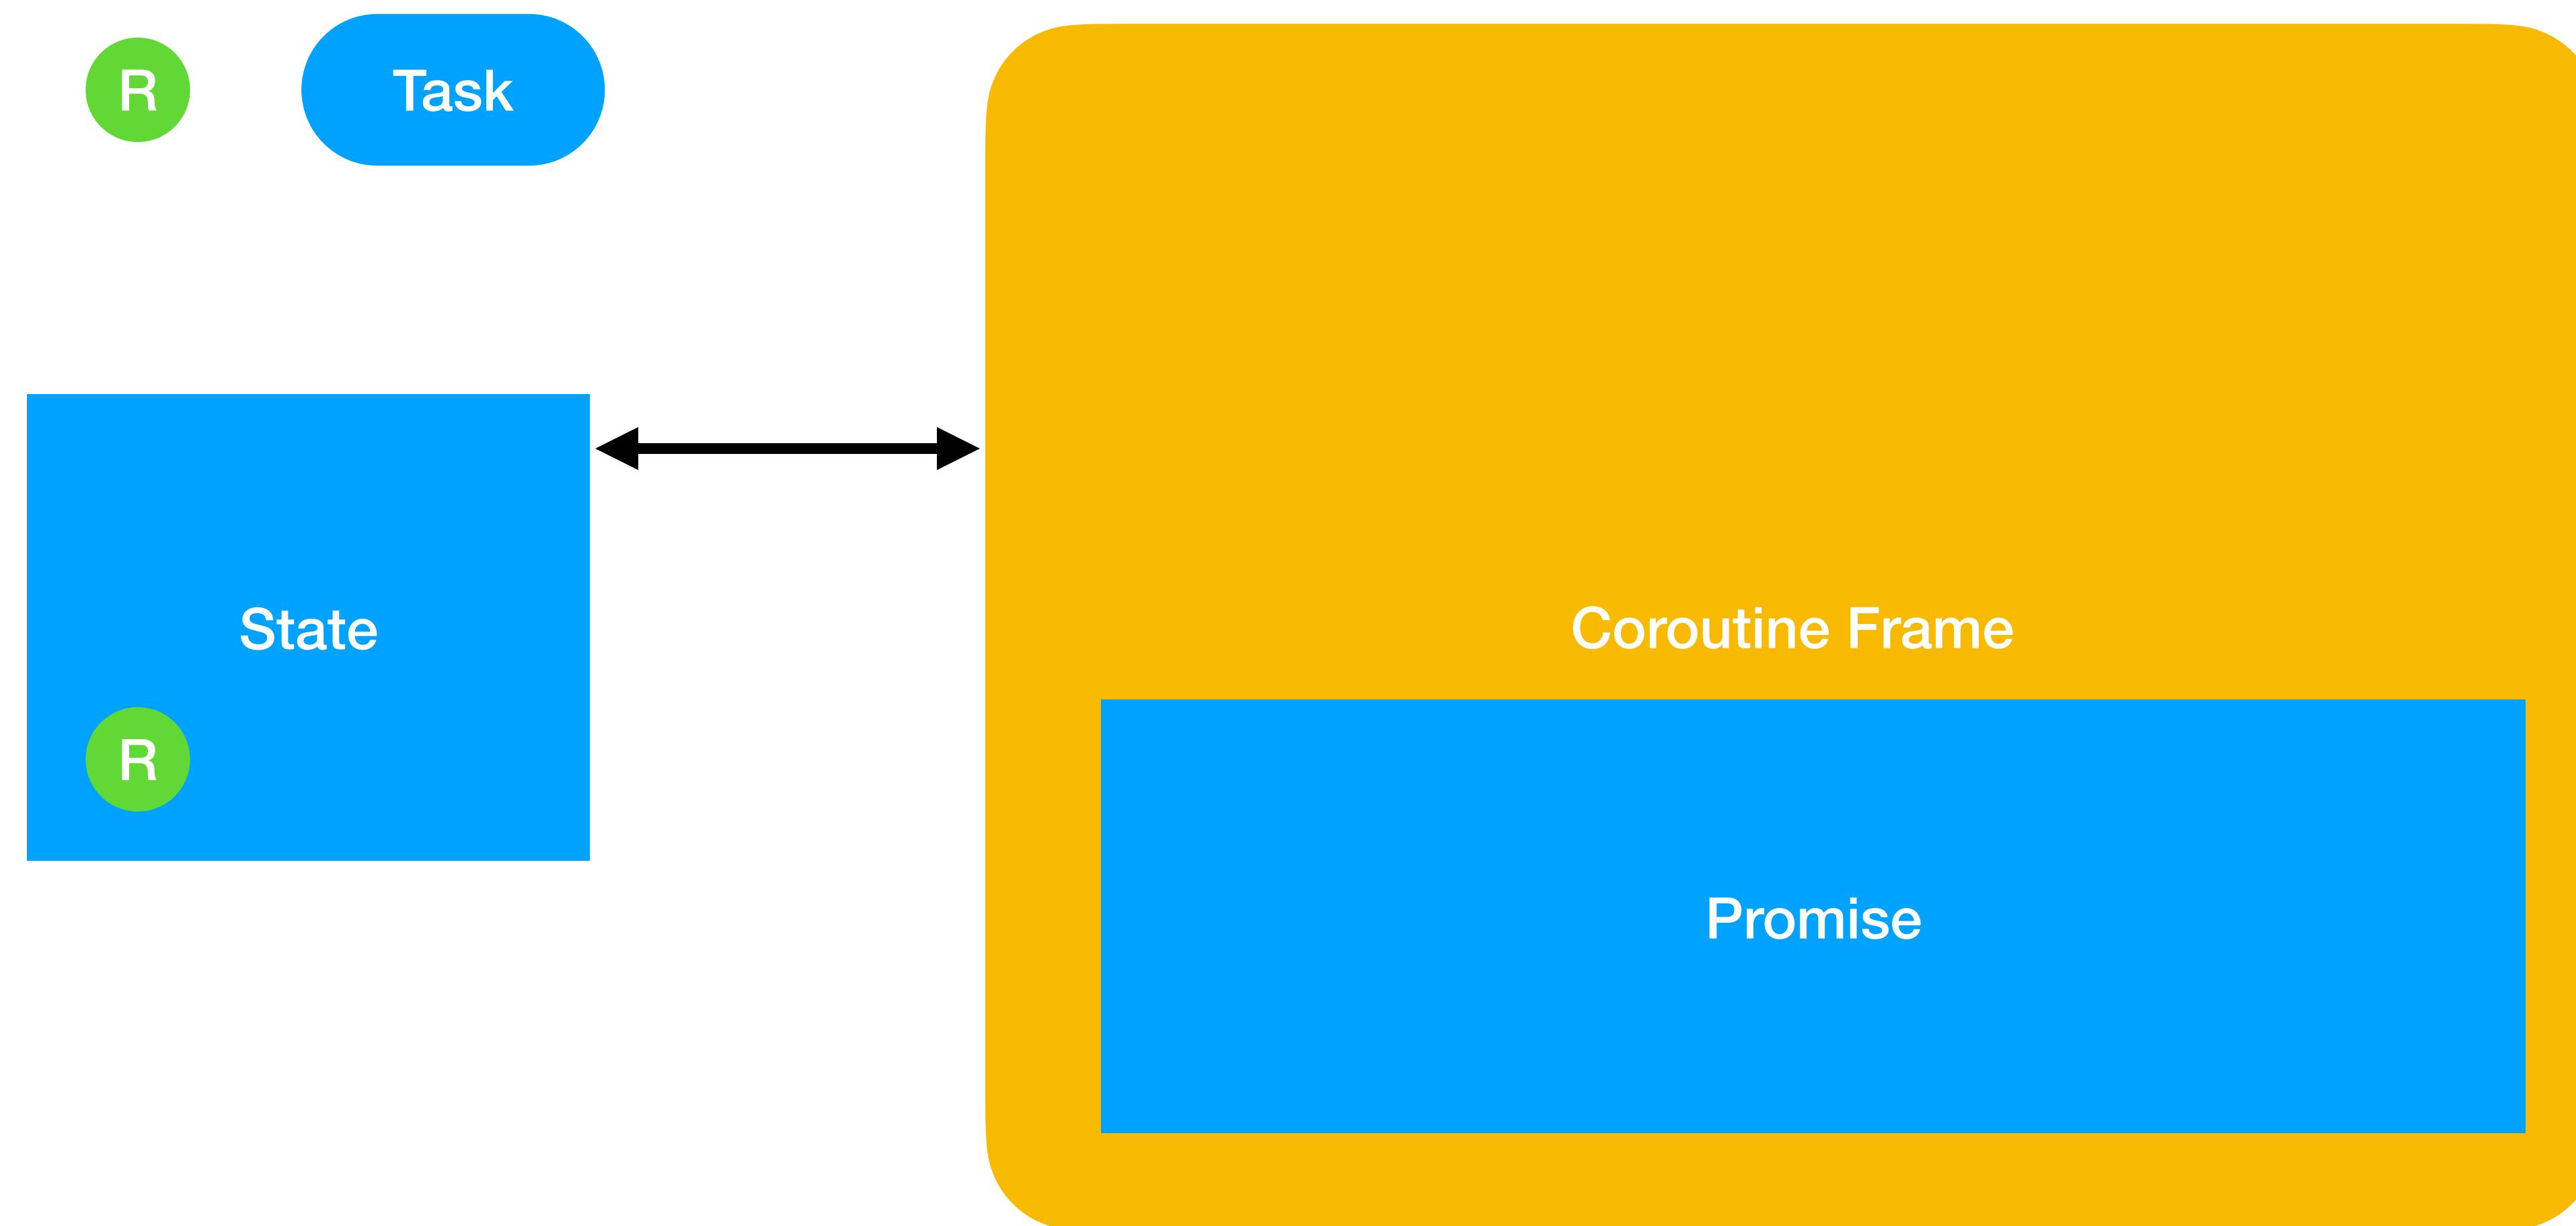
Task is a Sender



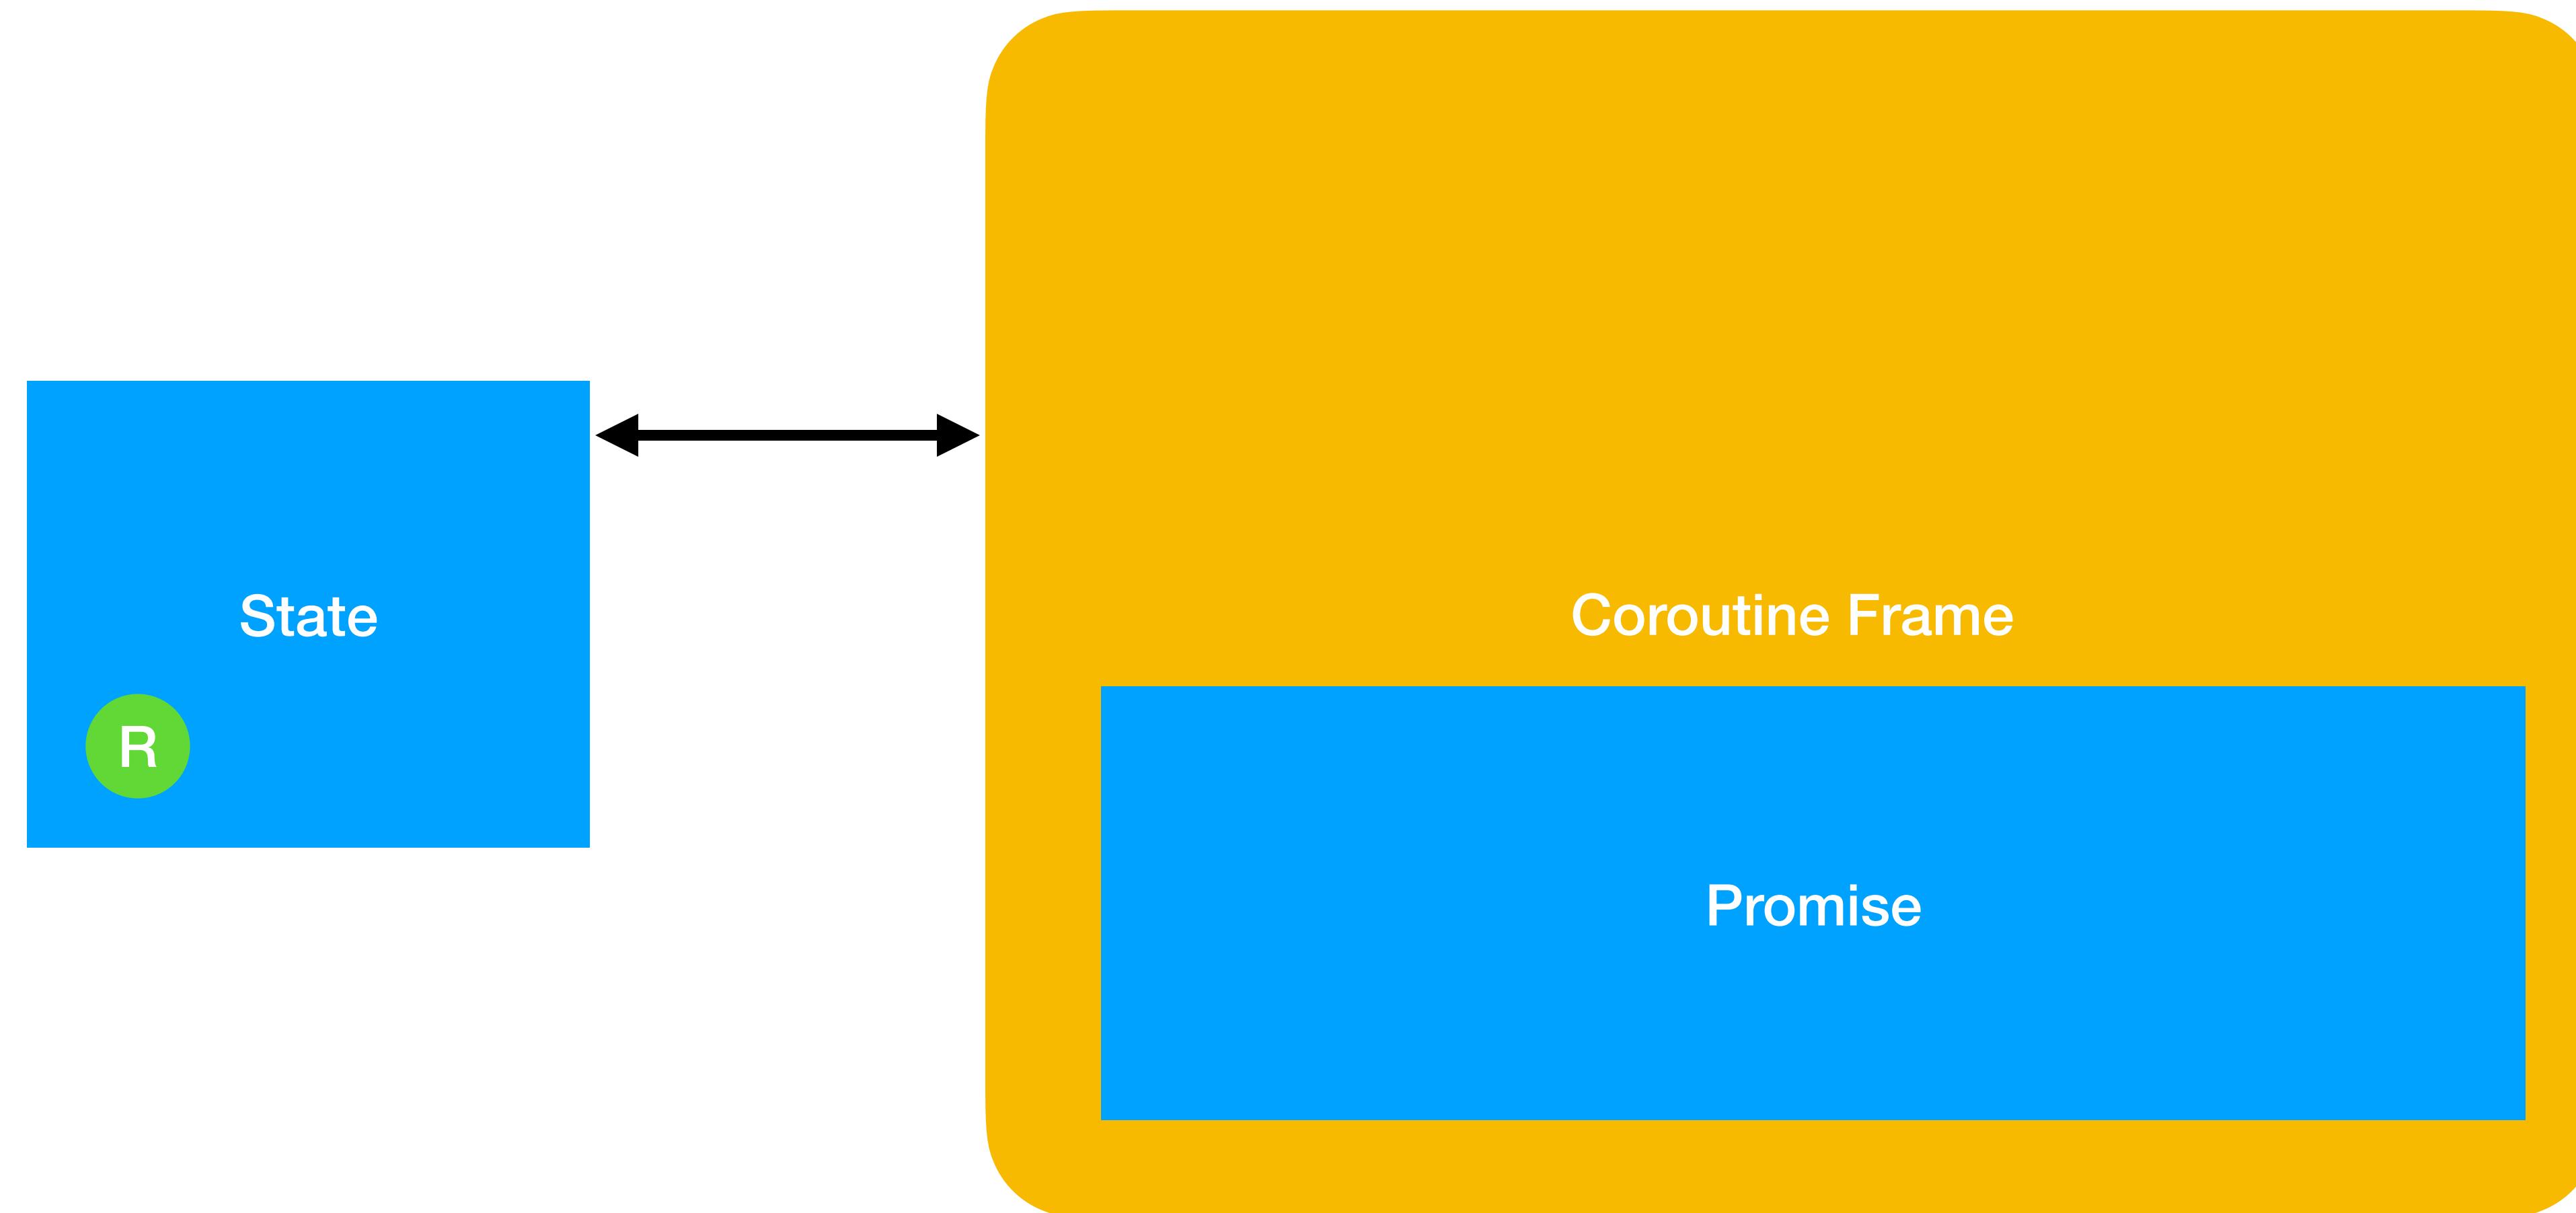
Task is a Sender



Task is a Sender



Task is a Sender



Result Types

```
#include <execution>

std::execution::task<> result_types() {
    /* void */    co_await std::execution::just();
    int          n = co_await std::execution::just(17);
    std::tuple   t = co_await std::execution::just(17, true);
}

int main() { std::execution::sync_wait(result_types()); }
```

Errors

```
#include <execution>
```

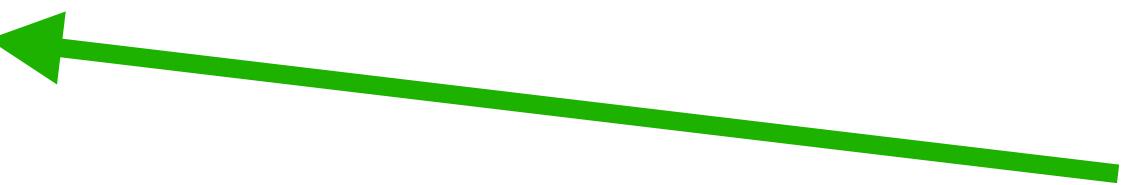
```
std::execution::task<> error_result() {
    try { co_await std::execution::just_error(17); }
    catch (int n) { std::print("Error: {}\n", n); }
}
```

```
int main() { std::execution::sync_wait(error_result()); }
```

Errors Without Exception

```
#include <execution>
```

```
std::execution::task<> error_result() {
    auto e = co_await as_expected(std::execution::just(17));
    auto u = co_await as_expected(std::execution::just_error(17));
}
```



not in the C++ standard library

```
int main() { std::execution::sync_wait(error_result()); }
```

Cancellation

```
#include <execution>

std::execution::task<> cancel() {
    co_await std::execution::just_stopped(); // never resumes
}

int main() { std::execution::sync_wait(cancel()); }
```

as_awaitable(sender, promise)

- The results of awaiting expressions in a task come from `as_awaitable()`
- `std::execution::as_awaitable(sender, promise)` came from P2300
- Should there be additional/different transformations?
 - Problem: everybody wants their specific transformations
 - Conclusion: the transformations should be explicit

Task's Return Type

```
#include <execution>

std::execution::task<> result() {
    co_return;
}

int main() {
    std::execution::sync_wait(result());
}
```

Task's Return Type

```
#include <execution>

std::execution::task<void> result() {
    co_return;
}

int main() {
    std::execution::sync_wait(result());
}
```

Task's Return Type

```
#include <execution>

std::execution::task<int> result() {
    co_return 17;
}

int main() {
    auto[n] = *std::execution::sync_wait(result());
}
```

Task's Completion Signatures

task<void>

- std::execution::set_value_t()
- std::execution::set_error_t(std::exception_ptr)
- std::execution::set_stopped_t()

Task's Completion Signatures

task<T>

- std::execution::set_value_t(T)
- std::execution::set_error_t(std::exception_ptr)
- std::execution::set_stopped_t()

Task's Completion Signatures

```
struct context {  
    using error_types = comp_sigs<set_error_t(E), set_error_t(F)>;  
};  
task<T, context>
```

- std::execution::set_value_t(T)
- std::execution::set_error_t(E), std::execution::set_error_t(F)
- std::execution::set_stopped_t()

Task's Completion Signatures

```
struct context {  
    using error_types = comp_sigs<set_error_t(E), set_error_t(F)>;  
};  
task<T, context>
```

No `set_error_t(exception_ptr)` but an exception \Rightarrow `terminate()`.

- It can't be statically determined if an exception may be thrown.
- `context::error_types` can be empty \Rightarrow no error completions.

Completing With An Error

```
struct ctxt { using error_types = comp_sigs<set_error_t(int)>; }

task<int, ctxt> call(int v) {
    if (v == 1) co_yield with_error(-1);
    co_return 2 * v;
}
int main(int ac, char* {
    auto[n] = *sync_wait(
        call(ac) | upon_error([](int e){ print("error({})\n", e); return -1; }));
}
```

Why Use co_yield?

- `co_return with_error(e)` can't be used for coroutines returning `void`:
 - promise can't have both `return_value(T)` and `return_void()`.
- `co_await` would work but is used to await outstanding work.
- `co_yield` conceptually yields a result.
- Other than that, only exceptions would work.
 - The whole point is to allow avoiding exceptions.

Allocator Support

- Memory for the coroutine may need to be allocated
- Let C fun(A&&...) be the signature used to create a coroutine
- There are two helpful interfaces which are used when defined:
 - C::promise_type::operator new(size_t, A&&...)
 - C::promise_type::promise_type(A&&...)

Allocator Support Interface

1. `context::allocator_type` can be used to define an allocator type
 - By default the type is `std::allocator<std::byte>`
2. If there is an `std::allocator_arg_t` parameter
 - use the next parameter to initialise an `allocator_type` object
3. Use the `allocator_type` object to manage memory

`std::destroying_delete_t`

- `operator delete(T* ptr, std::destroying_delete_t):`
 - `*ptr` needs to be destroyed and memory needs to be released
- The allocator can be retrieved from `*ptr` before destroying `*ptr`
- Sadly, doesn't work for coroutine promise_types!
 - The allocated object has `promise_type` just as a member

Allocator Example

```
struct with_alloc {
    using allocator_type = pmr::polymorphic_allocator<std::byte>;
};

task<int, with_alloc> coro(int value, auto&&...) {
    co_await just(value);
}

int main() {
    sync_wait(coro(0));
    sync_wait(coro(1, allocator_arg, pmr::new_delete_resource()));
}
```

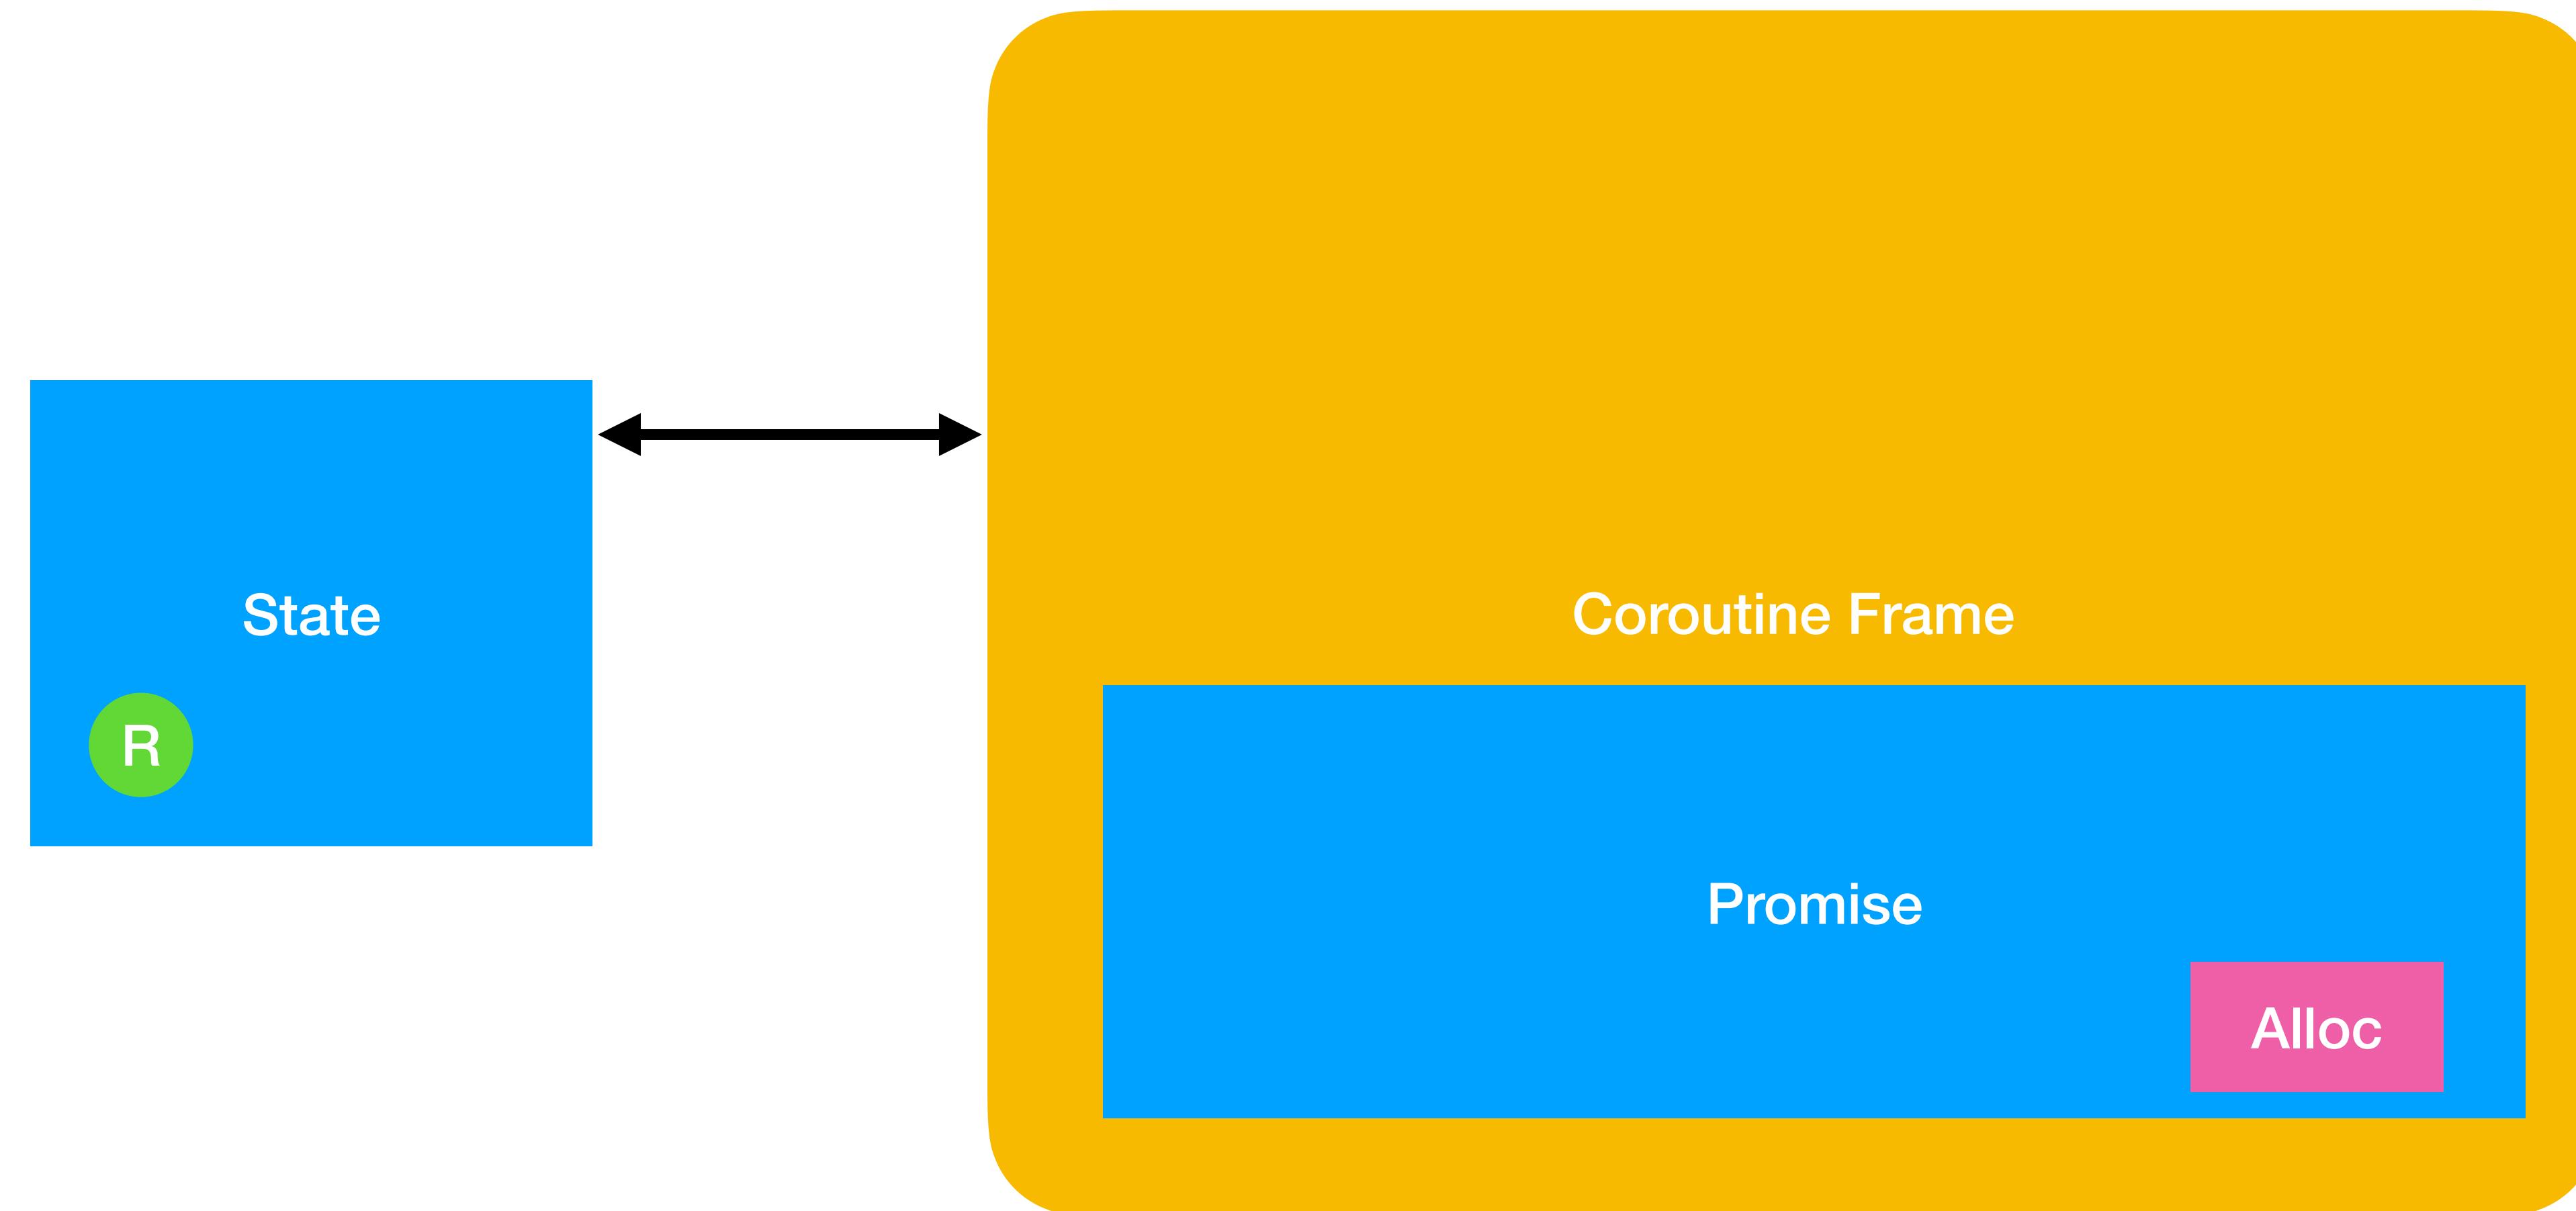
Allocator From Environment

```
struct with_alloc {
    using allocator_type = pmr::polymorphic_allocator<std::byte>;
};

task<int, with_alloc> coro(int value, auto&&...) {
    auto a = co_await execution::read_env(execution::get_allocator);
}

int main() {
    sync_wait(coro(0));
    sync_wait(coro(1, allocator_arg, pmr::new_delete_resource()));
}
```

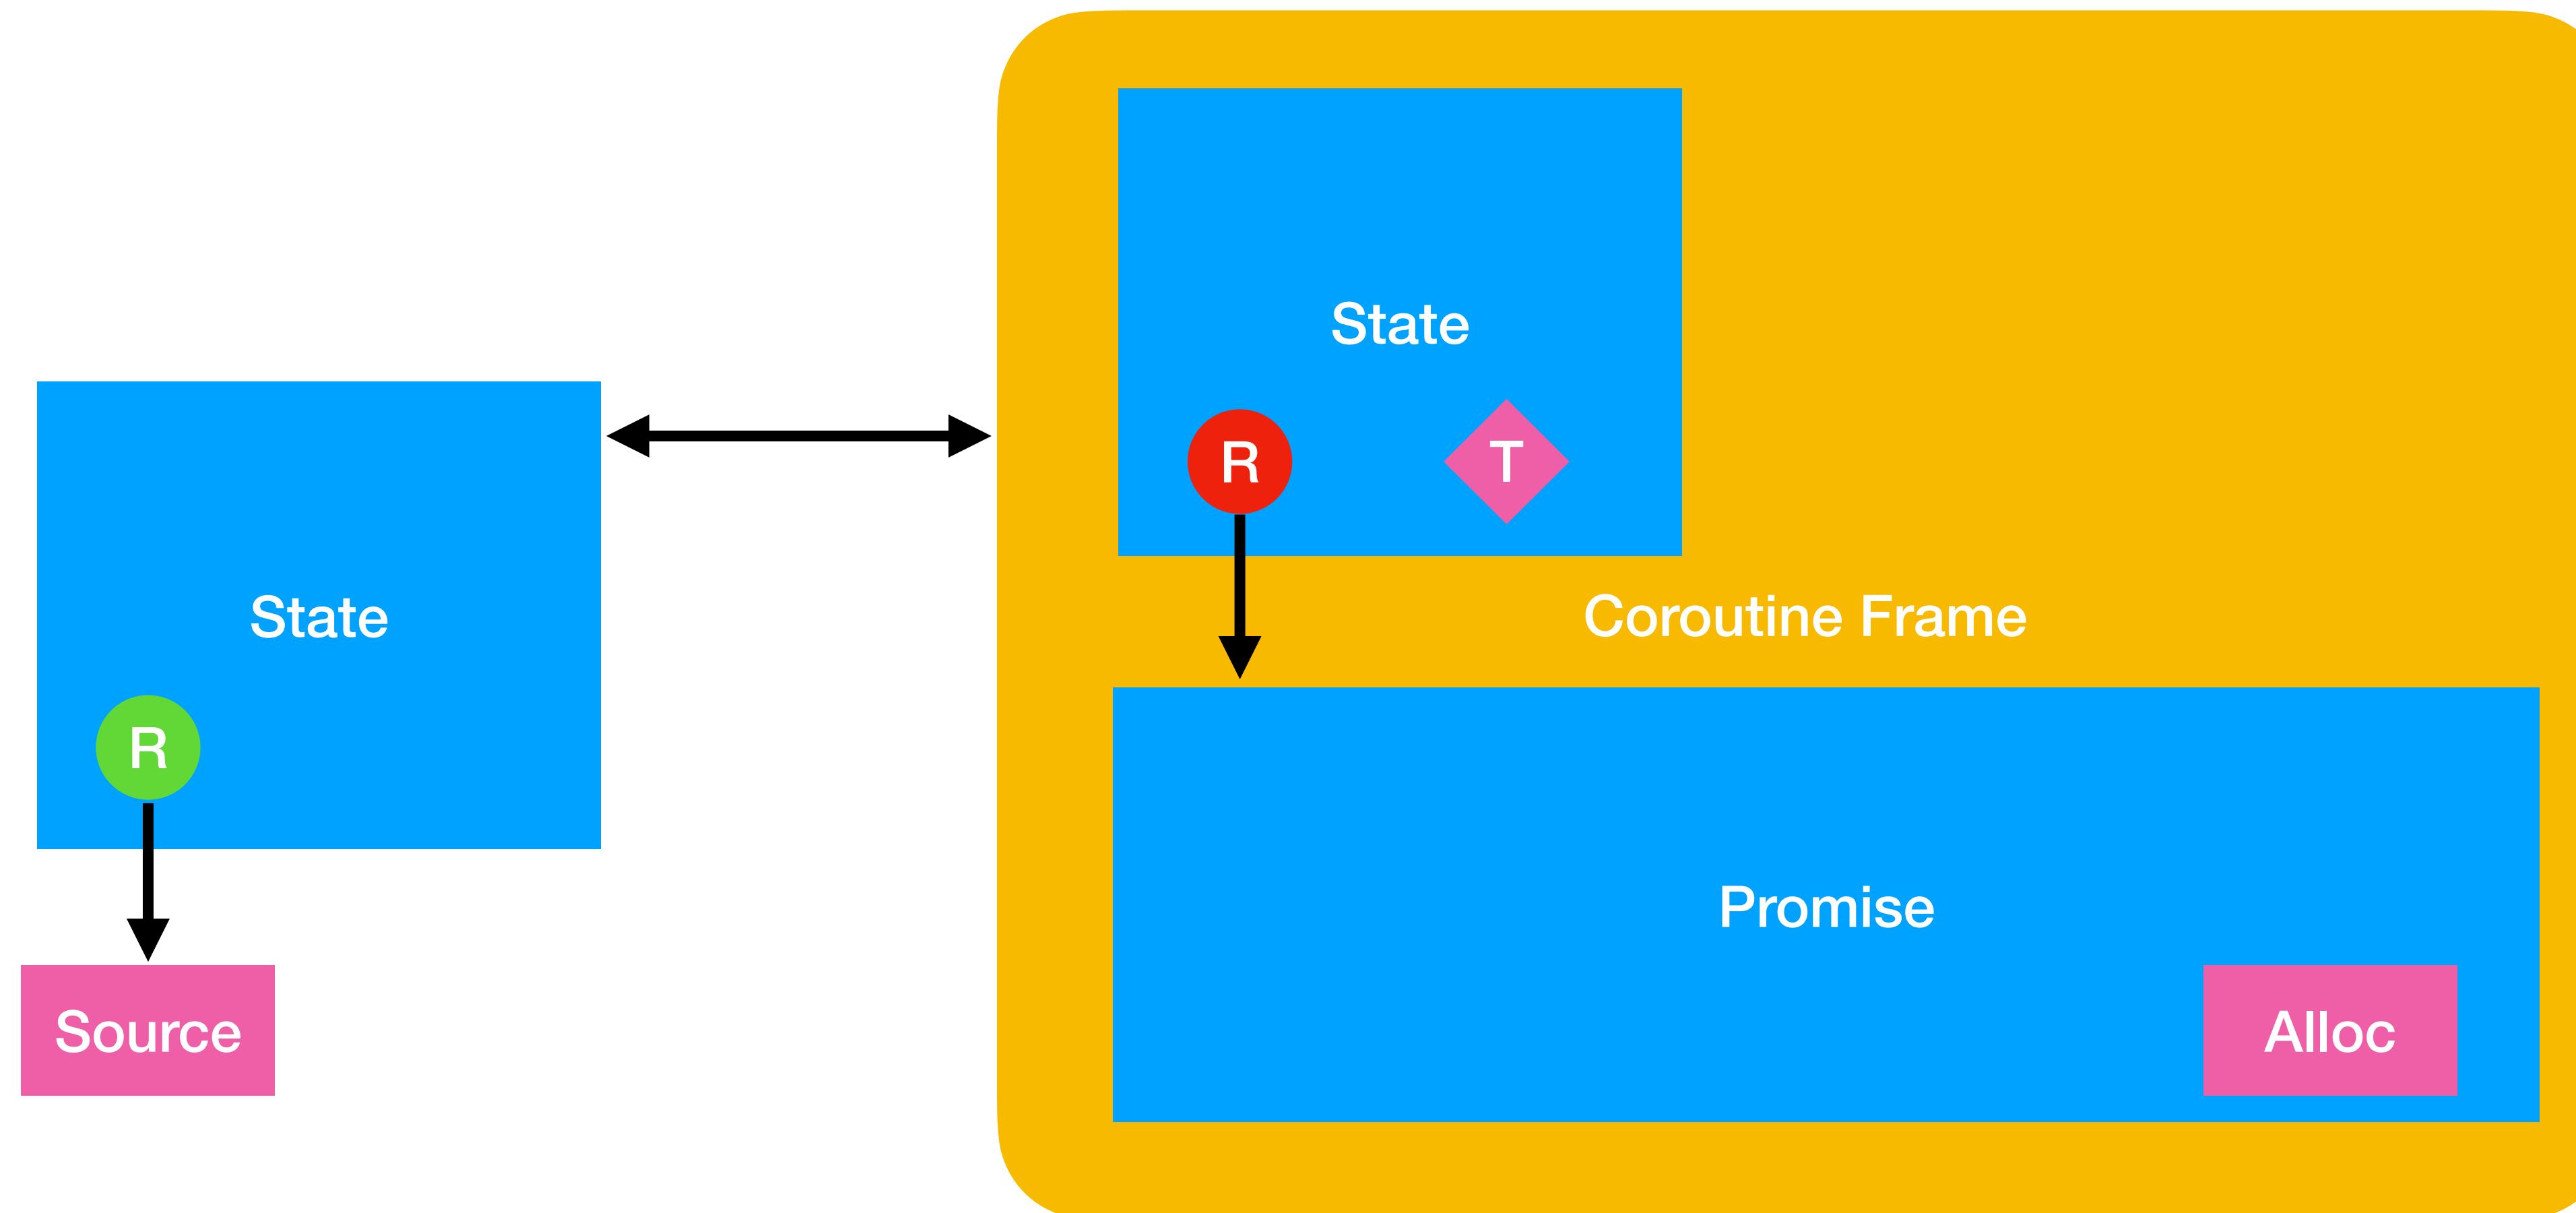
Task Holds an Allocator



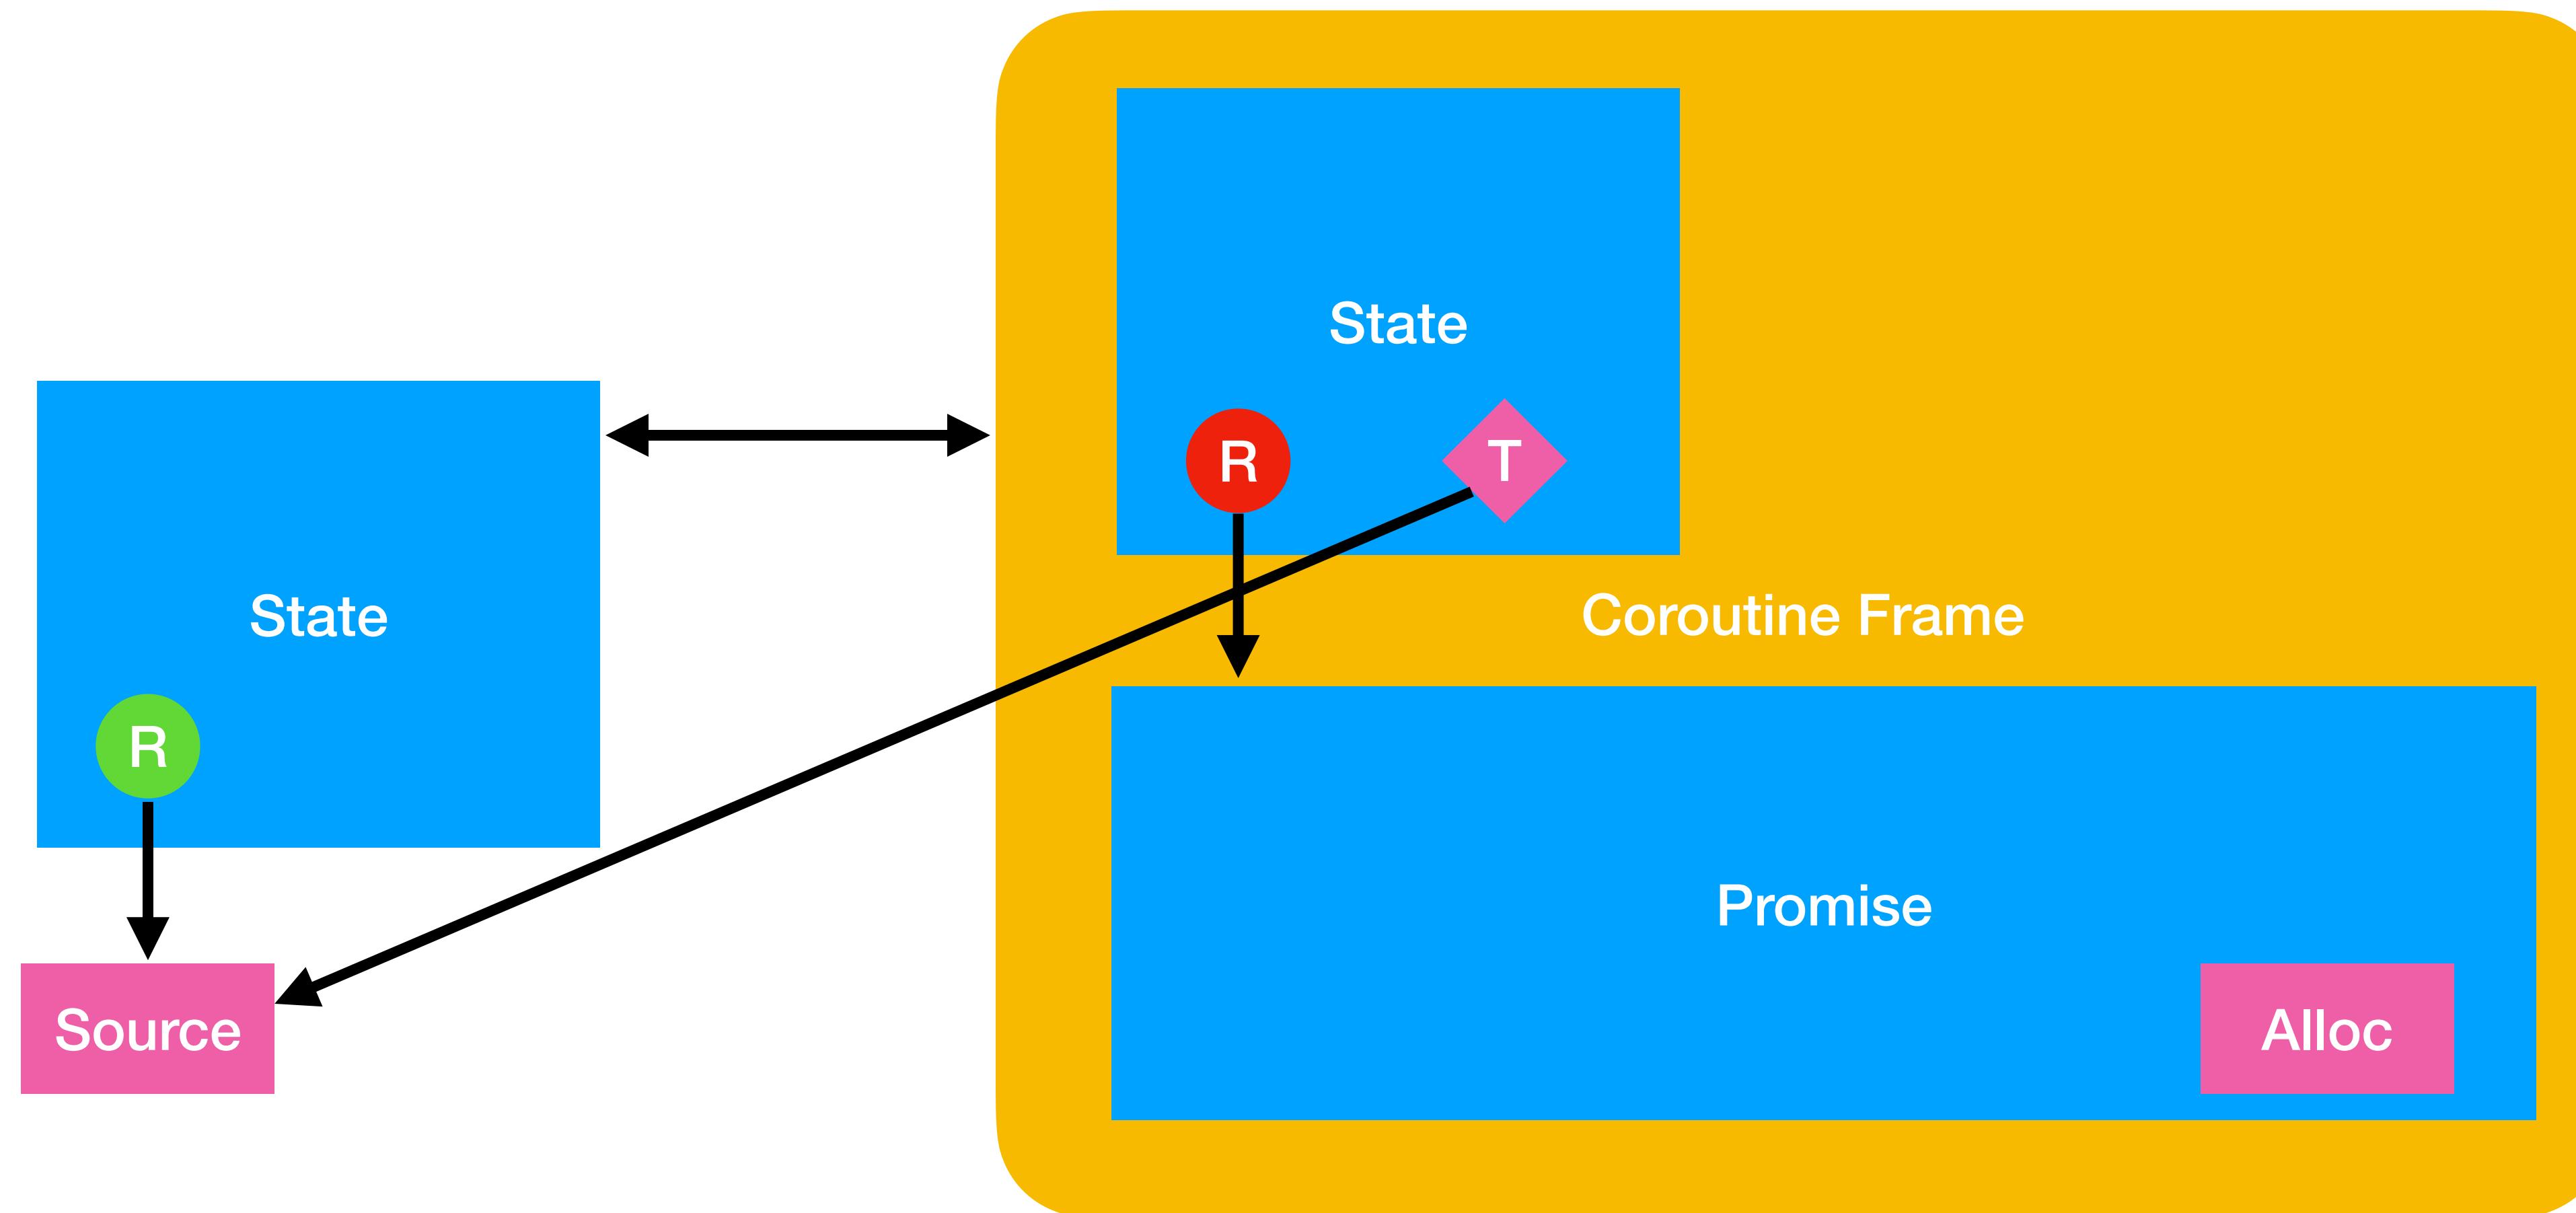
Stop Token From Environment

```
task<> stopping() {
    auto token = co_await read_env(get_stop_token);
    std::size_t count{};
    while (not token.stop_requested()) {
        ++count;
    }
}
```

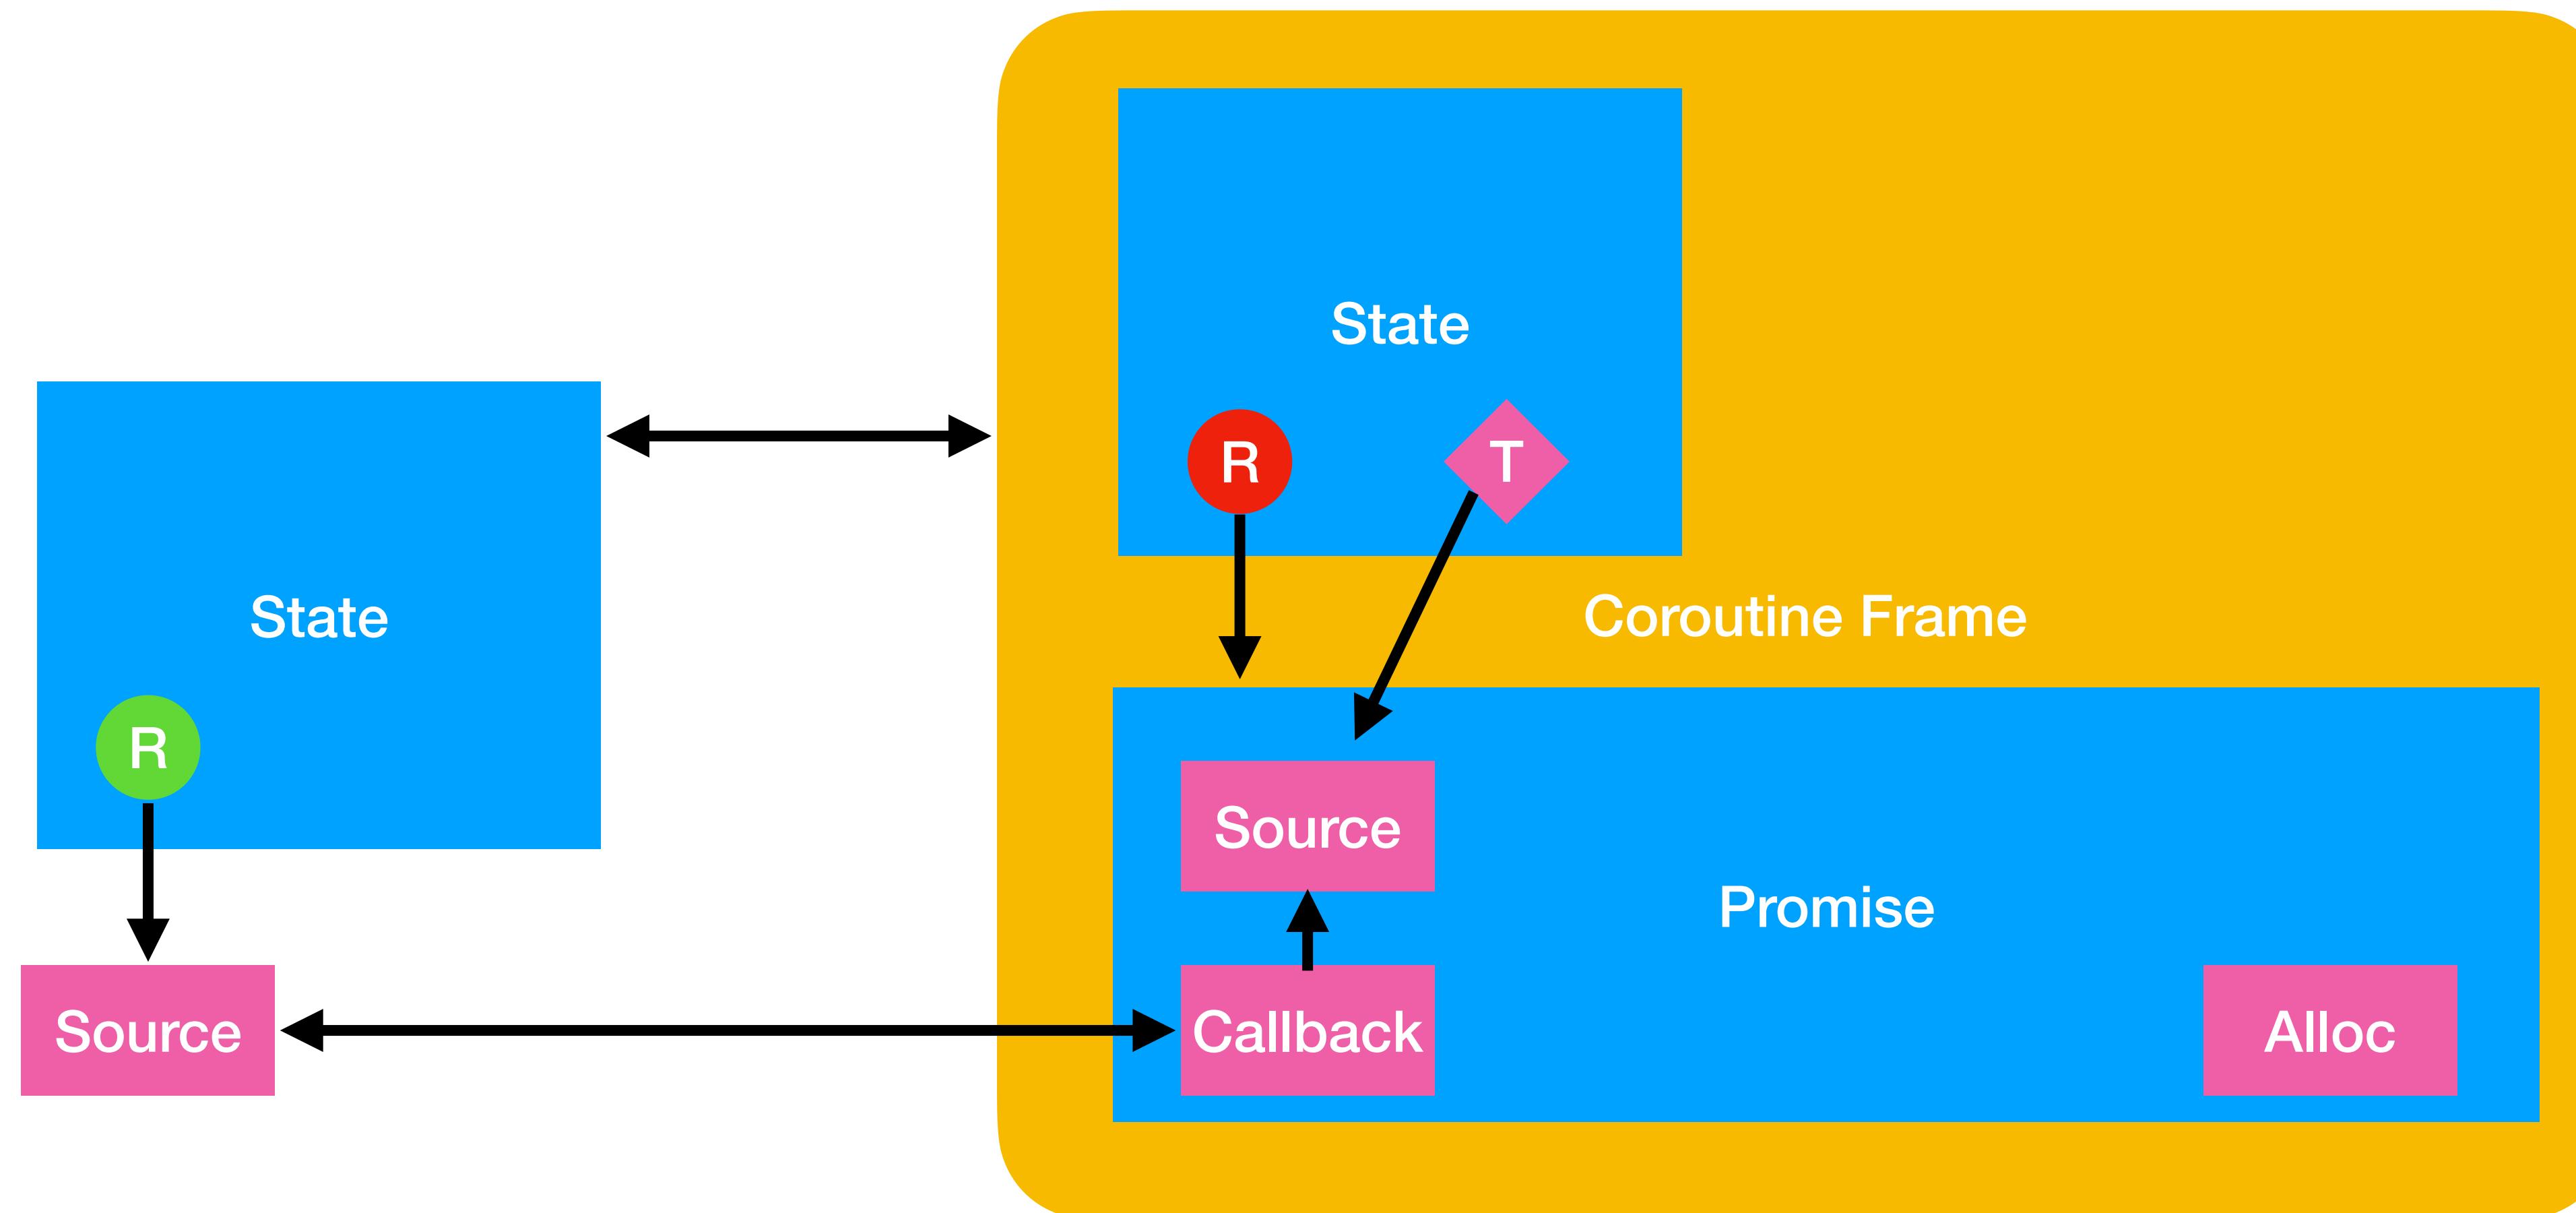
Task Forwards the Stop Token



Task Forwards the Stop Token



Task Forwards the Stop Token



Stop Token From Environment

```
struct stop_token_context {  
    using stop_source_type = std::stop_source;  
};  
task<void, stop_token_context> stopping() {  
    auto token = co_await read_env(get_stop_token);  
    std::size_t count{};  
    while (not token.stop_requested()) {  
        ++count;  
    }  
}
```

Stop Token Support

- Cancellation is very important; thus, stop token support is important
- The stop token type comes from the receiver's environment:
 - `task<...>` type-erases the stop token type
 - Doing so requires de-/registering callbacks
- Use the receiver's stop token if the types match

User-Defined Query

```
constexpr struct get_v_t {
    template <typename E>
        requires requires(const get_v_t& self, const E& e) {
            e.query(self);
        }
    decltype(auto) operator()(const E& e) const {return e.query(*this);}
    constexpr bool query(const forwarding_query_t&)const noexcept
        { return true; }
} get_v{};
```

User-Defined Context

```
task<void> with_env() {
    decltype(auto) v = co_await read_env(get_v); // Error!
}
int main() { ex::sync_wait(with_env()); }
```

User-Defined Context

```
struct context {  
    int value{};  
  
    int query(get_v_t const&) const { return value; }  
};  
  
task<void, context> with_env() {  
    decltype(auto) v = co_await read_env(get_v);  
}  
int main() { ex::sync_wait(with_env()); }
```

User-Defined Context

```
struct context {  
    int value{};  
    context(auto const& e): value(get_v(e)) {} // Error!  
    int query(get_v_t const&) const { return value; }  
};
```

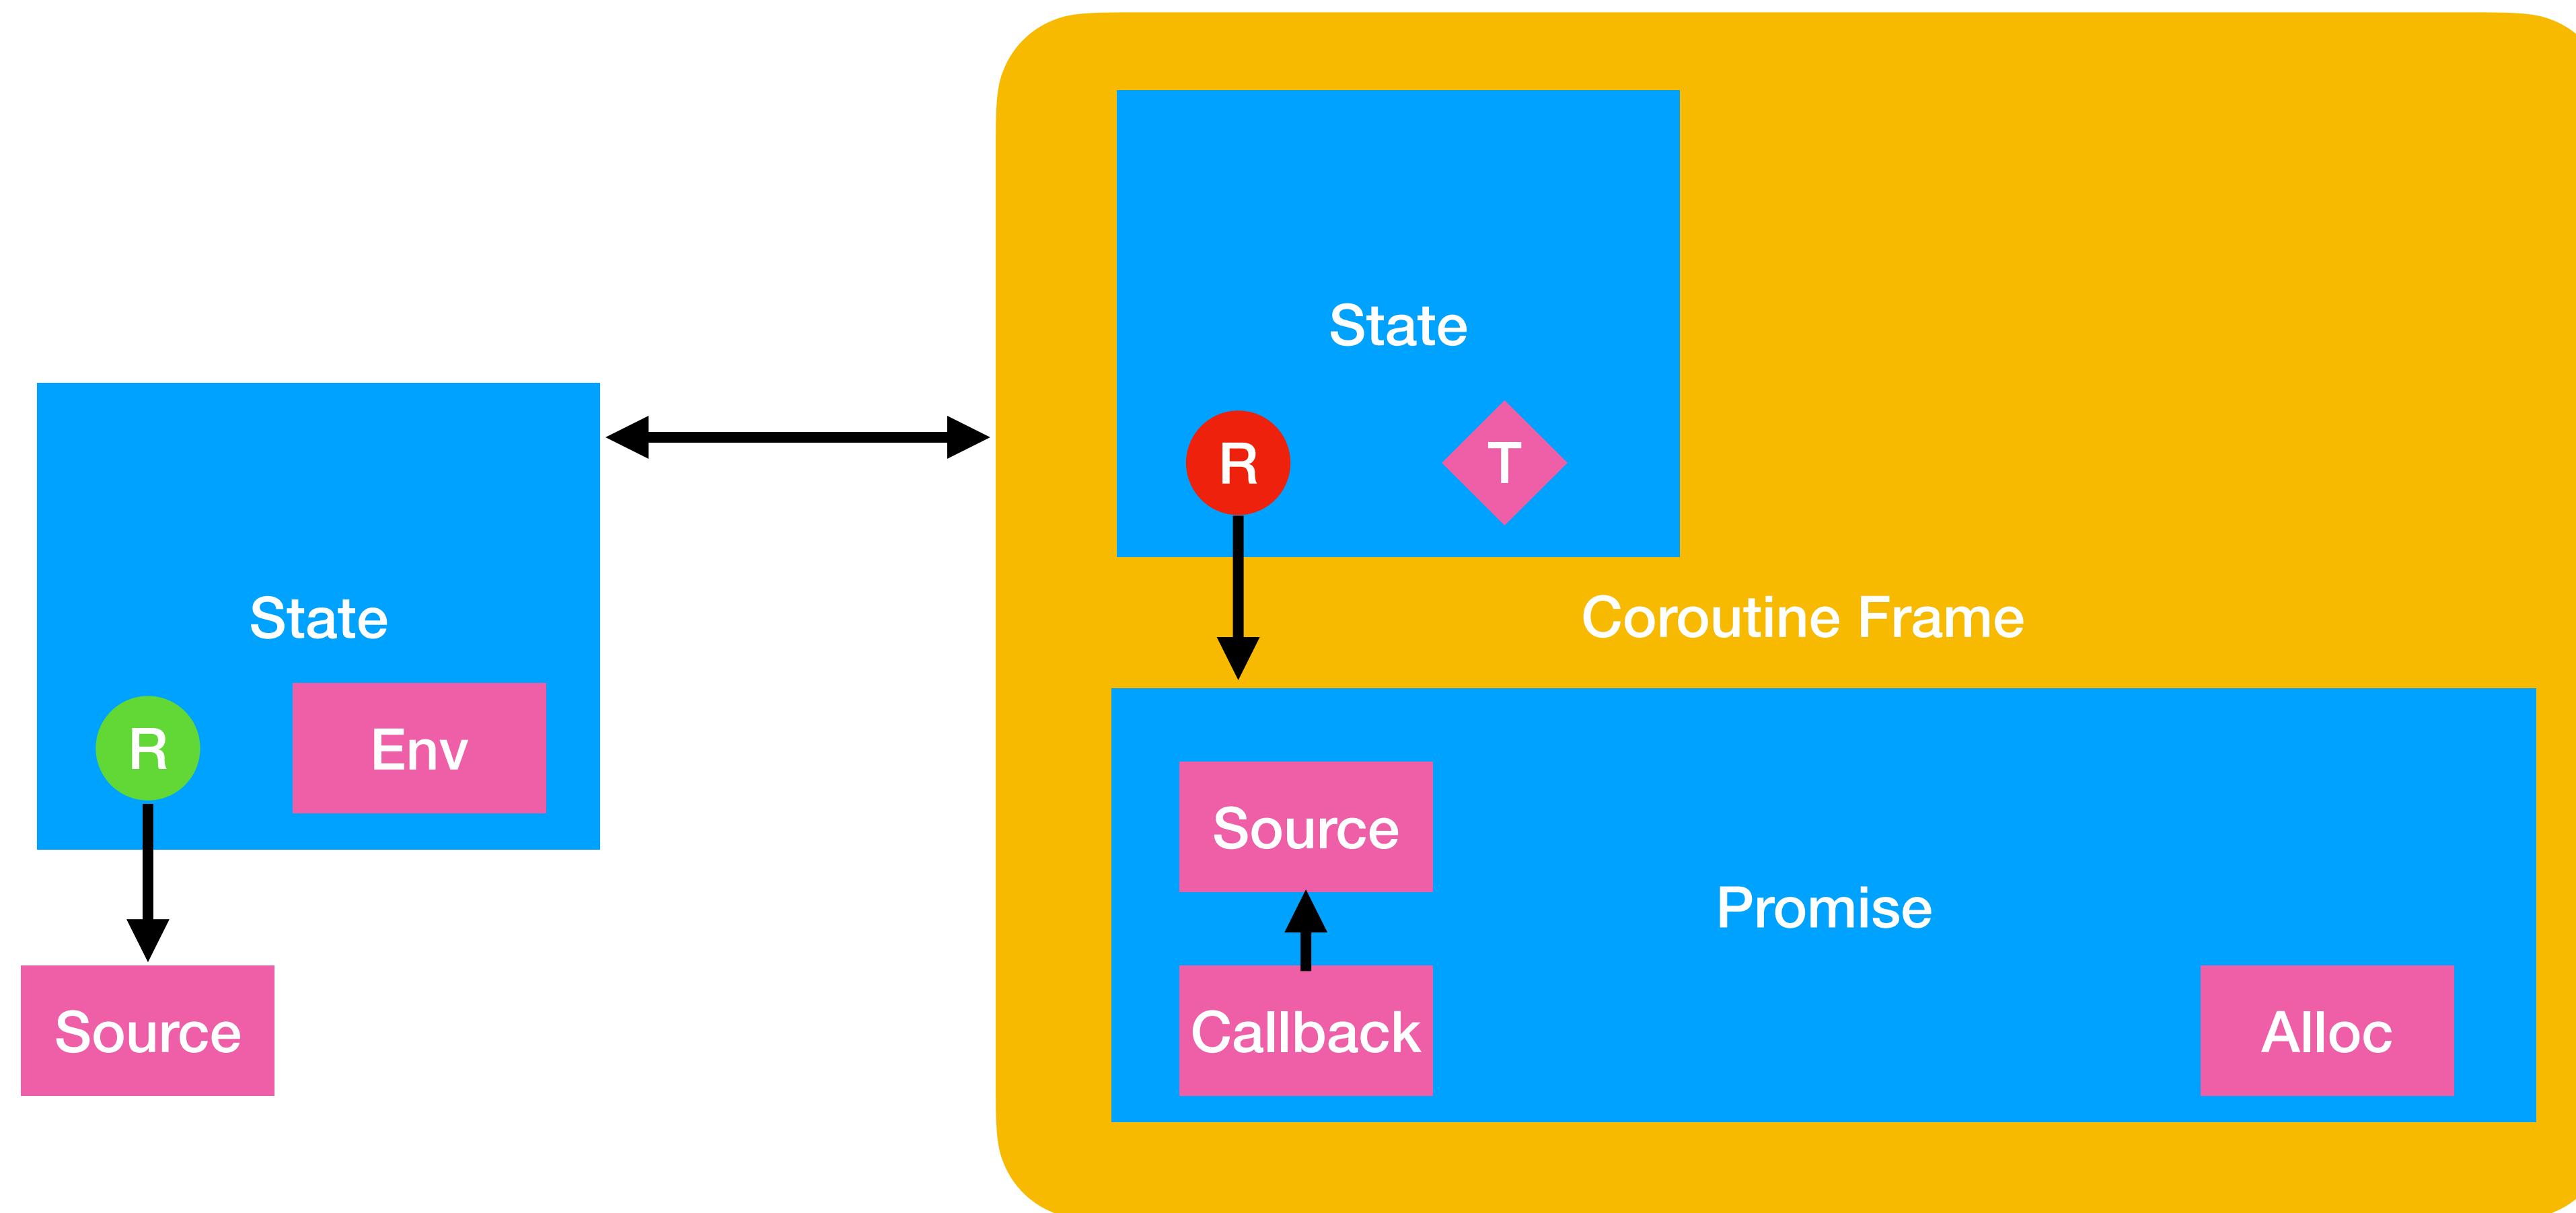
```
task<void, context> with_env() {  
    decltype(auto) v = co_await read_env(get_v);  
}  
int main() { ex::sync_wait(with_env()); }
```

User-Defined Context

```
struct context {  
    int value{};  
    context(auto const& e): value(get_v(e)) {}  
    int query(get_v_t const&) const { return value; }  
};
```

```
task<void, context> with_env() {  
    decltype(auto) v = co_await read_env(get_v);  
}  
int main() { sync_wait(write_env(with_env(), make_env(get_v, 17))); }
```

Task Holds an Environment



User-Defined Context

```
struct fancy {  
    struct env_base { virtual int get() const = 0; };  
    template <typename Env> struct env_type : env_base {  
        Env env;  
        env_type(const Env& e) : env(e) {}  
        int get() const override { return get_v(env); }  
    };  
    fancy(const env_base& own) : env(own) {} env_base const& env;  
    int query(const get_v_t&) const { return this->env.get(); }  
};
```

Where Does a Task Resume?

```
task<> work(auto sched) {
    // executing on orig_sched
    bool before = orig_sched == co_await read_env(get_scheduler);

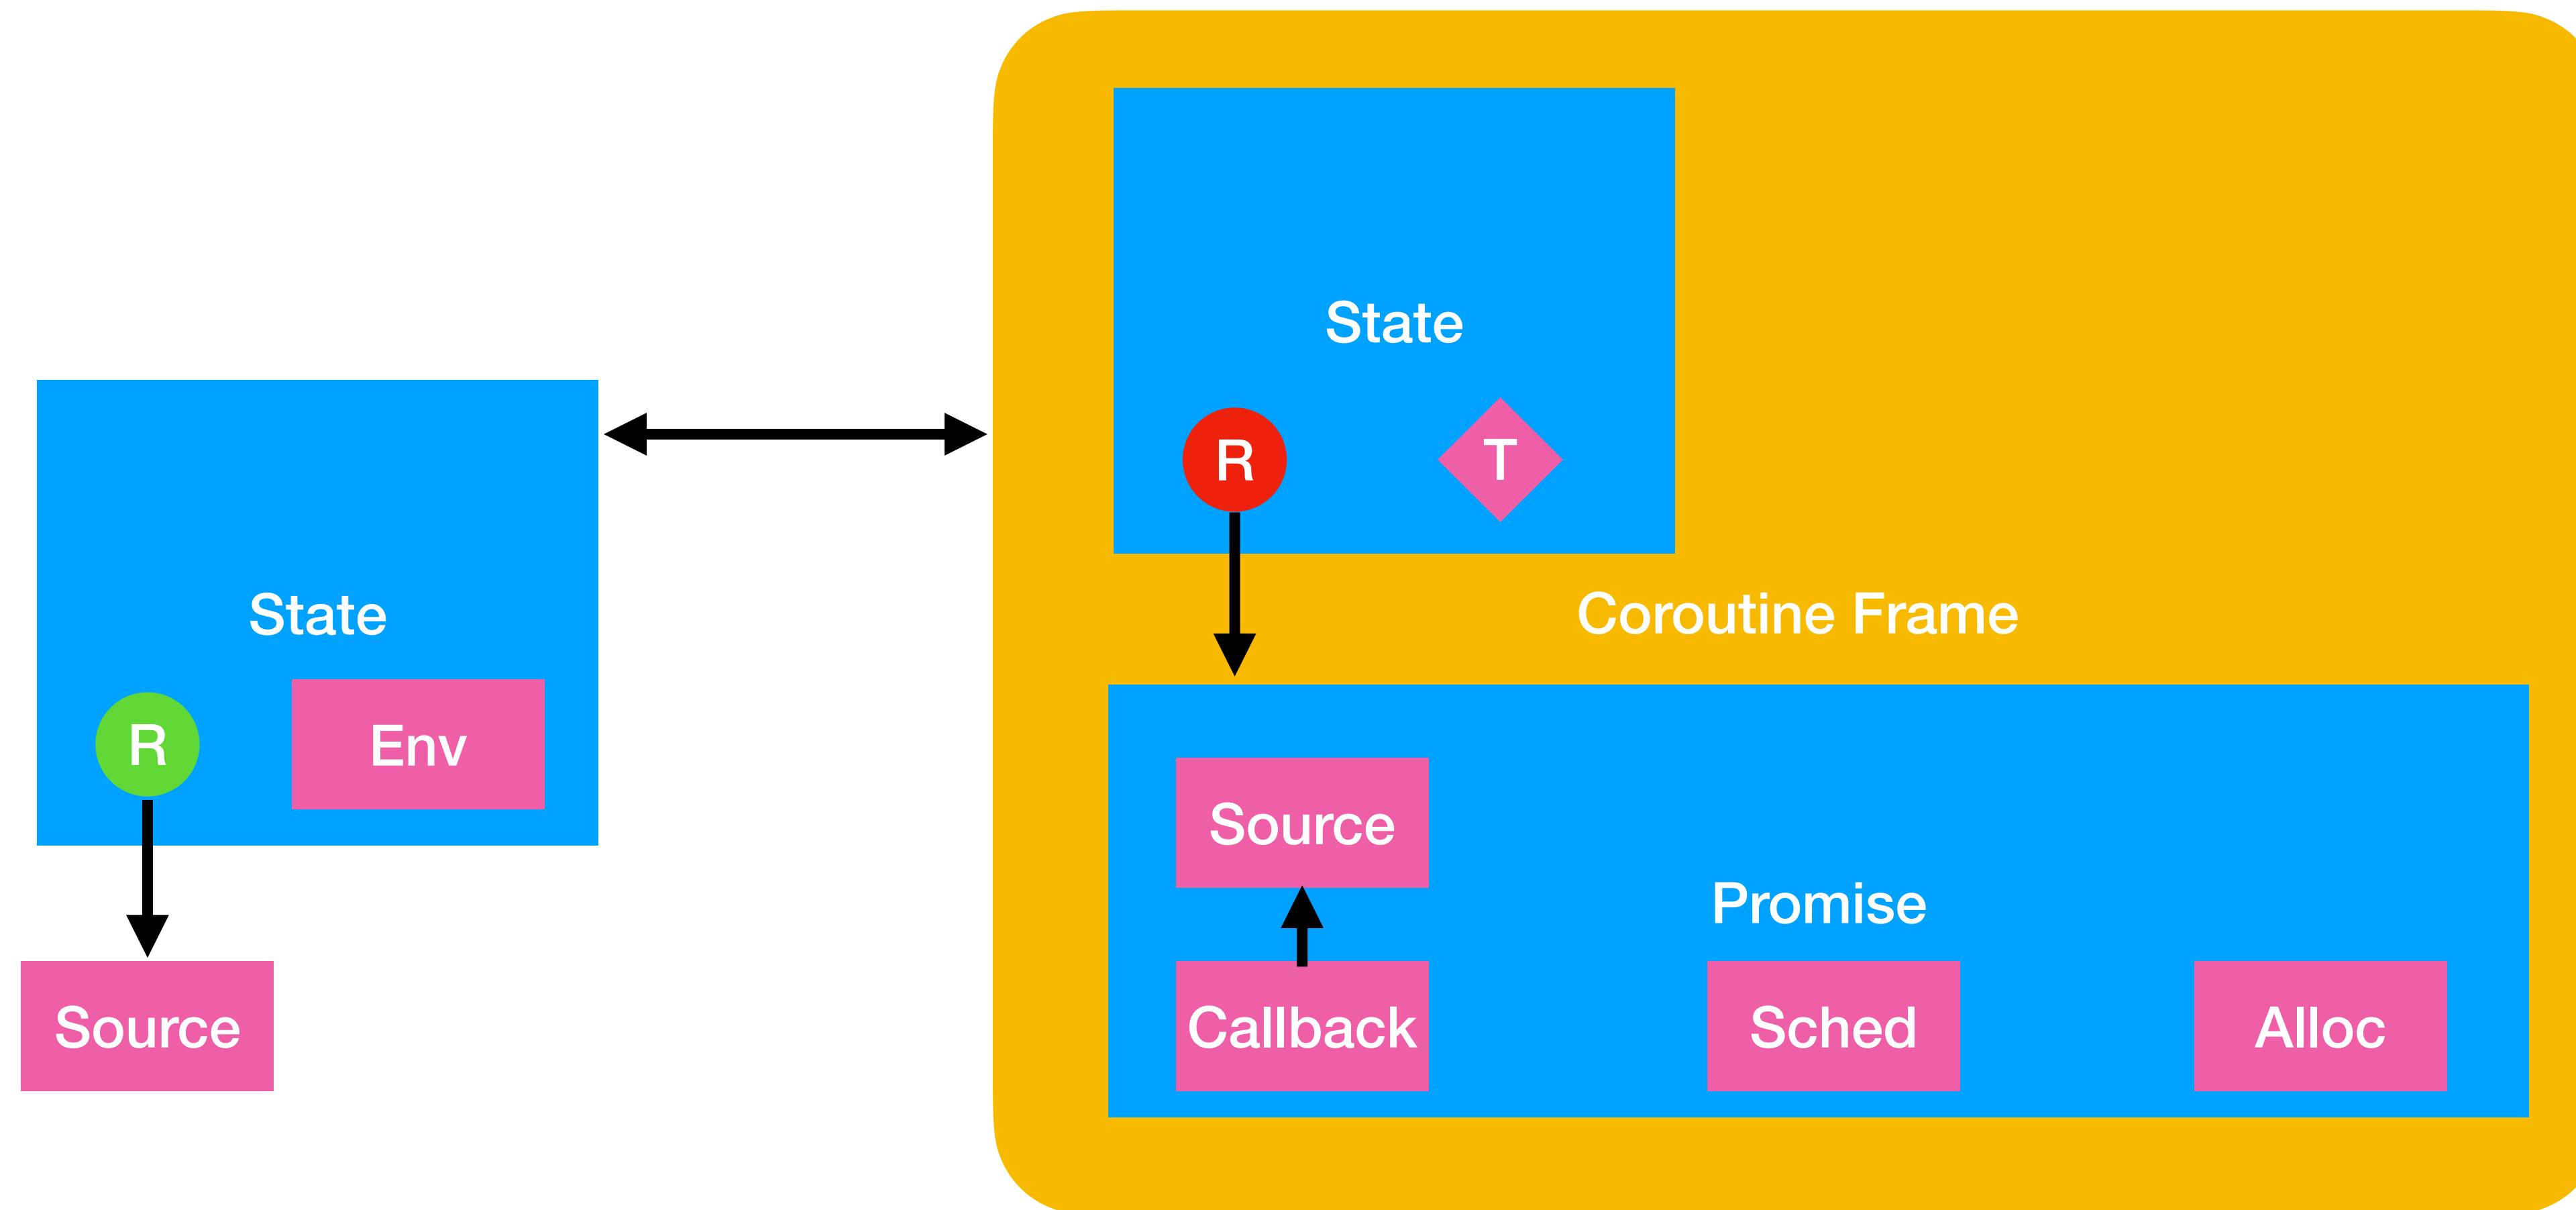
    co_await (schedule(sched) | then([]{ /* executing on sched */ }));
    // executing on orig_sched or on sched?
}
```

Scheduler Affinity

```
task<> work(auto sched) {
    // executing on orig_sched
    bool before = orig_sched == co_await read_env(get_scheduler);

    co_await (schedule(sched) | then([]{ /* executing on sched */ }));
    // executing on orig_sched or on sched?
}
```

Task Holds a Scheduler



Inhibiting Scheduler Affinity

```
task<> work(auto sched) {
    // executing on orig_sched
    bool before = orig_sched == co_await read_env(get_scheduler);
    co_await changeCoroutineScheduler(inline_scheduler{});

    co_await (schedule(sched) | then([]{ /* executing on sched */ }));
    // executing on orig_sched or on sched
}
```

Inhibiting Scheduler Affinity

```
struct inline_context { using scheduler_type = inline_scheduler; };
task<void, inline_context> work(auto sched) {
    // executing on orig_sched
    bool b = inline_scheduler == co_await read_env(get_scheduler);

    co_await (schedule(sched) | then([]{ /* executing on sched */ }));
    // executing on orig_sched or on sched
}
```

Resources

- Implementation: <https://github.com/bemanproject/task>
- Proposal: <https://wg21.link/P3552>



Thank you!