

# Parallel Range Algorithms

*The Evolution of Parallelism in C++*

Ruslan Arutyunyan

# About myself

- Working for Intel

# About myself

- Working for Intel
- oneAPI Data Parallel C++ (oneDPL) lead developer

# About myself

- Working for Intel
- oneAPI Data Parallel C++ (oneDPL) lead developer
- Significant contributions to other threading engines including oneAPI Threading Building Blocks (oneTBB)

# About myself

- Working for Intel
- oneAPI Data Parallel C++ (oneDPL) lead developer
- Significant contributions to other threading engines including oneAPI Threading Building Blocks (oneTBB)
- Contributor to SYCL (in the past)

# About myself

- Working for Intel
- oneAPI Data Parallel C++ (oneDPL) lead developer
- Significant contributions to other threading engines including oneAPI Threading Building Blocks (oneTBB)
- Contributor to SYCL (in the past)
- Contributions to C++ standard including `std::simd`, `std::execution`, and more

# About myself

- Working for Intel
- oneAPI Data Parallel C++ (oneDPL) lead developer
- Significant contributions to other threading engines including oneAPI Threading Building Blocks (oneTBB)
- Contributor to SYCL (in the past)
- Contributions to C++ standard including `std::simd`, `std::execution`, and more
- SG1: Concurrency and Parallelism co-chair in the C++ committee

# About my

- Working for
- oneAPI Data
- Significant
- oneAPI Thr
- Contributor
- Contribution
- `std::exe`
- SGI: Conc



lead developer

leading engines including  
(oneTBB)

ing `std::simd`,

o-chair



# About my

- Working for
- oneAPI Data
- Significant
- oneAPI Th
- Contributor
- Contribution
- `std::exe`
- SGI: Conc













# Parallel Range Algorithms

# Parallel Algorithms (C++17)

## Algorithms:

- With the first `ExecutionPolicy` template parameter
- In `std` namespace
- Taking iterators

# Parallel Algorithms (C++17)

## Algorithms:

- With the first **ExecutionPolicy** template parameter
- In **std** namespace
- Taking iterators

```
// serial
```

```
auto res = std::find_if(std::begin(input), std::end(input), pred);
```

```
// parallel
```

```
auto res = std::find_if(std::execution::par, std::begin(input), std::end(input), pred);
```

# Parallel Range algorithms (P3179 proposal)

## Algorithms:

- With the first template parameter constrained by the *execution-policy* concept
- In **ranges** namespace
- Taking
  - ranges
  - iterators and sentinels

# Parallel Range algorithms (P3179 proposal)

## Algorithms:

- With the first template parameter constrained by the *execution-policy* concept
- In **ranges** namespace
- Taking
  - ranges
  - iterators and sentinels

```
// serial  
auto res = std::ranges::find_if(input, pred);
```

```
// parallel with P3179  
auto res = std::ranges::find_if(std::execution::par, input, pred);
```



# Motivation

Combining the powerful ranges API with parallelism:

- Opportunity to fuse several parallel algorithm invocations into one
- Better expressiveness and productivity for parallel code
- Ease of use

# Example with C++17 Parallel Algorithms

```
std::transform(policy, std::begin(data), std::end(data), std::begin(result),  
              [](auto i){ return i + 1; });  
std::reverse(policy, std::begin(result), std::end(result));  
auto res = std::find_if(policy, std::begin(result), std::end(result), pred);
```

# Example with C++17 Parallel Algorithms

```
std::transform(policy, std::begin(data), std::end(data), std::begin(result),  
              [](auto i){ return i + 1; });  
std::reverse(policy, std::begin(result), std::end(result));  
auto res = std::find_if(policy, std::begin(result), std::end(result), pred);
```

- Three algorithm invocations, each invocation adds its own overhead
- The unnecessary work might be skipped only for the third algorithm call

# Example with fancy iterators and Parallel Algorithms

```
auto res = std::find_if(policy,
    std::make_reverse_iterator(
        dpl::make_transform_iterator(std::end(data), [](auto i){ return i + 1; })),
    std::make_reverse_iterator(
        dpl::make_transform_iterator(std::begin(data), [](auto i){ return i + 1; })),
    pred);
```

# Example with fancy iterators and Parallel Algorithms

```
auto res = std::find_if(policy,
    std::make_reverse_iterator(
        dpl::make_transform_iterator(std::end(data), [](auto i){ return i + 1; })),
    std::make_reverse_iterator(
        dpl::make_transform_iterator(std::begin(data), [](auto i){ return i + 1; })),
    pred);
```

- One algorithm invocation, less overhead for parallelism
- The unnecessary work might be skipped

# Example with fancy iterators and Parallel Algorithms

```
auto res = std::find_if(policy,
    std::make_reverse_iterator(
        dpl::make_transform_iterator(std::end(data), [](auto i){ return i + 1; })),
    std::make_reverse_iterator(
        dpl::make_transform_iterator(std::begin(data), [](auto i){ return i + 1; })),
    pred);
```

- One algorithm invocation, less overhead for parallelism
- The unnecessary work might be skipped

But:

- A lot of verbosity
- `end(data)` is a `reverse_iterator` begin

# Example with fancy iterators and Parallel Algorithms

```
auto res = std::find_if(policy,
    std::make_reverse_iterator(
        dpl::make_transform_iterator(std::end(data), [](auto i){ return i + 1; })),
    std::make_reverse_iterator(
        dpl::make_transform_iterator(std::begin(data), [](auto i){ return i + 1; })),
    pred);
```

- One algorithm invocation, less overhead for parallelism
- The unnecessary work might be skipped

But:

- A lot of verbosity
- `end(data)` is a `reverse_iterator` begin
- The code does not compile as is

# Example with fancy iterators and Parallel Algorithms

```
template<class ExecutionPolicy, class ForwardIt, class UnaryPred>  
ForwardIt find_if(ExecutionPolicy&& policy, ForwardIt first, ForwardIt last, UnaryPred p);
```



# Example with fancy iterators and Parallel Algorithms

```
template<class ExecutionPolicy, class ForwardIt, class UnaryPred>
ForwardIt find_if(ExecutionPolicy&& policy, ForwardIt first, ForwardIt last, UnaryPred p);

template<class It, class Func>
class dpl::transform_iterator;
```

# Example with fancy iterators and Parallel Algorithms

```
template<class ExecutionPolicy, class ForwardIt, class UnaryPred>
ForwardIt find_if(ExecutionPolicy&& policy, ForwardIt first, ForwardIt last, UnaryPred p);

template<class It, class Func>
class dpl::transform_iterator;

auto begin = dpl::make_transform_iterator(std::begin(data), [](auto x) { return x + 1; });
auto end = dpl::make_transform_iterator(std::end(data), [](auto x) { return x + 1; });
```

# Example with fancy iterators and Parallel Algorithms

```
template<class ExecutionPolicy, class ForwardIt, class UnaryPred>  
ForwardIt find_if(ExecutionPolicy&& policy, ForwardIt first, ForwardIt last, UnaryPred p);
```

```
template<class It, class Func>  
class dpl::transform_iterator;
```

```
auto begin = dpl::make_transform_iterator(std::begin(data), [](auto x) { return x + 1; });  
auto end = dpl::make_transform_iterator(std::end(data), [](auto x) { return x + 1; });
```

```
std::find_if(policy, begin, end, pred);    // compile-time error
```

# Example with fancy iterators and Parallel Algorithms

```
template<class ExecutionPolicy, class ForwardIt, class UnaryPred>
ForwardIt find_if(ExecutionPolicy&& policy, ForwardIt first, ForwardIt last, UnaryPred p);

template<class It, class Func>
class dpl::transform_iterator;

auto begin = dpl::make_transform_iterator(std::begin(data), [](auto x) { return x + 1; });
auto end = dpl::make_transform_iterator(std::end(data), [](auto x) { return x + 1; });

// begin type: dpl::transform_iterator<base, lambda1>
// end type: dpl::transform_iterator<base, lambda2>

std::find_if(policy, begin, end, pred);    // compile-time error
```

# Example with fancy iterators and Parallel Algorithms

```
auto res = std::find_if(policy,
    std::make_reverse_iterator(
        dpl::make_transform_iterator(std::end(data), [](auto i){ return i + 1; })),
    std::make_reverse_iterator(
        dpl::make_transform_iterator(std::begin(data), [](auto i){ return i + 1; })),
    pred);
```

- One algorithm invocation, less overhead for parallelism
- The unnecessary work might be skipped

But:

- A lot of verbosity
- `end(data)` is a `reverse_iterator` begin
- The code does not compile as is

# Example with fancy iterators and Parallel Algorithms

```
auto add_one = [](auto i){ return i + 1; };  
auto res = std::find_if(policy,  
    std::make_reverse_iterator(  
        dpl::make_transform_iterator(std::end(data), add_one)),  
    std::make_reverse_iterator(  
        dpl::make_transform_iterator(std::begin(data), add_one)),  
    pred);
```

- One algorithm invocation, less overhead for parallelism
- The unnecessary work might be skipped

But:

- A lot of verbosity
- `end(data)` is a `reverse_iterator` begin

# Example with C++17 parallel algorithms and ranges

```
auto pipeline = data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse;  
auto res = std::find_if(policy, std::ranges::begin(pipeline), std::ranges::end(pipeline),  
                        pred);
```

- One algorithm invocation, less overhead for parallelism
- The unnecessary work might be skipped

# Example with C++17 parallel algorithms and ranges

```
auto pipeline = data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse;  
auto res = std::find_if(policy, std::ranges::begin(pipeline), std::ranges::end(pipeline),  
                        pred);
```

- One algorithm invocation, less overhead for parallelism
- The unnecessary work might be skipped

But:

- Still unnecessary verbosity



# Example with P3179

```
auto res = std::ranges::find_if(policy,  
    data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse,  
    pred);
```

# Example with P3179

```
auto res = std::ranges::find_if(policy,  
    data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse,  
    pred);
```

- One algorithm invocation, less overhead for parallelism
- The unnecessary work might be skipped

# Example with P3179

```
auto res = std::ranges::find_if(policy,  
    data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse,  
    pred);
```

```
*res;    // compile-time error
```

- One algorithm invocation, less overhead for parallelism
- The unnecessary work might be skipped

But:

- `res` is unusable since `transform_view` is not a `borrowed_range`

# Example with P3179

```
auto pipeline = data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse;  
auto res = std::ranges::find_if(policy, pipeline, pred);
```

# Example with P3179

```
auto pipeline = data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse;  
auto res = std::ranges::find_if(policy, pipeline, pred);
```

- One algorithm invocation, less overhead for parallelism
- The unnecessary work might be skipped
- Concise

# Example with P3179

```
auto pipeline = data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse;  
auto res = std::ranges::find_if(policy, pipeline, pred, proj);
```

- One algorithm invocation, less overhead for parallelism
- The unnecessary work might be skipped
- Concise
- Ability to use projections

# Key differences to existing algorithms

# Key differences to existing algorithms

Comparing to the existing algorithms, we propose the following modifications:

- a) The execution policy parameter is added



# Key differences to existing algorithms

Comparing to the existing algorithms, we propose the following modifications:

- a) The execution policy parameter is added
- b) Parallel algorithms require `random_access_{iterator, range}`

# Key differences to existing algorithms

Comparing to the existing algorithms, we propose the following modifications:

- a) The execution policy parameter is added
- b) Parallel algorithms require `random_access_{iterator, range}`
- c) Parallel algorithms require sized ranges

# Key differences to existing algorithms

Comparing to the existing algorithms, we propose the following modifications:

- a) The execution policy parameter is added
- b) Parallel algorithms require `random_access_{iterator, range}`
- c) Parallel algorithms require sized ranges
- d) Parallel range algorithms take a range, not an iterator, as the output for the overloads with ranges, and additionally take an output sentinel for the "iterator and sentinel" overloads

# Starting from the serial signature

```
template<std::input_iterator I, std::sentinel_for<I> S, std::weakly_increamentable O,  
        std::copy_constructible F, class Proj = std::identity>  
    requires std::indirectly_writable<O, std::indirect_result_t<F&, std::projected<I, Proj>>>  
unary_transform_result<I, O>  
    transform(I first1, S last1, O result, F op, Proj proj = {});
```

```
template<ranges::input_range R, std::weakly_increamentable O,  
        std::copy_constructible F, class Proj = std::identity>  
    requires std::indirectly_writable<O,  
        std::indirect_result_t<F&, std::projected<ranges::iterator_t<R>, Proj>>>  
unary_transform_result<ranges::borrowed_iterator_t<R>, O>  
    transform(R&& r, O result, F op, Proj proj = {});
```

## a) Adding an execution policy parameter

```
template<execution-policy Ep, std::input_iterator I, std::sentinel_for<I> S,  
        std::weakly_incremtable O, std::copy_constructible F, class Proj = std::identity>  
    requires std::indirectly_writable<O, std::indirect_result_t<F&, std::projected<I, Proj>>>  
    unary_transform_result<I, O>  
    transform(Ep&& exec, I first1, S last1, O result, F op, Proj proj = {});
```

```
template<execution-policy Ep, ranges::input_range R, std::weakly_incremtable O,  
        std::copy_constructible F, class Proj = std::identity>  
    requires std::indirectly_writable<O,  
        std::indirect_result_t<F&, std::projected<ranges::iterator_t<R>, Proj>>>  
    unary_transform_result<ranges::borrowed_iterator_t<R>, O>  
    transform(Ep&& exec, R&& r, O result, F op, Proj proj = {});
```

```
template<class Ep>  
concept execution-policy = // exposition only  
    std::is_execution_policy_v<std::remove_cvref_t<Ep>>;
```

## b) Requiring random access

```
template<execution-policy Ep, std::random_access_iterator I, std::sentinel_for<I> S,  
        std::random_access_iterator O, std::copy_constructible F, class Proj = std::identity>  
    requires std::indirectly_writable<O, std::indirect_result_t<F&, std::projected<I, Proj>>>  
    unary_transform_result<I, O>  
    transform(Ep&& exec, I first1, S last1, O result, F op, Proj proj = {});
```

```
template<execution-policy Ep, ranges::random_access_range R, std::weakly_increamentable O,  
        std::copy_constructible F, class Proj = std::identity>  
    requires std::indirectly_writable<O,  
        std::indirect_result_t<F&, std::projected<ranges::iterator_t<R>, Proj>>>  
    unary_transform_result<ranges::borrowed_iterator_t<R>, O>  
    transform(Ep&& exec, R&& r, O result, F op, Proj proj = {});
```

## c) Requiring sized range

```
template<execution-policy Ep, std::random_access_iterator I, std::sized_sentinel_for<I> S,  
        std::random_access_iterator O, std::copy_constructible F, class Proj = std::identity>  
    requires std::indirectly_writable<O, std::indirect_result_t<F&, std::projected<I, Proj>>>  
    unary_transform_result<I, O>  
    transform(Ep&& exec, I first1, S last1, O result, F op, Proj proj = {});
```

```
template<execution-policy Ep, sized-random-access-range R, std::weakly_increamentable O,  
        std::copy_constructible F, class Proj = std::identity>  
    requires std::indirectly_writable<O,  
        std::indirect_result_t<F&, std::projected<ranges::iterator_t<R>, Proj>>>  
    unary_transform_result<ranges::borrowed_iterator_t<R>, O>  
    transform(Ep&& exec, R&& r, O result, F op, Proj proj = {});
```

```
template<class R>  
concept sized-random-access-range = // exposition only  
    ranges::random_access_range<R> && ranges::sized_range<R>;
```

## d) Using a range for the output

```
template<execution-policy Ep, std::random_access_iterator I, std::sized_sentinel_for<I> S,  
        std::random_access_iterator O, std::sized_sentinel_for<O> OutS,  
        std::copy_constructible F, class Proj = std::identity>  
    requires std::indirectly_writable<O, std::indirect_result_t<F&, std::projected<I, Proj>>>  
    unary_transform_result<I, O>  
    transform(Ep&& exec, I first1, S last1, O result, OutS result_last, F op, Proj proj = {});
```

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR,  
        std::copy_constructible F, class Proj = std::identity>  
    requires std::indirectly_writable<ranges::iterator_t<OutR>,  
        std::indirect_result_t<F&, std::projected<ranges::iterator_t<R>, Proj>>>  
    unary_transform_result<ranges::borrowed_iterator_t<R>, ranges::borrowed_iterator_t<OutR>>  
    transform(Ep&& exec, R&& r, OutR&& result_r, F op, Proj proj = {});
```



# Comparing side-by-side (Iterator and Sentinel overload)

```
// serial
template<std::input_iterator I, std::sentinel_for<I> S, std::weakly_increamentable O,
        std::copy_constructible F, class Proj = std::identity>
    requires std::indirectly_writable<O, std::indirect_result_t<F&, std::projected<I, Proj>>>
    unary_transform_result<I, O>
    transform(I first1, S last1, O result, F op, Proj proj = {});

// parallel
template<execution_policy Ep, std::random_access_iterator I, std::sized_sentinel_for<I> S,
        std::random_access_iterator O, std::sized_sentinel_for<O> OutS,
        std::copy_constructible F, class Proj = std::identity>
    requires std::indirectly_writable<O, std::indirect_result_t<F&, std::projected<I, Proj>>>
    unary_transform_result<I, O>
    transform(Ep&& exec, I first1, S last1, O result, OutS result_last, F op, Proj proj = {});
```

# Comparing side-by-side (range overload)

```
// serial
template<ranges::input_range R, std::weakly_increamentable O,
        std::copy_constructible F, class Proj = std::identity>
    requires std::indirectly_writable<O,
        std::indirect_result_t<F&, std::projected<ranges::iterator_t<R>, Proj>>>
ranges::unary_transform_result<ranges::borrowed_iterator_t<R>, O>
    ranges::transform(R&& r, O result, F op, Proj proj = {});

// parallel
template<execution_policy Ep, sized-random-access-range R, sized-random-access-range OutR,
        std::copy_constructible F, class Proj = std::identity>
    requires indirectly_writable<ranges::iterator_t<OutR>,
        std::indirect_result_t<F&, std::projected<ranges::iterator_t<R>, Proj>>>
ranges::unary_transform_result<ranges::borrowed_iterator_t<R>, ranges::borrowed_iterator_t<OutR>>
    ranges::transform(Ep&& exec, R&& r, OutR&& result_r, F op, Proj proj = {});
```

# More about design

## *sized-random-access-range*

- C++17 parallel algorithms require *Cpp17ForwardIterator*
  - Intel oneDPL, Nvidia Thrust, GNU libstdc++ implementations are based on random access
  - Only Microsoft STL supports forward iterators\*

## *sized-random-access-range*

- C++17 parallel algorithms require *Cpp17ForwardIterator*
  - Intel oneDPL, Nvidia Thrust, GNU libstdc++ implementations are based on random access
  - Only Microsoft STL supports forward iterators\*
- Random access is the best abstraction in the standard (for now)
  - Some potentially useful views are not supported (e.g., **filter\_view**)
  - Might be relaxed in the future

## *sized-random-access-range*

- C++17 parallel algorithms require *Cpp17ForwardIterator*
  - Intel oneDPL, Nvidia Thrust, GNU libstdc++ implementations are based on random access
  - Only Microsoft STL supports forward iterators\*
- Random access is the best abstraction in the standard (for now)
  - Some potentially useful views are not supported (e.g., **filter\_view**)
  - Might be relaxed in the future
- Size is necessary to know in advance for parallelization
  - memory safety: everything is bounded, including the output
  - performance: no need to do unnecessary work

# Range-as-the-output

**Category** – algorithms with output

# Range-as-the-output

- Pros:
  - Ease of use
  - Better memory safety
  - Potentially better performance
  - Ability to detect an error



# Range-as-the-output

- Pros:
  - Ease of use
  - Better memory safety
  - Potentially better performance
  - Ability to detect an error
- The objection was:
  - More complicated switch between serial and parallel algorithms
  - Inconsistency with serial range algorithms
  - Unclear semantics

# Input evolution example

```
template<class InputIterator1, class InputIterator2,  
        class OutputIterator, class BinaryOperation>  
constexpr OutputIterator  
    transform(InputIterator1 first1, InputIterator1 last1,  
              InputIterator2 first2, OutputIterator result,  
              BinaryOperation binary_op);
```

# Input evolution example

```
template<class InputIterator1, class InputIterator2,
        class OutputIterator, class BinaryOperation>
constexpr OutputIterator
    transform(InputIterator1 first1, InputIterator1 last1,
              InputIterator2 first2, OutputIterator result,
              BinaryOperation binary_op);

template<input_iterator I1, sentinel_for<I1> S1, input_iterator I2, sentinel_for<I2> S2,
        weakly_increamentable O, copy_constructible F, class Proj1 = identity,
        class Proj2 = identity>
    requires indirectly_writable<O, indirect_result_t<F&, projected<I1, Proj1>, projected<I2, Proj2>>>
constexpr ranges::binary_transform_result<I1, I2, O>
    ranges::transform(I1 first1, S1 last1, I2 first2, S2 last2, O result,
                     F binary_op, Proj1 proj1 = {}, Proj2 proj2 = {});

template<input_range R1, input_range R2, weakly_increamentable O,
        copy_constructible F, class Proj1 = identity, class Proj2 = identity>
    requires indirectly_writable<O, indirect_result_t<F&, projected<iterator_t<R1>, Proj1>,
        projected<iterator_t<R2>, Proj2>>>
constexpr ranges::binary_transform_result<borrowed_iterator_t<R1>, borrowed_iterator_t<R2>, O>
    ranges::transform(R1&& r1, R2&& r2, O result, F binary_op, Proj1 proj1 = {}, Proj2 proj2 = {});
```

# Output inconsistency

```
template<input_iterator I, sentinel_for<I> S, weakly_incremtable O>  
    requires indirectly_copyable<I, O>  
constexpr ranges::copy_result<I, O> ranges::copy(I first, S last, O result);  
  
template<input_range R, weakly_incremtable O>  
    requires indirectly_copyable<iterator_t<R>, O>  
constexpr ranges::copy_result<borrowed_iterator_t<R>, O> ranges::copy(R&& r, O result);
```

# Output inconsistency

```
template<input_iterator I, sentinel_for<I> S, weakly_incremtable O>  
    requires indirectly_copyable<I, O>  
constexpr ranges::copy_result<I, O> ranges::copy(I first, S last, O result);
```

```
template<input_range R, weakly_incremtable O>  
    requires indirectly_copyable<iterator_t<R>, O>  
constexpr ranges::copy_result<borrowed_iterator_t<R>, O> ranges::copy(R&& r, O result);
```

```
template<input_iterator I, sentinel_for<I> S1, nothrow-forward-iterator O, nothrow-sentinel-for<O> S2>  
    requires constructible_from<iter_value_t<O>, iter_reference_t<I>>  
constexpr uninitialized_copy_result<I, O>  
    uninitialized_copy(I ifirst, S1 ilast, O ofirst, S2 olast);
```

```
template<input_range IR, nothrow-forward-range OR>  
    requires constructible_from<range_value_t<OR>, range_reference_t<IR>>  
constexpr uninitialized_copy_result<borrowed_iterator_t<IR>, borrowed_iterator_t<OR>>  
    uninitialized_copy(IR&& in_range, OR&& out_range);
```

# Unclear semantics

```
std::vector<int> v1{1,2,3,4,5};
```

```
std::vector<int> v2(3);
```

```
std::ranges::copy(v1, v2); // might appear that copy allocates for v2
```

# Addressing unclear semantics

```
std::vector<int> v1{1,2,3,4,5};  
std::vector<int> v2(3);  
std::ranges::copy(v1, v2); // might appear that copy allocates for v2
```

Our proposal: Execute an algorithm until any of the ranges ends

# Addressing unclear semantics

```
std::vector<int> v1{1,2,3,4,5};  
std::vector<int> v2(3);  
std::ranges::copy(v1, v2); // might appear that copy allocates for v2
```

Our proposal: Execute an algorithm until any of the ranges ends

Algorithm with the same semantics:

- uninitialized\_copy



# Addressing unclear semantics

```
std::vector<int> v1{1,2,3,4,5};  
std::vector<int> v2(3);  
std::ranges::copy(v1, v2); // might appear that copy allocates for v2
```

Our proposal: Execute an algorithm until any of the ranges ends

Algorithms with the same semantics:

- uninitialized\_copy
- uninitialized\_move
- partial\_sort\_copy

# Range-as-the-output

- Pros:
  - Ease of use
  - Better memory safety
  - Potentially better performance
  - Ability to detect an error
- The objection was:
  - More complicated switch between serial and parallel algorithms
  - Inconsistency with serial range algorithms
  - Unclear semantics

# Range-as-the-output

- Pros:
  - Ease of use
  - Better memory safety
  - Potentially better performance
  - Ability to detect an error
- The objection was:
  - More complicated switch between serial and parallel algorithms
  - ~~■ Inconsistency with serial range algorithms~~
  - ~~■ Unclear semantics~~

# Range-as-the-output

- Pros:
  - Ease of use
  - Better memory safety
  - Potentially better performance
  - Ability to detect an error
- The objection was:
  - More complicated switch between serial and parallel algorithms
  - ~~▪ Inconsistency with serial range algorithms~~
  - ~~▪ Unclear semantics~~

[P3490](#) proposal has more detail about range-as-the-output design aspect

# reverse\_copy and rotate\_copy

Category – special algorithms

# reverse\_copy

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR>  
    requires indirectly_copyable<iterator_t<R>, iterator_t<OutR>>  
ranges::reverse_copy_result<borrowed_iterator_t<R>, borrowed_iterator_t<OutR>>  
    reverse_copy(Ep&& exec, R&& r, OutR&& result_r);
```

# reverse\_copy

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR>  
    requires indirectly_copyable<iterator_t<R>, iterator_t<OutR>>  
ranges::reverse_copy_result<borrowed_iterator_t<R>, borrowed_iterator_t<OutR>>  
    reverse_copy(Ep&& exec, R&& r, OutR&& result_r);
```

```
std::vector input{1,2,3,4};  
std::vector<int> output(2);
```

```
// Takes 2 elements from input  
auto res = std::ranges::reverse_copy(policy, input, output);
```

# reverse\_copy

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR>
    requires indirectly_copyable<iterator_t<R>, iterator_t<OutR>>
ranges::reverse_copy_result<borrowed_iterator_t<R>, borrowed_iterator_t<OutR>>
    reverse_copy(Ep&& exec, R&& r, OutR&& result_r);
```

```
std::vector input{1,2,3,4};
std::vector<int> output(2);
```

```
// Takes 2 elements from input
auto res = std::ranges::reverse_copy(policy, input, output);
// After the algorithm execution:
// - output is {4,3}
// - res.in == (input.end() - 2) is true, res.in points to the element equal to 3
// - [ input.begin(), res.in ) gives uncopied range
```



# reverse\_copy

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR>
    requires indirectly_copyable<iterator_t<R>, iterator_t<OutR>>
ranges::reverse_copy_result<borrowed_iterator_t<R>, borrowed_iterator_t<OutR>>
    reverse_copy(Ep&& exec, R&& r, OutR&& result_r);
```

```
std::vector input{1,2,3,4};
std::vector<int> output(4);
```

```
// Takes 2 elements from input
auto res = std::ranges::reverse_copy(policy, input, output);
// After the algorithm execution:
// - output is {4,3,2,1}
// - res.in == (input.end() - 4) is true, res.in points to the element equal to 1
// - [ input.begin(), res.in ) is an empty range
```

# reverse\_copy emulation

```
std::vector input{1,2,3,4};
std::vector<int> output(10);

// Emulates reverse_copy
auto res = std::ranges::copy(input | std::views::reverse | std::views::take(2), output.begin());
// After the algorithm execution:
// - output is {4,3}
// - res.in points to the element equal to 2, however it is a reverse_iterator,
//   thus is not comparable with input.end()
// - res.in.base() has the same type as input.end(), and res.in.base() == input.end() - 2 is true
```

# reverse\_copy emulation for output

```
std::vector input{1,2,3,4};
std::unique_ptr mem = std::make_unique_for_overwrite<int[]>(10);
std::ranges::subrange<int*, int*> mem_range{mem.get(), mem.get() + 10};

auto res = std::ranges::uninitialized_copy(input | std::views::reverse, mem_range | std::views::take(2));
// After the algorithm execution:
// - output is {4,3}
// - res.in points to the element equal to 2, however it is a reverse_iterator,
//   thus is not comparable with input.end()
// - res.in.base() has the same type as input.end(), and res.in.base() == input.end() - 2 is true
```

# rotate\_copy

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR>  
    requires indirectly_copyable<iterator_t<R>, iterator_t<OutR>>  
ranges::rotate_copy_result<borrowed_iterator_t<R>, borrowed_iterator_t<OutR>>  
    ranges::rotate_copy(Ep&& exec, R&& r, iterator_t<R> middle, OutR&& result_r);
```

# rotate\_copy

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR>  
    requires indirectly_copyable<iterator_t<R>, iterator_t<OutR>>  
ranges::rotate_copy_result<borrowed_iterator_t<R>, borrowed_iterator_t<OutR>>  
    ranges::rotate_copy(Ep&& exec, R&& r, iterator_t<R> middle, OutR&& result_r);
```

## Scenarios:

1. `result_r` has space for some elements
2. `result_r` has space for all elements
3. `result_r` has space equal to “tail”
4. `result_r` is empty

# rotate\_copy (case 1a)

```
std::vector input{1,2,3,4,5,6};  
std::vector<int> output(2);  
  
auto res = std::ranges::rotate_copy(policy, input, input.begin() + 3, output);  
// After the algorithm execution:  
// - output is {4,5}  
// - res.in points to the element equal to 6, past the last element written to the output
```

# rotate\_copy (case 1b)

```
std::vector input{1,2,3,4,5,6};  
std::vector<int> output(5);  
  
auto res = std::ranges::rotate_copy(policy, input, input.begin() + 3, output);  
// After the algorithm execution:  
// - output is {4,5,6,1,2}  
// - res.in points to the element equal to 3, past the last element written to the output
```

# rotate\_copy (case 2)

```
std::vector input{1,2,3,4,5,6};  
std::vector<int> output(6);  
  
auto res = std::ranges::rotate_copy(policy, input, input.begin() + 3, output);  
// After the algorithm execution:  
// - output is {4,5,6,1,2,3}  
// - res.in points to the element equal to 4, past the last element written to the output
```



# rotate\_copy (case 3)

```
std::vector input{1,2,3,4,5,6};  
std::vector<int> output(3);  
  
auto res = std::ranges::rotate_copy(policy, input, input.begin() + 3, output);  
// After the algorithm execution:  
// - output is {4,5,6}  
// - res.in points to the element equal to 1, past the last element written to the output
```

# rotate\_copy (case 4)

```
std::vector input{1,2,3,4,5,6};  
std::vector<int> output;  
  
auto res = std::ranges::rotate_copy(policy, input, input.begin() + 3, output);  
// After the algorithm execution:  
// - output is {}  
// - res.in points to the element equal to 4, middle, because no elements were written  
// - res.out == output.begin() is true
```

# rotate\_copy emulation

```
std::vector input{1,2,3,4,5,6};  
std::vector<int> output(6);  
  
// rotate_view emulation  
auto middle_to_last = input | std::views::drop(3);  
auto first_to_middle = input | std::views::take(3);  
auto rotate_view = ::ranges::views::concat(middle_to_last, first_to_middle); // from range-v3
```

# rotate\_copy emulation

```
std::vector input{1,2,3,4,5,6};
std::vector<int> output(6);

// rotate_view emulation
auto middle_to_last = input | std::views::drop(3);
auto first_to_middle = input | std::views::take(3);
auto rotate_view = ::ranges::views::concat(middle_to_last, first_to_middle); // from range-v3

auto res = std::ranges::copy(rotate_view, output.begin());
// After the algorithm execution:
// - output is {4,5,6,1,2,3}
// - res.in points to the element equal to 4, middle, which is past the last element written to the output
```

# rotate\_copy emulation

```
std::vector input{1,2,3,4,5,6};
std::vector<int> output(6);

// rotate_view emulation
auto middle_to_last = input | std::views::drop(3);
auto first_to_middle = input | std::views::take(3);
auto rotate_view = ::ranges::views::concat(middle_to_last, first_to_middle); // from range-v3

auto res = std::ranges::copy(rotate_view | std::views::take(3), output.begin());
// After the algorithm execution:
// - output is {4,5,6}
// - res.in points to the element equal to 1, past the last element written to the output
```

# copy\_if

Category – algorithms with output and gaps

## copy\_if parallel signature

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR,
        class Proj = std::identity,
        std::indirect_unary_predicate<std::projected<std::ranges::iterator_t<R>, Proj>> Pred>
    requires std::indirectly_copyable<std::ranges::iterator_t<R>, std::ranges::iterator_t<OutR>>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
    copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {})
{
}
}
```

# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR,  
        class Proj = std::identity,  
        std::indirect_unary_predicate<std::projected<std::ranges::iterator_t<R>, Proj>> Pred>  
    requires std::indirectly_copyable<std::ranges::iterator_t<R>, std::ranges::iterator_t<OutR>>  
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>  
    copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {})  
{  
    std::size_t size = std::ranges::size(r);  
  
    std::ranges::for_each(exec, std::views::iota(std::size_t(0), size), [=, &r, &result_r](auto i) {  
        if (std::invoke(pred, std::invoke(proj, r[i])))  
            result_r[i] = r[i];  
    });  
  
    return {std::ranges::begin(r) + size, std::ranges::begin(result_r) + size};  
}
```



# copy\_if parallel implementation (wrong)

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR,  
        class Proj = std::identity,  
        std::indirect_unary_predicate<std::projected<std::ranges::iterator_t<R>, Proj>> Pred>  
    requires std::indirectly_copyable<std::ranges::iterator_t<R>, std::ranges::iterator_t<OutR>>  
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>  
    copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {})  
{  
    std::size_t size = std::ranges::size(r);  
  
    std::ranges::for_each(exec, std::views::iota(std::size_t(0), size), [=, &r, &result_r](auto i) {  
        if (std::invoke(pred, std::invoke(proj, r[i])))  
            result_r[i] = r[i];  
    });  
  
    return {std::ranges::begin(r) + size, std::ranges::begin(result_r) + size};  
}  
  
// Input:  {1,7,4,4,4,6,5,2,3,7,9}, pred: [](auto x) { return (x & 0x01) == 1; }  
// Output: {0,0,0,0,0,0,0,0,0,0,0}  
// Result: {1,7,0,0,0,0,5,0,3,7,9}
```

# copy\_if example

```
std::vector data{1,1,2,3,5,8,8,8,4,9,7,2,2,2,2,3};  
std::vector<int> out(data.size());
```

```
auto res = std::ranges::copy_if(std::execution::par, data, out, [](auto x) { return (x & 0x01) == 1; });
```

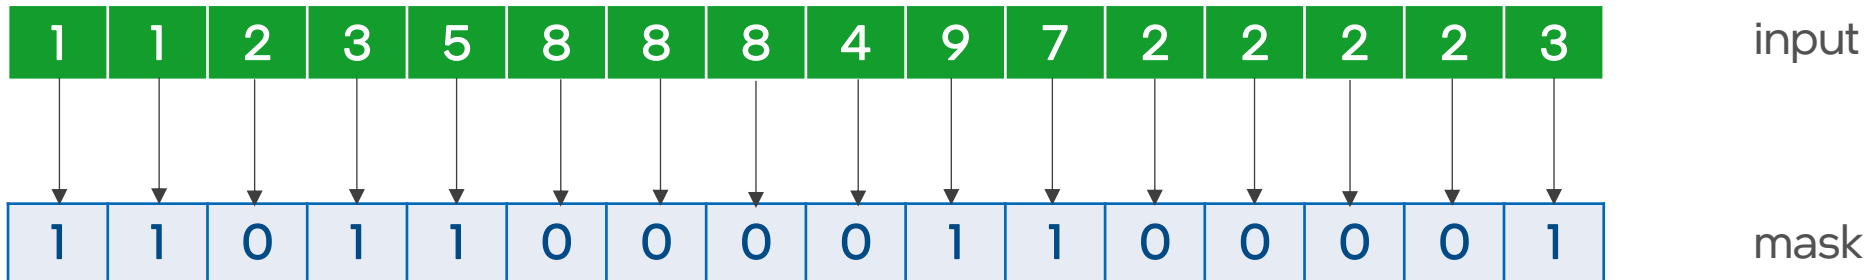
1	1	2	3	5	8	8	8	4	9	7	2	2	2	2	3
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

input

# copy\_if example

```
std::vector data{1,1,2,3,5,8,8,8,4,9,7,2,2,2,2,3};  
std::vector<int> out(data.size());
```

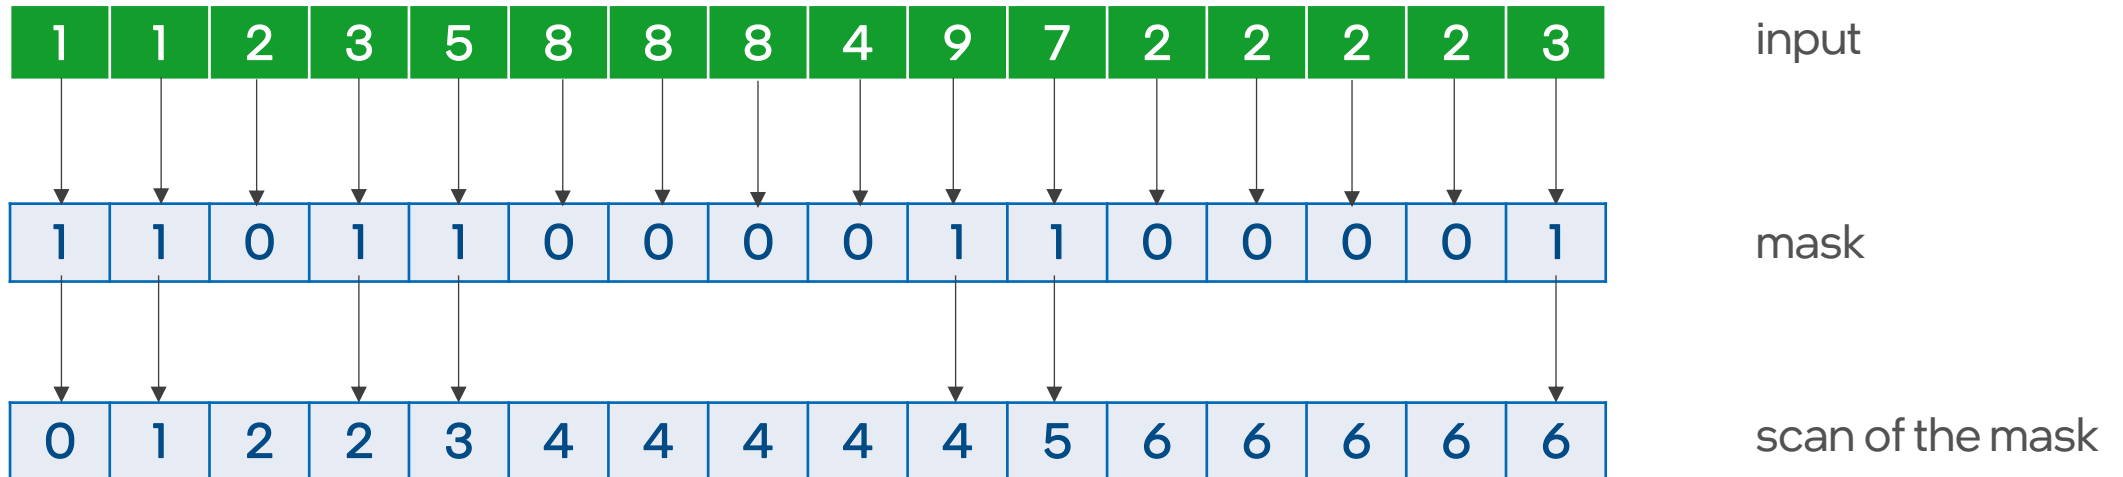
```
auto res = std::ranges::copy_if(std::execution::par, data, out, [](auto x) { return (x & 0x01) == 1; });
```



# copy\_if example

```
std::vector data{1,1,2,3,5,8,8,8,4,9,7,2,2,2,2,3};  
std::vector<int> out(data.size());
```


```
auto res = std::ranges::copy_if(std::execution::par, data, out, [](auto x) { return (x & 0x01) == 1; });
```



# copy\_if example

```
std::vector data{1,1,2,3,5,8,8,8,4,9,7,2,2,2,2,3};  
std::vector<int> out(data.size());
```


```
auto res = std::ranges::copy_if(std::execution::par, data, out, [](auto x) { return (x & 0x01) == 1; });
```



# copy\_if example

```
std::vector data{1,1,2,3,5,8,8,8,4,9,7,2,2,2,2,3};  
std::vector<int> out(data.size());
```

```
auto res = std::ranges::copy_if(std::execution::par, data, out, [](auto x) { return (x & 0x01) == 1; });
```



# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR, //, Pred, Proj>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
    copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {}) {

}
}
```

# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR, //, Pred, Proj>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
    copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {}) {
    std::vector<std::size_t> mask(std::ranges::size(r));

    std::ranges::transform(exec, r, mask, [pred = std::move(pred), proj = std::move(proj)](auto x) {
        return std::size_t(std::invoke(pred, std::invoke(proj, x))); });

    }
}
```



# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR, //, Pred, Proj>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {}) {
    std::vector<std::size_t> mask(std::ranges::size(r));
    std::vector<std::size_t> scan_result(std::ranges::size(r));

    std::ranges::transform(exec, r, mask, [pred = std::move(pred), proj = std::move(proj)](auto x) {
        return std::size_t(std::invoke(pred, std::invoke(proj, x))); });

    std::exclusive_scan(exec, mask.begin(), mask.end(), scan_result.begin(), 0);

}
```

# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR, //, Pred, Proj>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {}) {
    std::vector<std::size_t> mask(std::ranges::size(r));
    std::vector<std::size_t> scan_result(std::ranges::size(r));

    std::ranges::transform(exec, r, mask, [pred = std::move(pred), proj = std::move(proj)](auto x) {
        return std::size_t(std::invoke(pred, std::invoke(proj, x))); });

    std::exclusive_scan(exec, mask.begin(), mask.end(), scan_result.begin(), 0);

    auto zip = std::views::zip(r, mask, scan_result);
    std::ranges::for_each(exec, zip, [&result_r](auto tuple) {
        auto [in, mask, scan_result] = tuple;
        if (mask == 1) result_r[scan_result] = in;
    });
}
```

# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR, //, Pred, Proj>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {}) {
    std::vector<std::size_t> mask(std::ranges::size(r));
    std::vector<std::size_t> scan_result(std::ranges::size(r));

    std::ranges::transform(exec, r, mask, [pred = std::move(pred), proj = std::move(proj)](auto x) {
        return std::size_t(std::invoke(pred, std::invoke(proj, x))); });


    std::exclusive_scan(exec, mask.begin(), mask.end(), scan_result.begin(), 0);

    auto zip = std::views::zip(r, mask, scan_result);
    std::ranges::for_each(exec, zip, [&result_r](auto tuple) {
        auto [in, mask, scan_result] = tuple;
        if (mask == 1) result_r[scan_result] = in;
    });
    auto copied_elements = scan_result.back() + mask.back();
    return {std::ranges::begin(r) + std::ranges::size(r), std::ranges::begin(result_r) + copied_elements};
}
```

# copy\_if example

```
std::vector data{1,1,2,3,5,8,8,8,4,9,7,2,2,2,2,3};  
std::vector<int> out(data.size());
```


```
auto res = std::ranges::copy_if(std::execution::par, data, out, [](auto x) { return (x & 0x01) == 1; });
```



# copy\_if example

```
std::vector data{1,1,2,3,5,8,8,8,4,9,7,2,2,2,2,3};  
std::vector<int> out(7);
```


```
auto res = std::ranges::copy_if(std::execution::par, data, out, [](auto x) { return (x & 0x01) == 1; });
```



# copy\_if example

```
std::vector data{1,1,2,3,5,8,8,8,4,9,7,2,2,2,2,3};  
std::vector<int> out(4);
```


```
auto res = std::ranges::copy_if(std::execution::par, data, out, [](auto x) { return (x & 0x01) == 1; });
```



# copy\_if example

```
std::vector data{1,1,2,3,5,8,8,8,4,9,7,2,2,2,2,3};  
std::vector<int> out(4);
```


```
auto res = std::ranges::copy_if(std::execution::par, data, out, [](auto x) { return (x & 0x01) == 1; });
```



# copy\_if example

```
std::vector data{1,1,2,3,5,8,8,8,4,9,7,2,2,2,2,3};  
std::vector<int> out(4);
```

```
auto res = std::ranges::copy_if(std::execution::par, data, out, [](auto x) { return (x & 0x01) == 1; });
```





# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR, //, Pred, Proj>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {}) {
    std::vector<std::size_t> mask(std::ranges::size(r));
    std::vector<std::size_t> scan_result(std::ranges::size(r));

    std::ranges::transform(exec, r, mask, [pred = std::move(pred), proj = std::move(proj)](auto x) {
        return std::size_t(std::invoke(pred, std::invoke(proj, x))); });
    std::exclusive_scan(exec, mask.begin(), mask.end(), scan_result.begin(), 0);

    auto zip = std::views::zip(r, mask, scan_result);
    std::ranges::for_each(exec, zip, [&result_r](auto tuple) {
        auto [in, mask, scan_result] = tuple;
        if (mask == 1) result_r[scan_result] = in;
    });
    auto copied_elements = scan_result.back() + mask.back();
    return {std::ranges::begin(r) + std::ranges::size(r), std::ranges::begin(result_r) + copied_elements};
}
```

# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR, //, Pred, Proj>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {}) {
    std::vector<std::size_t> mask(std::ranges::size(r));
    std::vector<std::size_t> scan_result(std::ranges::size(r));

    std::ranges::transform(exec, r, mask, [pred = std::move(pred), proj = std::move(proj)](auto x) {
        return std::size_t(std::invoke(pred, std::invoke(proj, x))); });
    std::exclusive_scan(exec, mask.begin(), mask.end(), scan_result.begin(), 0);

    auto last_iterator = std::ranges::upper_bound(scan_result, std::ranges::size(result_r)) - 1;

    auto zip = std::views::zip(r, mask, scan_result);
    std::ranges::for_each(exec, zip, [&result_r](auto tuple) {
        auto [in, mask, scan_result] = tuple;
        if (mask == 1) result_r[scan_result] = in;
    });
    auto copied_elements = scan_result.back() + mask.back();
    return {std::ranges::begin(r) + std::ranges::size(r), std::ranges::begin(result_r) + copied_elements};
}
```

# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR, //, Pred, Proj>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {}) {
    std::vector<std::size_t> mask(std::ranges::size(r));
    std::vector<std::size_t> scan_result(std::ranges::size(r));

    std::ranges::transform(exec, r, mask, [pred = std::move(pred), proj = std::move(proj)](auto x) {
        return std::size_t(std::invoke(pred, std::invoke(proj, x))); });
    std::exclusive_scan(exec, mask.begin(), mask.end(), scan_result.begin(), 0);

    auto last_iterator = std::ranges::upper_bound(scan_result, std::ranges::size(result_r)) - 1;
    bool enough_space = std::ranges::size(result_r) > scan_result.back();

    auto zip = std::views::zip(r, mask, scan_result);
    std::ranges::for_each(exec, zip, [&result_r](auto tuple) {
        auto [in, mask, scan_result] = tuple;
        if (mask == 1) result_r[scan_result] = in;
    });
    auto copied_elements = scan_result.back() + mask.back();
    return {std::ranges::begin(r) + std::ranges::size(r), std::ranges::begin(result_r) + copied_elements};
}
```

# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR, //, Pred, Proj>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {}) {
    std::vector<std::size_t> mask(std::ranges::size(r));
    std::vector<std::size_t> scan_result(std::ranges::size(r));

    std::ranges::transform(exec, r, mask, [pred = std::move(pred), proj = std::move(proj)](auto x) {
        return std::size_t(std::invoke(pred, std::invoke(proj, x))); });
    std::exclusive_scan(exec, mask.begin(), mask.end(), scan_result.begin(), 0);

    auto last_iterator = std::ranges::upper_bound(scan_result, std::ranges::size(result_r)) - 1;
    bool enough_space = std::ranges::size(result_r) > scan_result.back();
    last_iterator += std::size_t(enough_space || (*last_iterator == scan_result.back() && mask.back() == 0));

    auto zip = std::views::zip(r, mask, scan_result);
    std::ranges::for_each(exec, zip, [&result_r](auto tuple) {
        auto [in, mask, scan_result] = tuple;
        if (mask == 1) result_r[scan_result] = in;
    });
    auto copied_elements = scan_result.back() + mask.back();
    return {std::ranges::begin(r) + std::ranges::size(r), std::ranges::begin(result_r) + copied_elements};
}
```

# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR, //, Pred, Proj>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {}) {
    std::vector<std::size_t> mask(std::ranges::size(r));
    std::vector<std::size_t> scan_result(std::ranges::size(r));

    std::ranges::transform(exec, r, mask, [pred = std::move(pred), proj = std::move(proj)](auto x) {
        return std::size_t(std::invoke(pred, std::invoke(proj, x))); });
    std::exclusive_scan(exec, mask.begin(), mask.end(), scan_result.begin(), 0);

    auto last_iterator = std::ranges::upper_bound(scan_result, std::ranges::size(result_r)) - 1;
    bool enough_space = std::ranges::size(result_r) > scan_result.back();
    last_iterator += std::size_t(enough_space || (*last_iterator == scan_result.back() && mask.back() == 0));
    std::size_t distance = std::ranges::distance(scan_result.begin(), last_iterator);

    auto zip = std::views::zip(r | std::views::take(distance), mask, scan_result);
    std::ranges::for_each(exec, zip, [&result_r](auto tuple) {
        auto [in, mask, scan_result] = tuple;
        if (mask == 1) result_r[scan_result] = in;
    });
    auto copied_elements = scan_result.back() + mask.back();
    return {std::ranges::begin(r) + std::ranges::size(r), std::ranges::begin(result_r) + copied_elements};
}
```

# copy\_if parallel implementation

```
template<execution-policy Ep, sized-random-access-range R, sized-random-access-range OutR, //, Pred, Proj>
std::ranges::copy_if_result<std::ranges::borrowed_iterator_t<R>, std::ranges::borrowed_iterator_t<OutR>>
copy_if(Ep&& exec, R&& r, OutR&& result_r, Pred pred, Proj proj = {}) {
    std::vector<std::size_t> mask(std::ranges::size(r));
    std::vector<std::size_t> scan_result(std::ranges::size(r));

    std::ranges::transform(exec, r, mask, [pred = std::move(pred), proj = std::move(proj)](auto x) {
        return std::size_t(std::invoke(pred, std::invoke(proj, x))); });
    std::exclusive_scan(exec, mask.begin(), mask.end(), scan_result.begin(), 0);


    auto last_iterator = std::ranges::upper_bound(scan_result, std::ranges::size(result_r)) - 1;
    bool enough_space = std::ranges::size(result_r) > scan_result.back();
    last_iterator += std::size_t(enough_space || (*last_iterator == scan_result.back() && mask.back() == 0));
    std::size_t distance = std::ranges::distance(scan_result.begin(), last_iterator);

    auto zip = std::views::zip(r | std::views::take(distance), mask, scan_result);
    std::ranges::for_each(exec, zip, [&result_r](auto tuple) {
        auto [in, mask, scan_result] = tuple;
        if (mask == 1) result_r[scan_result] = in;
    });
    auto copied_elements = scan_result[distance - 1] + mask[distance - 1];
    return {std::ranges::begin(r) + distance, std::ranges::begin(result_r) + copied_elements};
}
```

# copy\_if example

```
std::vector data{1,1,2,3,5,8,8,8,4,9,7,2,2,2,2};  
std::vector<int> out(6);
```

```
auto res = std::ranges::copy_if(std::execution::par, data, out, [](auto x) { return (x & 0x01) == 1; });
```



input

mask

scan of the mask

output

# for\_each (Category - return type)

API	Return type
<code>std::for_each</code>	Function
Parallel <code>std::for_each</code>	<code>void</code>
<code>std::for_each_n</code>	Iterator
Parallel <code>std::for_each_n</code>	Iterator
<code>std::ranges::for_each</code>	<code>for_each_result&lt;ranges::borrowed_iterator_t&lt;Range&gt;, Function&gt;</code>
<code>std::ranges::for_each, l + S overload</code>	<code>for_each_result&lt;Iterator, Function&gt;</code>
<code>std::ranges::for_each_n</code>	<code>for_each_n_result&lt;Iterator, Function&gt;</code>



# for\_each

API	Return type
Parallel <code>std::ranges::for_each</code>	<code>ranges::borrowed_iterator_t&lt;Range&gt;</code>
Parallel <code>std::ranges::for_each</code> , I + S overload	Iterator
Parallel <code>std::ranges::for_each_n</code>	Iterator

# Heterogeneity with oneDPL

## oneAPI DPC++ library (oneDPL):

- Implementation of standard Parallel Algorithms made by Intel
- Evolution of former Parallel STL
  - Donated to LLVM, used as GNU libstc++ parallel algorithms implementation
- Supports heterogeneous execution via SYCL
- Currently a part of UXL Foundation: both specification and source code
- ~40 parallel range algorithms available as experimental (hetero only)
- ~20 parallel range algorithms available as a product quality

# Heterogeneous example

```
auto pipeline = data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse;  
auto res = std::find_if(std::execution::par, pipeline, pred);
```

# Heterogeneous example

```
// Has an associated device  
sycl::queue q;
```

```
auto pipeline = data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse;  
auto res = std::find_if(std::execution::par, pipeline, pred);
```

# Heterogeneous example

```
// Has an associated device
```

```
sycl::queue q;
```

```
// Accessible memory on both the host and the device
```

```
using alloc_type = sycl::usm_allocator<std::size_t, sycl::usm::alloc::shared>;
```

```
// Creating an allocator object
```

```
alloc_type alloc(q);
```

```
auto pipeline = data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse;
```

```
auto res = std::find_if(std::execution::par, pipeline, pred);
```

# Heterogeneous example

```
// Has an associated device
sycl::queue q;

// Accessible memory on both the host and the device
using alloc_type = sycl::usm_allocator<std::size_t, sycl::usm::alloc::shared>;

// Creating an allocator object
alloc_type alloc(q);

// Creating data
std::vector<std::size_t, alloc_type> v(size, alloc);
std::ranges::subrange data{v.begin(), v.end()};

auto pipeline = data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse;
auto res = std::find_if(std::execution::par, pipeline, pred);
```

# Heterogeneous example

```
// Has an associated device
sycl::queue q;

// Accessible memory on both the host and the device
using alloc_type = sycl::usm_allocator<std::size_t, sycl::usm::alloc::shared>;

// Creating an allocator object
alloc_type alloc(q);


// Creating data
std::vector<std::size_t, alloc_type> v(size, alloc);
std::ranges::subrange data{v.begin(), v.end()};

auto pipeline = data | std::views::transform([](auto i){ return i + 1; }) | std::views::reverse;
auto res = dpl::find_if(dpl::make_device_policy(q), pipeline, pred);
```

# Performance

## Speedup Over Sequential (higher is better)

16M Elements, Target Element in Center of Range  
100 Iterations, Median, icpx 2025.1.0





# Scope and status

# 79 algorithms in namespace std already with policies

all_of	search[_n]	remove_copy	is_sorted	is_heap
any_of	copy[_n]	remove_copy_if	is_sorted_until	is_heap_until
none_of	copy_if	unique	nth_element	min_element
for_each[_n]	move	unique_copy	is_partitioned	max_element
find	swap_ranges	reverse	partition	minmax_element
find_if	transform	reverse_copy	stable_partition	lexicographical_compare
find_if_not	replace	rotate	partition_copy	uninitialized_default_construct[_n]
find_end	replace_if	rotate_copy	merge	uninitialized_value_construct[_n]
find_first_of	replace_copy	shift_left	inplace_merge	uninitialized_copy[_n]
adjacent_find	replace_copy_if	shift_right	includes	uninitialized_move[_n]
count	fill[_n]	sort	set_union	uninitialized_fill[_n]
count_if	generate[_n]	stable_sort	set_intersection	destroy[_n]
mismatch	remove	partial_sort	set_difference	
equal	remove_if	partial_sort_copy	set_symmetric_difference	

# 10 algorithms only in namespace `std::ranges`

<code>std::ranges</code> algorithms to add policies	<code>std</code> algorithms used as the guidance
<code>contains</code>	<i>find</i>
<code>contains_subrange</code>	<i>search</i>
<code>find_last</code>	<i>find</i>
<code>find_last_if</code>	<i>find_if</i>
<code>find_last_if_not</code>	<i>find_if_not</i>
<code>starts_with</code>	<i>mismatch</i>
<code>ends_with</code>	<i>equal</i>
<code>min</code>	<i>min_element</i>
<code>max</code>	<i>max_element</i>
<code>minmax</code>	<i>minmax_element</i>

# Out of scope algorithms

- All algorithms that do not have an **ExecutionPolicy** in their C++17 counterpart
- Algorithms from `<numeric>`
- Algorithms only in namespace `std::ranges` other than mentioned before:
  - `fold` algorithm family
  - `generate_random`

# “Algorithms with output” category

all_of	search[_n]	remove_copy	is_sorted	is_heap
any_of	copy[_n]	remove_copy_if	is_sorted_until	is_heap_until
none_of	copy_if	unique	nth_element	min_element
for_each[_n]	move	unique_copy	is_partitioned	max_element
find	swap_ranges	reverse	partition	minmax_element
find_if	transform	reverse_copy	stable_partition	lexicographical_compare
find_if_not	replace	rotate	partition_copy	uninitialized_default_construct[_n]
find_end	replace_if	rotate_copy	merge	uninitialized_value_construct[_n]
find_first_of	replace_copy	shift_left	inplace_merge	uninitialized_copy[_n]
adjacent_find	replace_copy_if	shift_right	includes	uninitialized_move[_n]
count	fill[_n]	sort	set_union	uninitialized_fill[_n]
count_if	generate[_n]	stable_sort	set_intersection	destroy[_n]
mismatch	remove	partial_sort	set_difference	
equal	remove_if	partial_sort_copy	set_symmetric_difference	

# “Special algorithms” category

all_of	search[_n]	remove_copy	is_sorted	is_heap
any_of	copy[_n]	remove_copy_if	is_sorted_until	is_heap_until
none_of	copy_if	unique	nth_element	min_element
for_each[_n]	move	unique_copy	is_partitioned	max_element
find	swap_ranges	reverse	partition	minmax_element
find_if	transform	reverse_copy	stable_partition	lexicographical_compare
find_if_not	replace	rotate	partition_copy	uninitialized_default_construct[_n]
find_end	replace_if	rotate_copy	merge	uninitialized_value_construct[_n]
find_first_of	replace_copy	shift_left	inplace_merge	uninitialized_copy[_n]
adjacent_find	replace_copy_if	shift_right	includes	uninitialized_move[_n]
count	fill[_n]	sort	set_union	uninitialized_fill[_n]
count_if	generate[_n]	stable_sort	set_intersection	destroy[_n]
mismatch	remove	partial_sort	set_difference	
equal	remove_if	partial_sort_copy	set_symmetric_difference	

# “Algorithms with output and gaps” category

all_of	search[_n]	remove_copy	is_sorted	is_heap
any_of	copy[_n]	remove_copy_if	is_sorted_until	is_heap_until
none_of	copy_if	unique	nth_element	min_element
for_each[_n]	move	unique_copy	is_partitioned	max_element
find	swap_ranges	reverse	partition	minmax_element
find_if	transform	reverse_copy	stable_partition	lexicographical_compare
find_if_not	replace	rotate	partition_copy	uninitialized_default_construct[_n]
find_end	replace_if	rotate_copy	merge	uninitialized_value_construct[_n]
find_first_of	replace_copy	shift_left	inplace_merge	uninitialized_copy[_n]
adjacent_find	replace_copy_if	shift_right	includes	uninitialized_move[_n]
count	fill[_n]	sort	set_union	uninitialized_fill[_n]
count_if	generate[_n]	stable_sort	set_intersection	destroy[_n]
mismatch	remove	partial_sort	set_difference	
equal	remove_if	partial_sort_copy	set_symmetric_difference	

# “Return type” category

all_of	search[_n]	remove_copy	is_sorted	is_heap
any_of	copy[_n]	remove_copy_if	is_sorted_until	is_heap_until
none_of	copy_if	unique	nth_element	min_element
for_each[_n]	move	unique_copy	is_partitioned	max_element
find	swap_ranges	reverse	partition	minmax_element
find_if	transform	reverse_copy	stable_partition	lexicographical_compare
find_if_not	replace	rotate	partition_copy	uninitialized_default_construct[_n]
find_end	replace_if	rotate_copy	merge	uninitialized_value_construct[_n]
find_first_of	replace_copy	shift_left	inplace_merge	uninitialized_copy[_n]
adjacent_find	replace_copy_if	shift_right	includes	uninitialized_move[_n]
count	fill[_n]	sort	set_union	uninitialized_fill[_n]
count_if	generate[_n]	stable_sort	set_intersection	destroy[_n]
mismatch	remove	partial_sort	set_difference	
equal	remove_if	partial_sort_copy	set_symmetric_difference	



# oneDPL: implemented

<b>all_of</b>	<b>search[_n]</b>	remove_copy	<b>is_sorted</b>	is_heap
<b>any_of</b>	<b>copy[_n]</b>	remove_copy_if	is_sorted_until	is_heap_until
<b>none_of</b>	<b>copy_if</b>	unique	nth_element	<b>min_element</b>
<b>for_each[_n]</b>	move	unique_copy	is_partitioned	<b>max_element</b>
<b>find</b>	swap_ranges	reverse	partition	minmax_element
<b>find_if</b>	<b>transform</b>	reverse_copy	stable_partition	lexicographical_compare
<b>find_if_not</b>	replace	rotate	partition_copy	uninitialized_default_construct[_n]
find_end	replace_if	rotate_copy	<b>merge</b>	uninitialized_value_construct[_n]
find_first_of	replace_copy	shift_left	inplace_merge	uninitialized_copy[_n]
<b>adjacent_find</b>	replace_copy_if	shift_right	includes	uninitialized_move[_n]
<b>count</b>	fill[_n]	<b>sort</b>	set_union	uninitialized_fill[_n]
<b>count_if</b>	generate[_n]	<b>stable_sort</b>	set_intersection	destroy[_n]
mismatch	remove	partial_sort	set_difference	
<b>equal</b>	remove_if	partial_sort_copy	set_symmetric_difference	

# oneDPL: implemented + coming

all_of	search[_n]	remove_copy	is_sorted	is_heap
any_of	copy[_n]	remove_copy_if	is_sorted_until	is_heap_until
none_of	copy_if	unique	nth_element	min_element
for_each[_n]	move	unique_copy	is_partitioned	max_element
find	swap_ranges	reverse	partition	minmax_element
find_if	transform	reverse_copy	stable_partition	lexicographical_compare
find_if_not	replace	rotate	partition_copy	uninitialized_default_construct[_n]
find_end	replace_if	rotate_copy	merge	uninitialized_value_construct[_n]
find_first_of	replace_copy	shift_left	inplace_merge	uninitialized_copy[_n]
adjacent_find	replace_copy_if	shift_right	includes	uninitialized_move[_n]
count	fill[_n]	sort	set_union	uninitialized_fill[_n]
count_if	generate[_n]	stable_sort	set_intersection	destroy[_n]
mismatch	remove	partial_sort	set_difference	
equal	remove_if	partial_sort_copy	set_symmetric_difference	

# oneDPL: coming

std::ranges algorithms to add policies	std algorithms used as the guidance
contains	<i>find</i>
contains_subrange	<i>search</i>
find_last	<i>find</i>
find_last_if	<i>find_if</i>
find_last_if_not	<i>find_if_not</i>
starts_with	<i>mismatch</i>
ends_with	<i>equal</i>
min	<i>min_element</i>
max	<i>max_element</i>
minmax	<i>minmax_element</i>

# P3179 proposal status

- Design approved
  - Successfully passed SG1: Concurrency and Parallelism, SG9: Ranges, and Library Evolution Working Group (LEWG)
- Is actively reviewed in Library Working Group (LWG)
- Has an opportunity to be included in C++26

# Further work

- Numeric range algorithms
- Synchronous parallel algorithms and Senders/Receivers (P2500)
- Asynchronous parallel algorithms (P3300)
- Stretch goal: Range-as-the-output for serial range algorithms

# Useful links

- Parallel Range Algorithms proposal:  
<https://wg21.link/P3179>
- Range-as-the-output paper:  
<https://wg21.link/P3490>
- oneDPL source code:  
<https://github.com/uxlfoundation/oneDPL>
- oneDPL specification:  
<https://oneapi-spec.uxlfoundation.org/specifications/oneapi/latest/elements/onedpl/source/>
- SYCL specification:  
<https://registry.khronos.org/SYCL/specs/sycl-2020/html/sycl-2020.html>

# Special thanks

- Alexey Kukanov

# Special thanks

- Alexey Kukanov
- Jonathan Muller



# Special thanks

- Alexey Kukanov
- Jonathan Muller
- Inbal Levi

# Special thanks

- Alexey Kukanov
- Jonathan Muller
- Inbal Levi
- Jeff Garland and Jonathan Wakely

# Special thanks

- Alexey Kukanov
- Jonathan Muller
- Inbal Levi
- Jeff Garland and Jonathan Wakely
- Zach Laine and Jonathan Muller for reviewing the abstract

# Special thanks

- Alexey Kukanov
- Jonathan Muller
- Inbal Levi
- Jeff Garland and Jonathan Wakely
- Zach Laine and Jonathan Muller for reviewing the abstract
  - Feedback: “The abstract looks good as is” 😊

# Bonus

# Count words example with C++ 17 parallel algorithms

```
int word_count(const std::string& text) {  
  
    // check for empty string  
    if (text.empty()) return 0;  
  
    // compute the number characters that start a new word  
    int wc = std::transform_reduce(std::execution::par,  
        text.begin(), text.end() - 1,           // sequence of left characters  
        text.begin() + 1,                       // sequence of right characters  
        0,                                       // initial value  
        [](int s1, int s2) { return s1 + s2; }, // sum values together  
        [](char s1, char s2) {  
            return int(!std::isalpha(s1)  
                && std::isalpha(s2)); // check if the right character starts the word  
        });  
  
    // if the first character is alphabetical, then it also begins a word  
    if (std::isalpha(*text.begin())) ++wc;  
  
    return wc;  
}
```

# Count words example with parallel range algorithms

```
int word_count(const std::string& text) {  
    using namespace std;  
    // check for empty string  
    if (text.empty()) return 0;  
  
    // compute the number characters that start a new word  
    int wc = ranges::count_if(std::execution::par,  
        views::zip(text | views::take(text.size() - 1), text | views::drop(1)),  
        [](auto v) {  
            auto [s1, s2] = v;  
            return !std::isalpha(s1)  
                && std::isalpha(s2);  
        });  
  
    // if the first character is alphabetical, then it also begins a word  
    if (std::isalpha(*text.begin())) wc++;  
  
    return wc;  
}
```

The Intel logo is centered on a solid blue background. It features the word "intel" in a white, lowercase, sans-serif typeface. A small, bright blue square is positioned above the first vertical stroke of the letter 'i'. To the right of the word "intel" is a small white registered trademark symbol (®).

intel®